

*FCEN, UBA*

---

# Inferencia Bayesian

---

Cecilia Garraffo

---

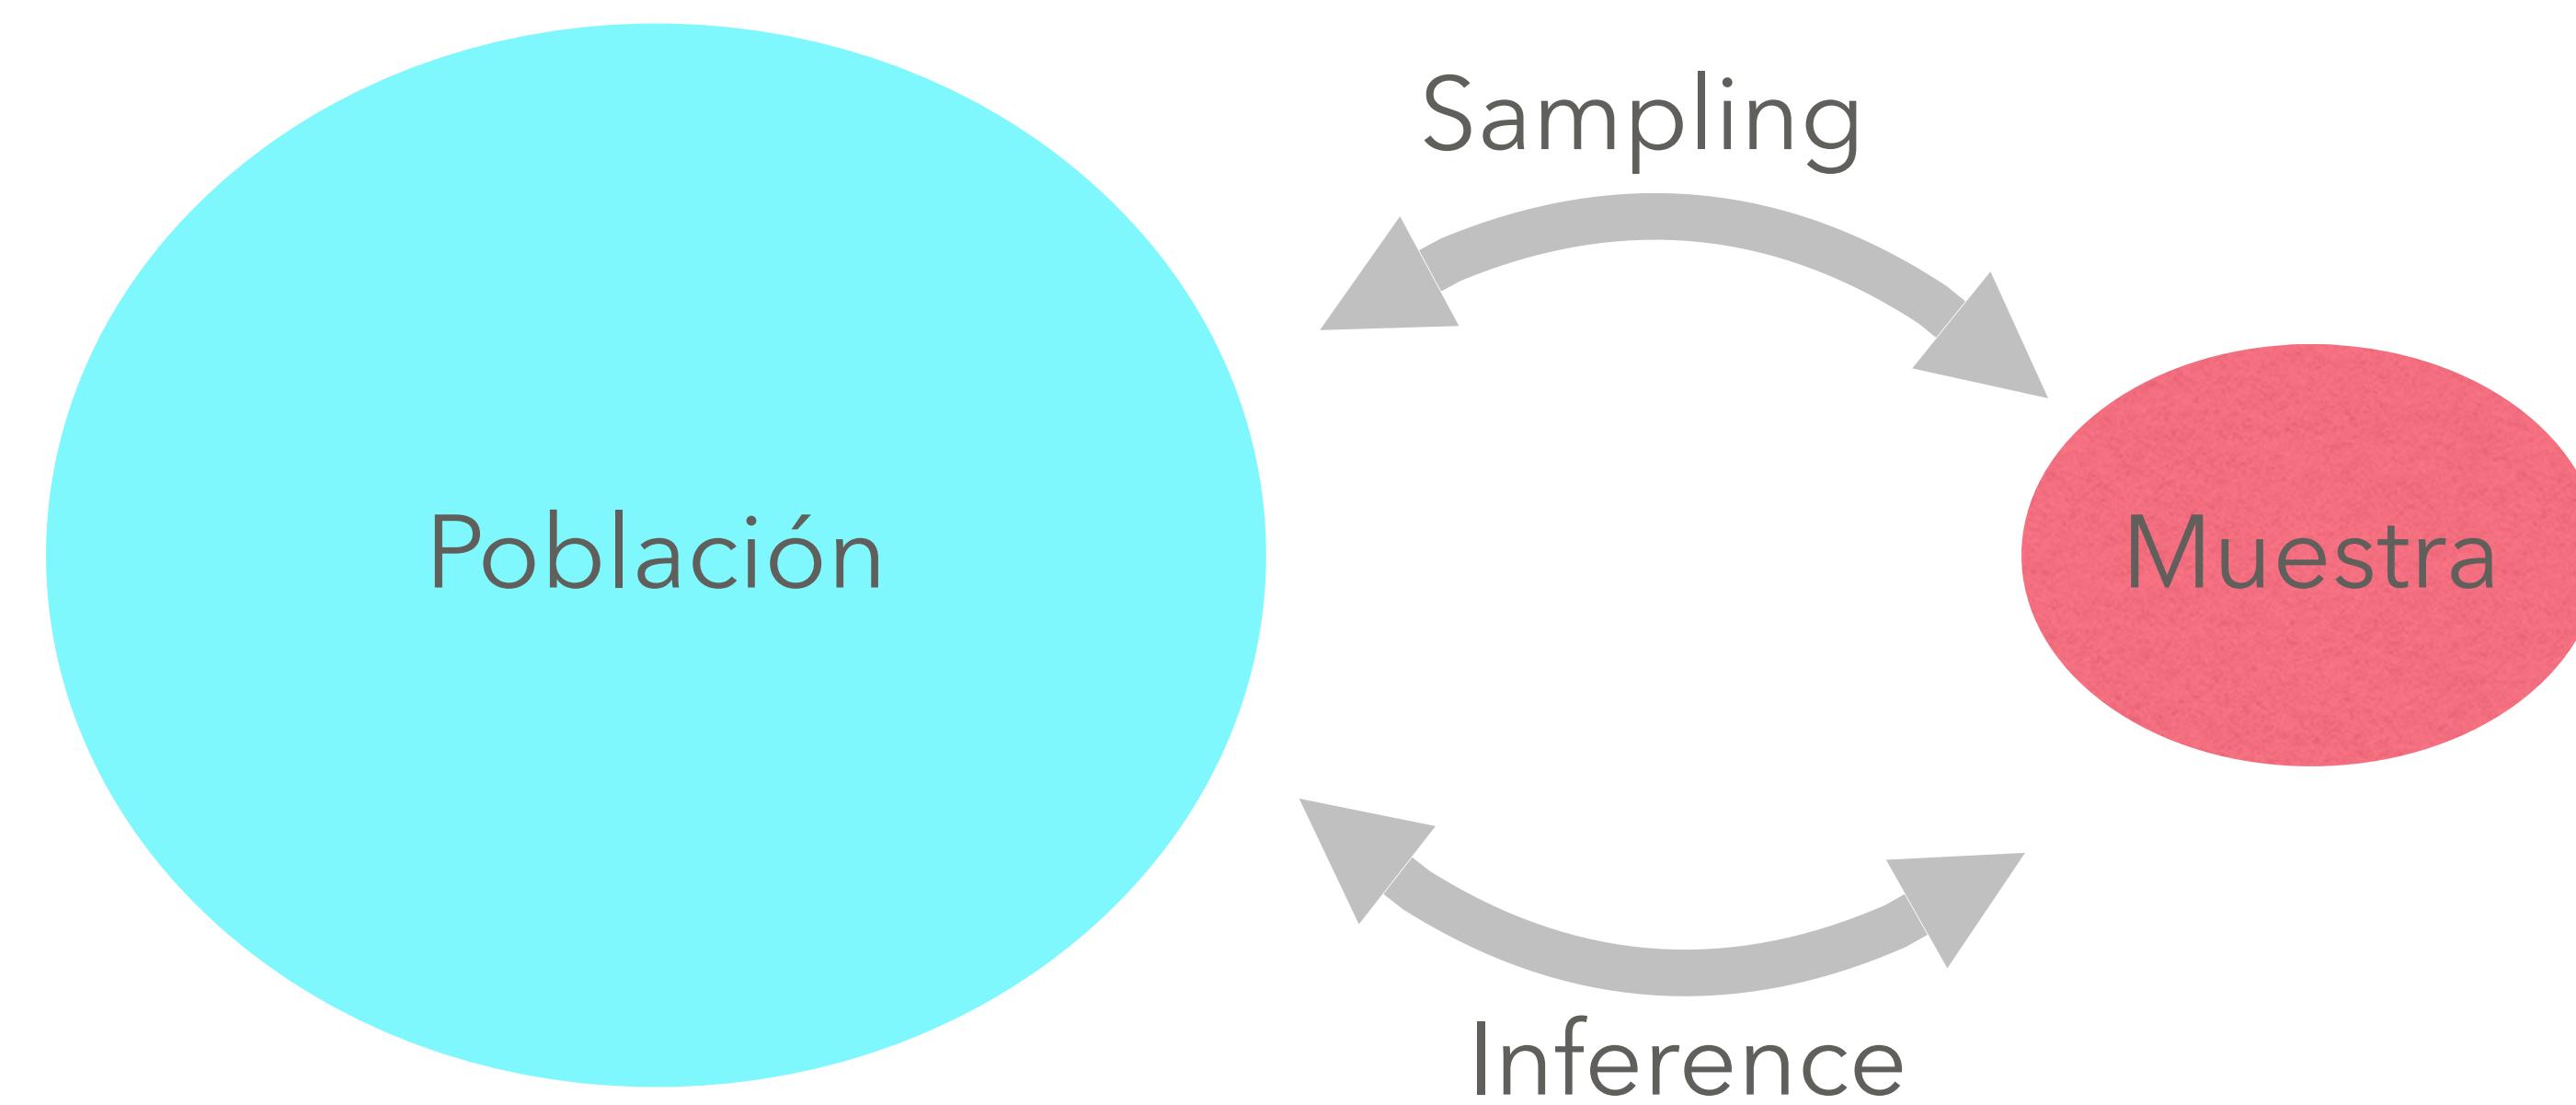
# Outline

---

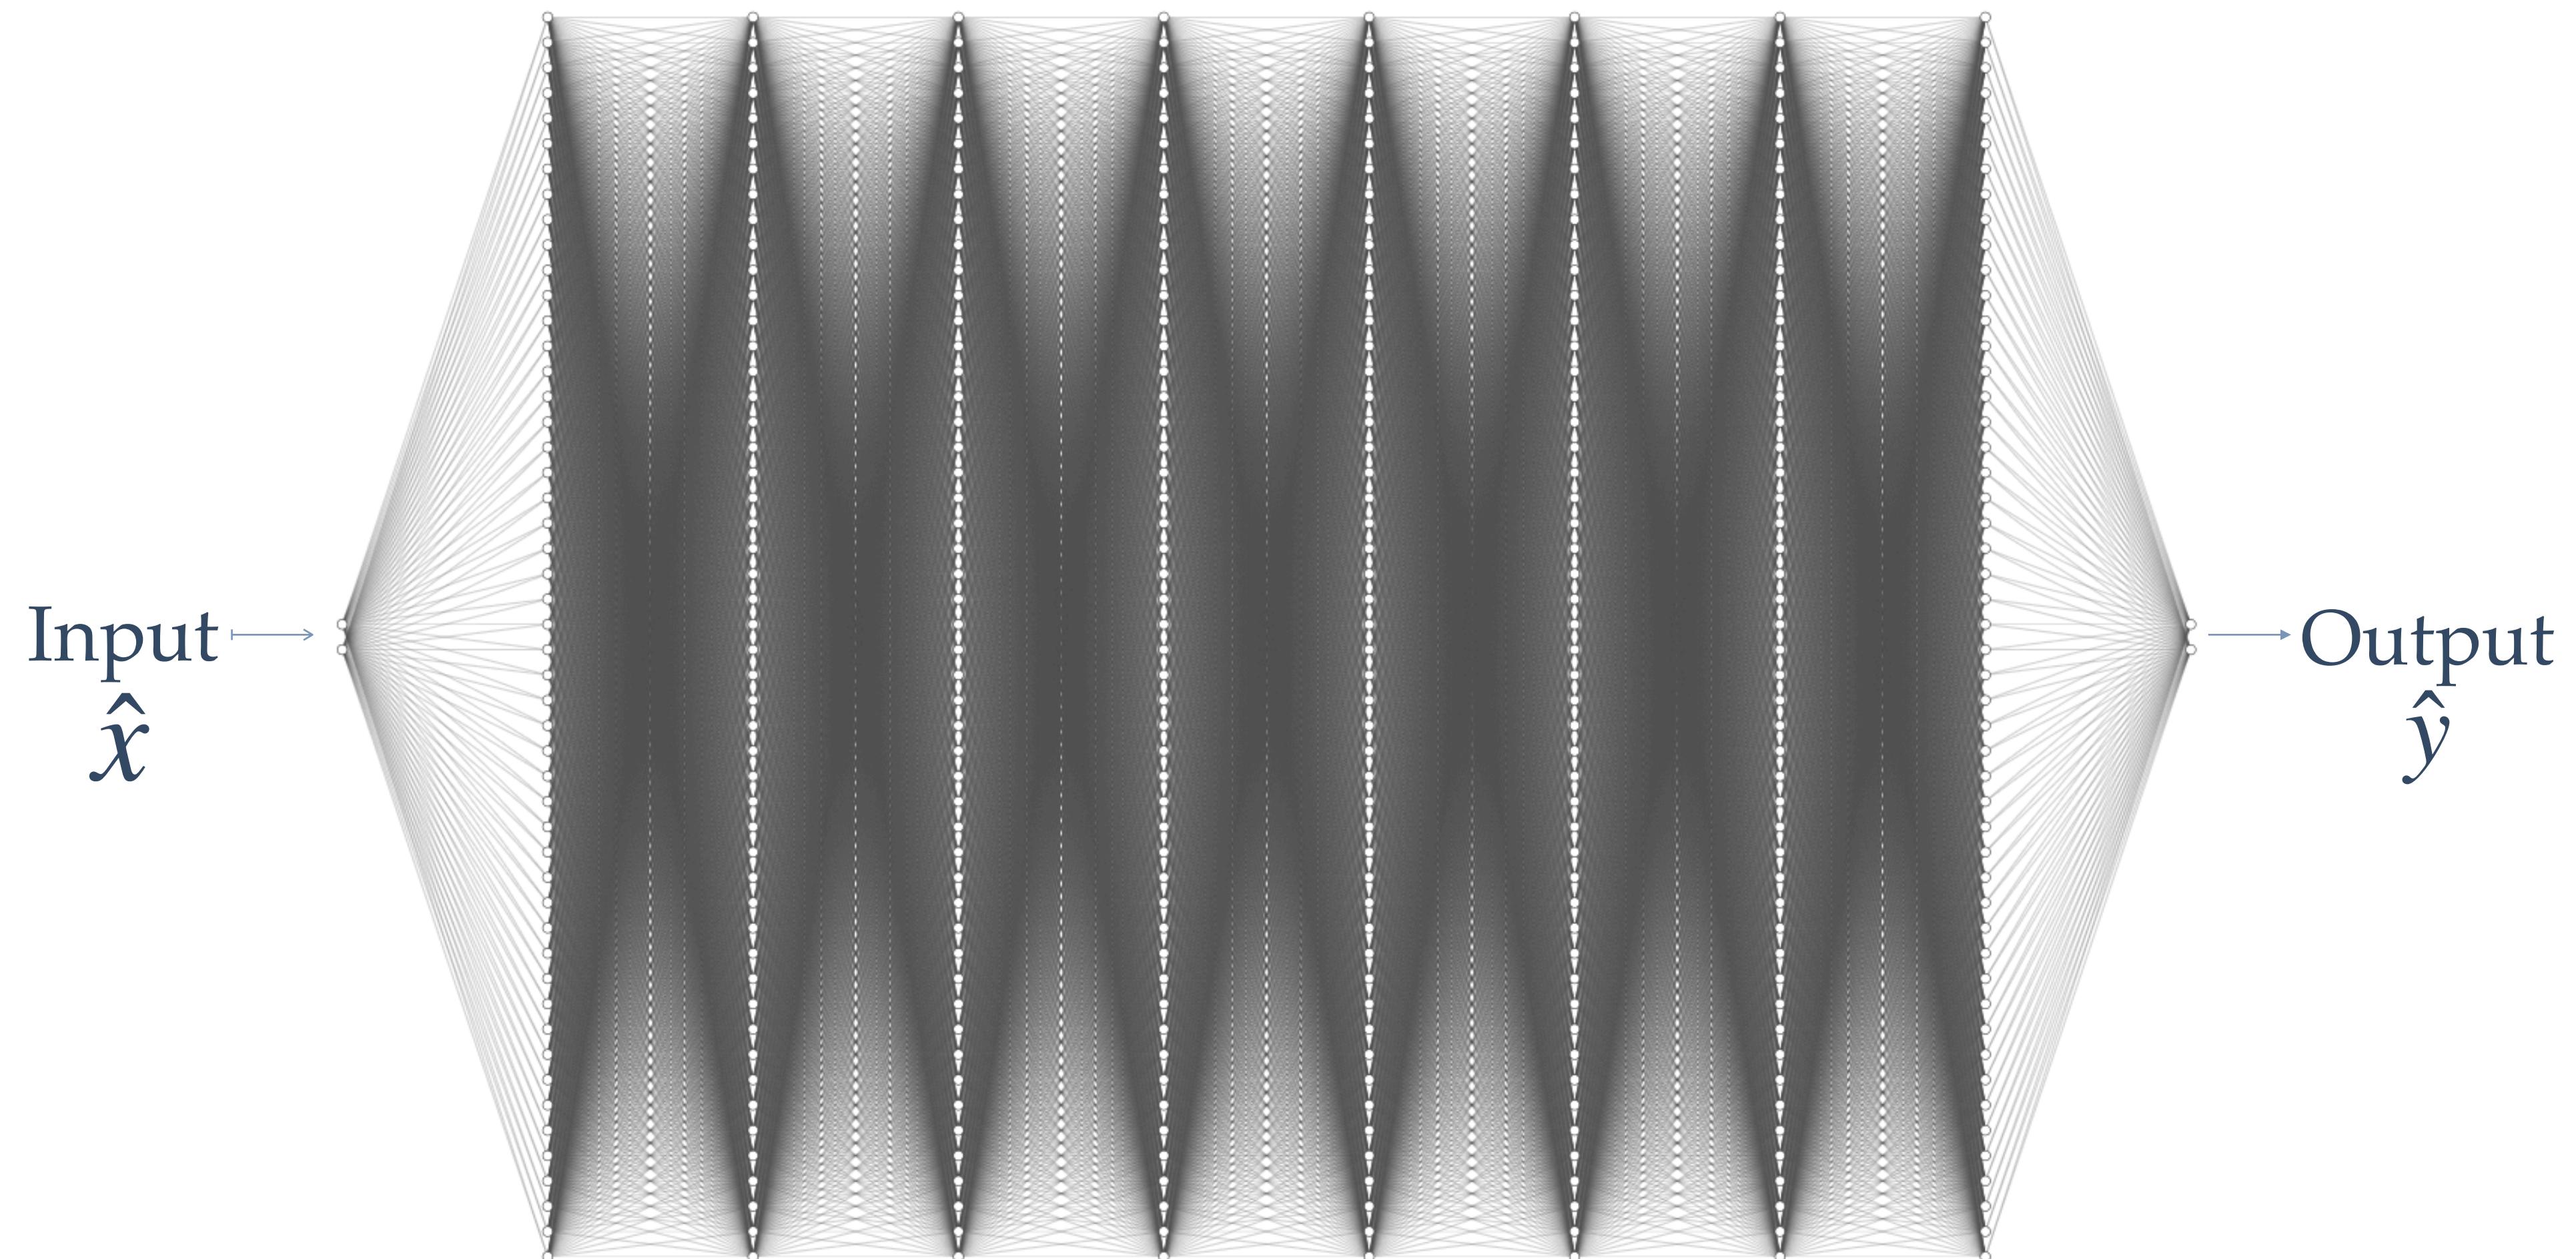
- ❖ Inferencia Bayesian
- ❖ Markov Chain Monte Carlo
- ❖ Inferencia Variacional
- ❖ Redes Neuronales Bayesianas
- ❖ Bayes con backprop y flipout
- ❖ Mean Field Variational Bayes

# Inferencia Estadística

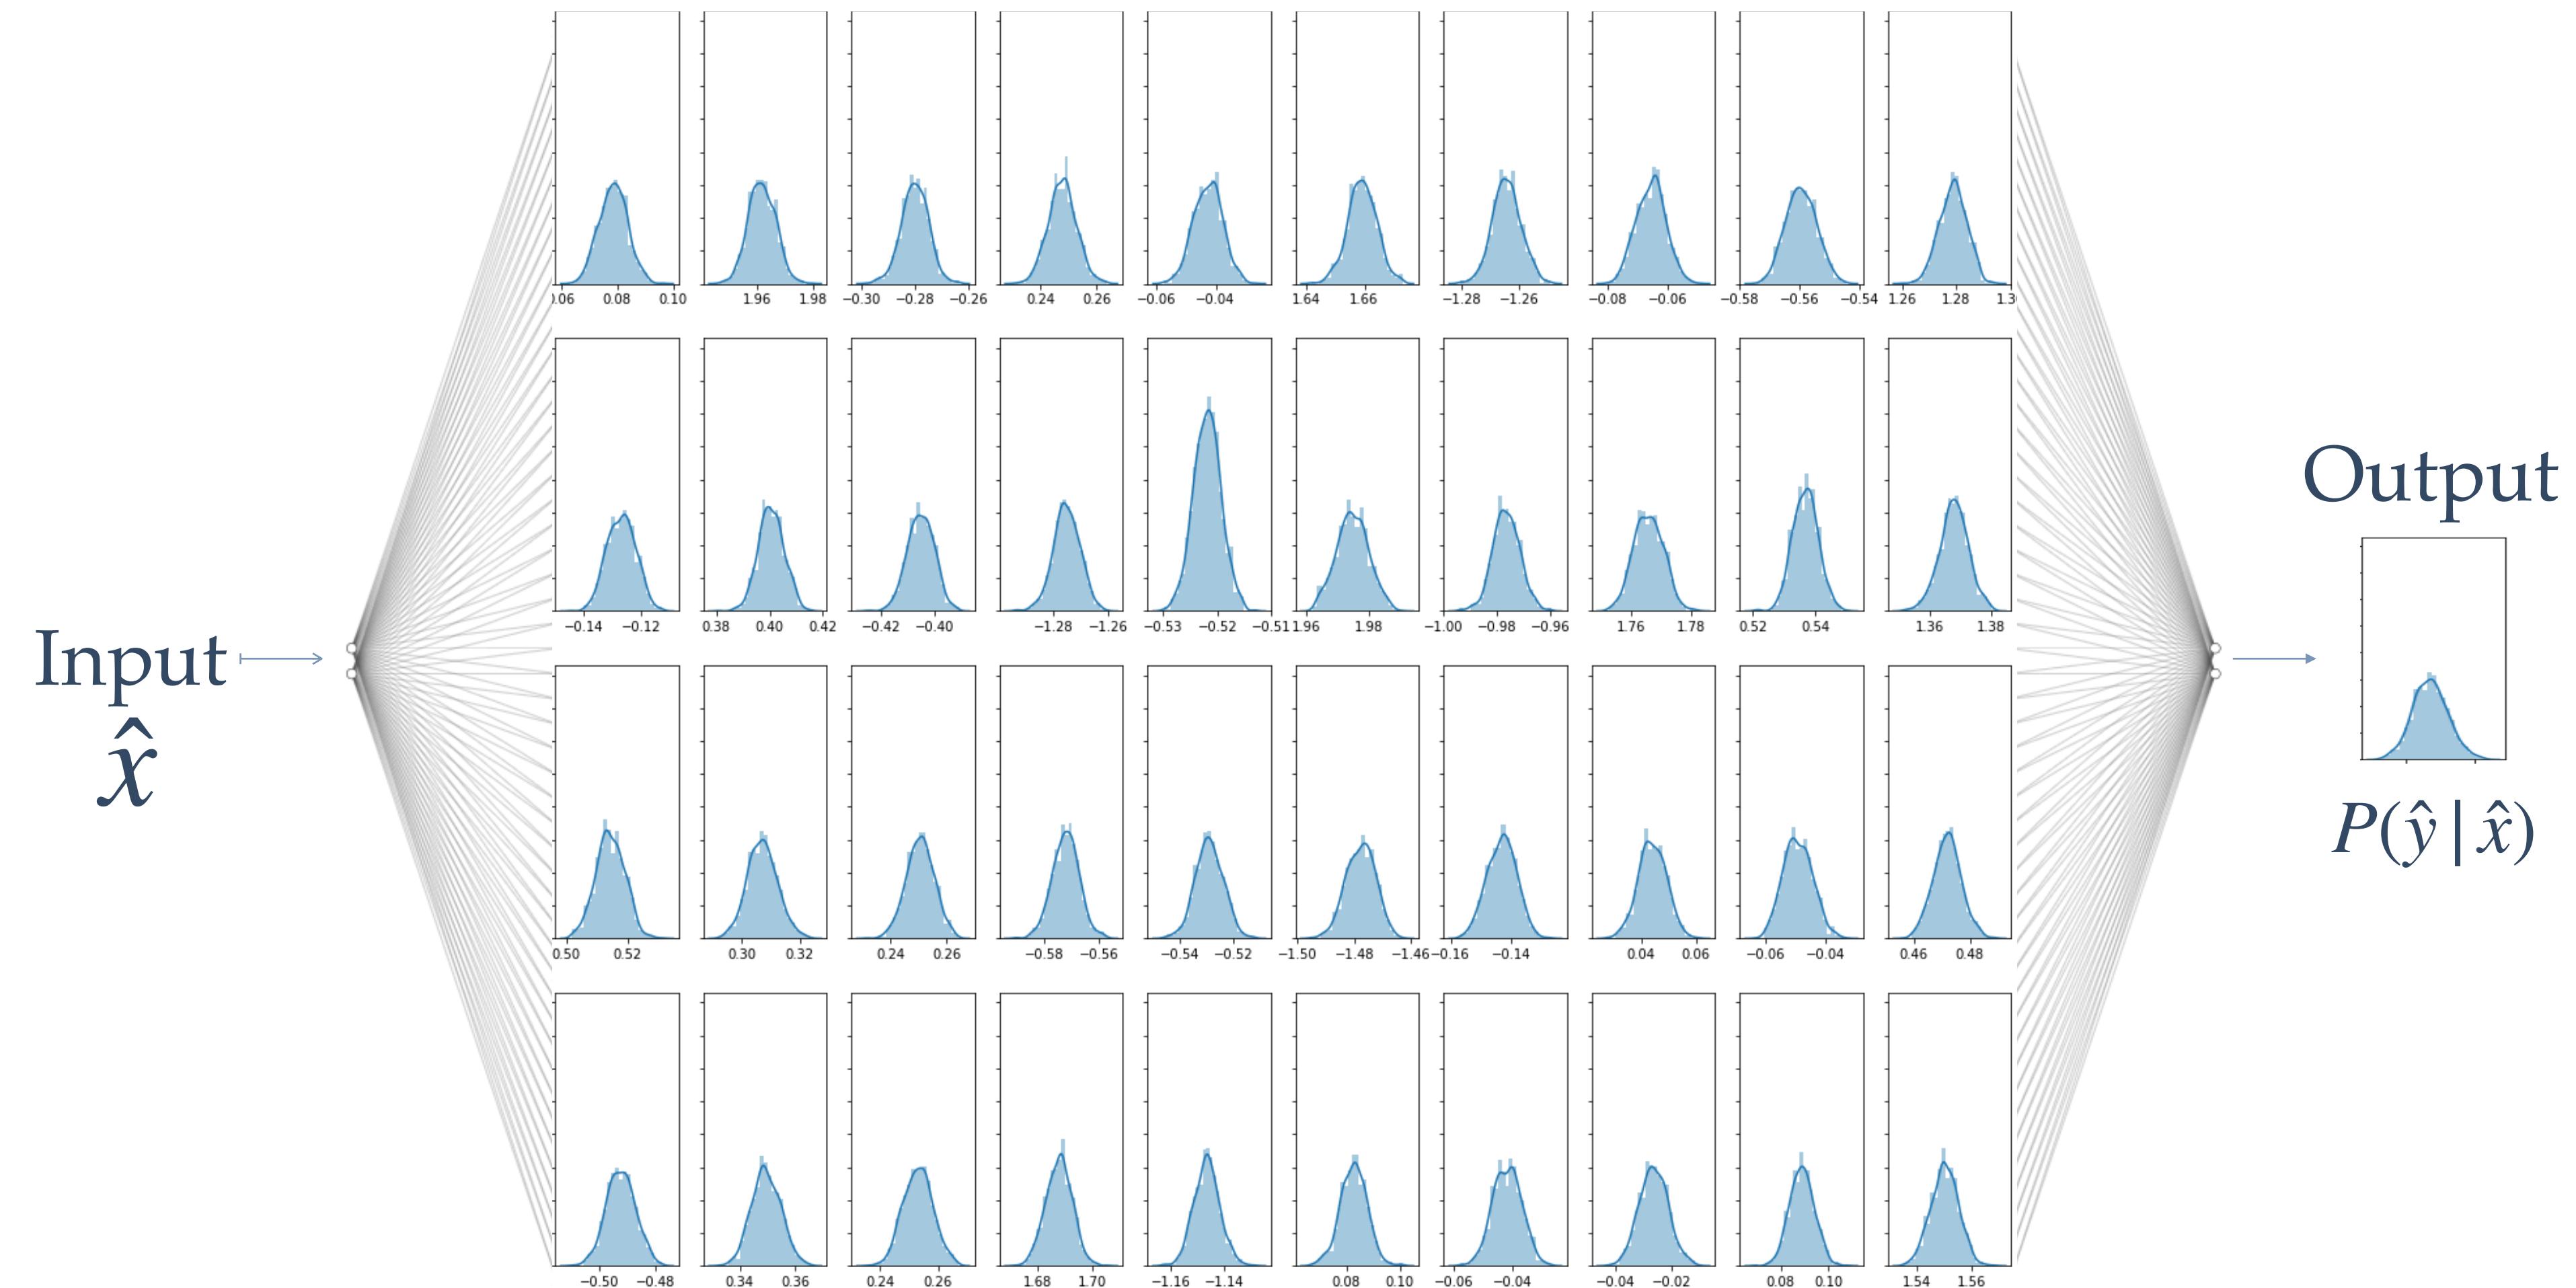
Sacar conclusiones de una distribución de probabilidades a partir de una muestra



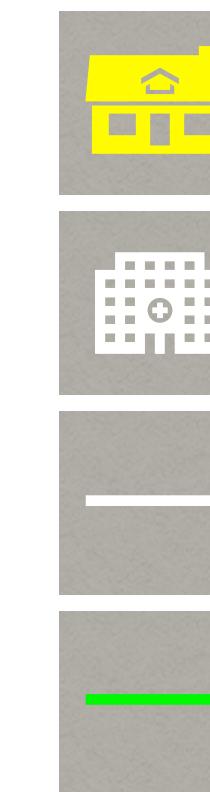
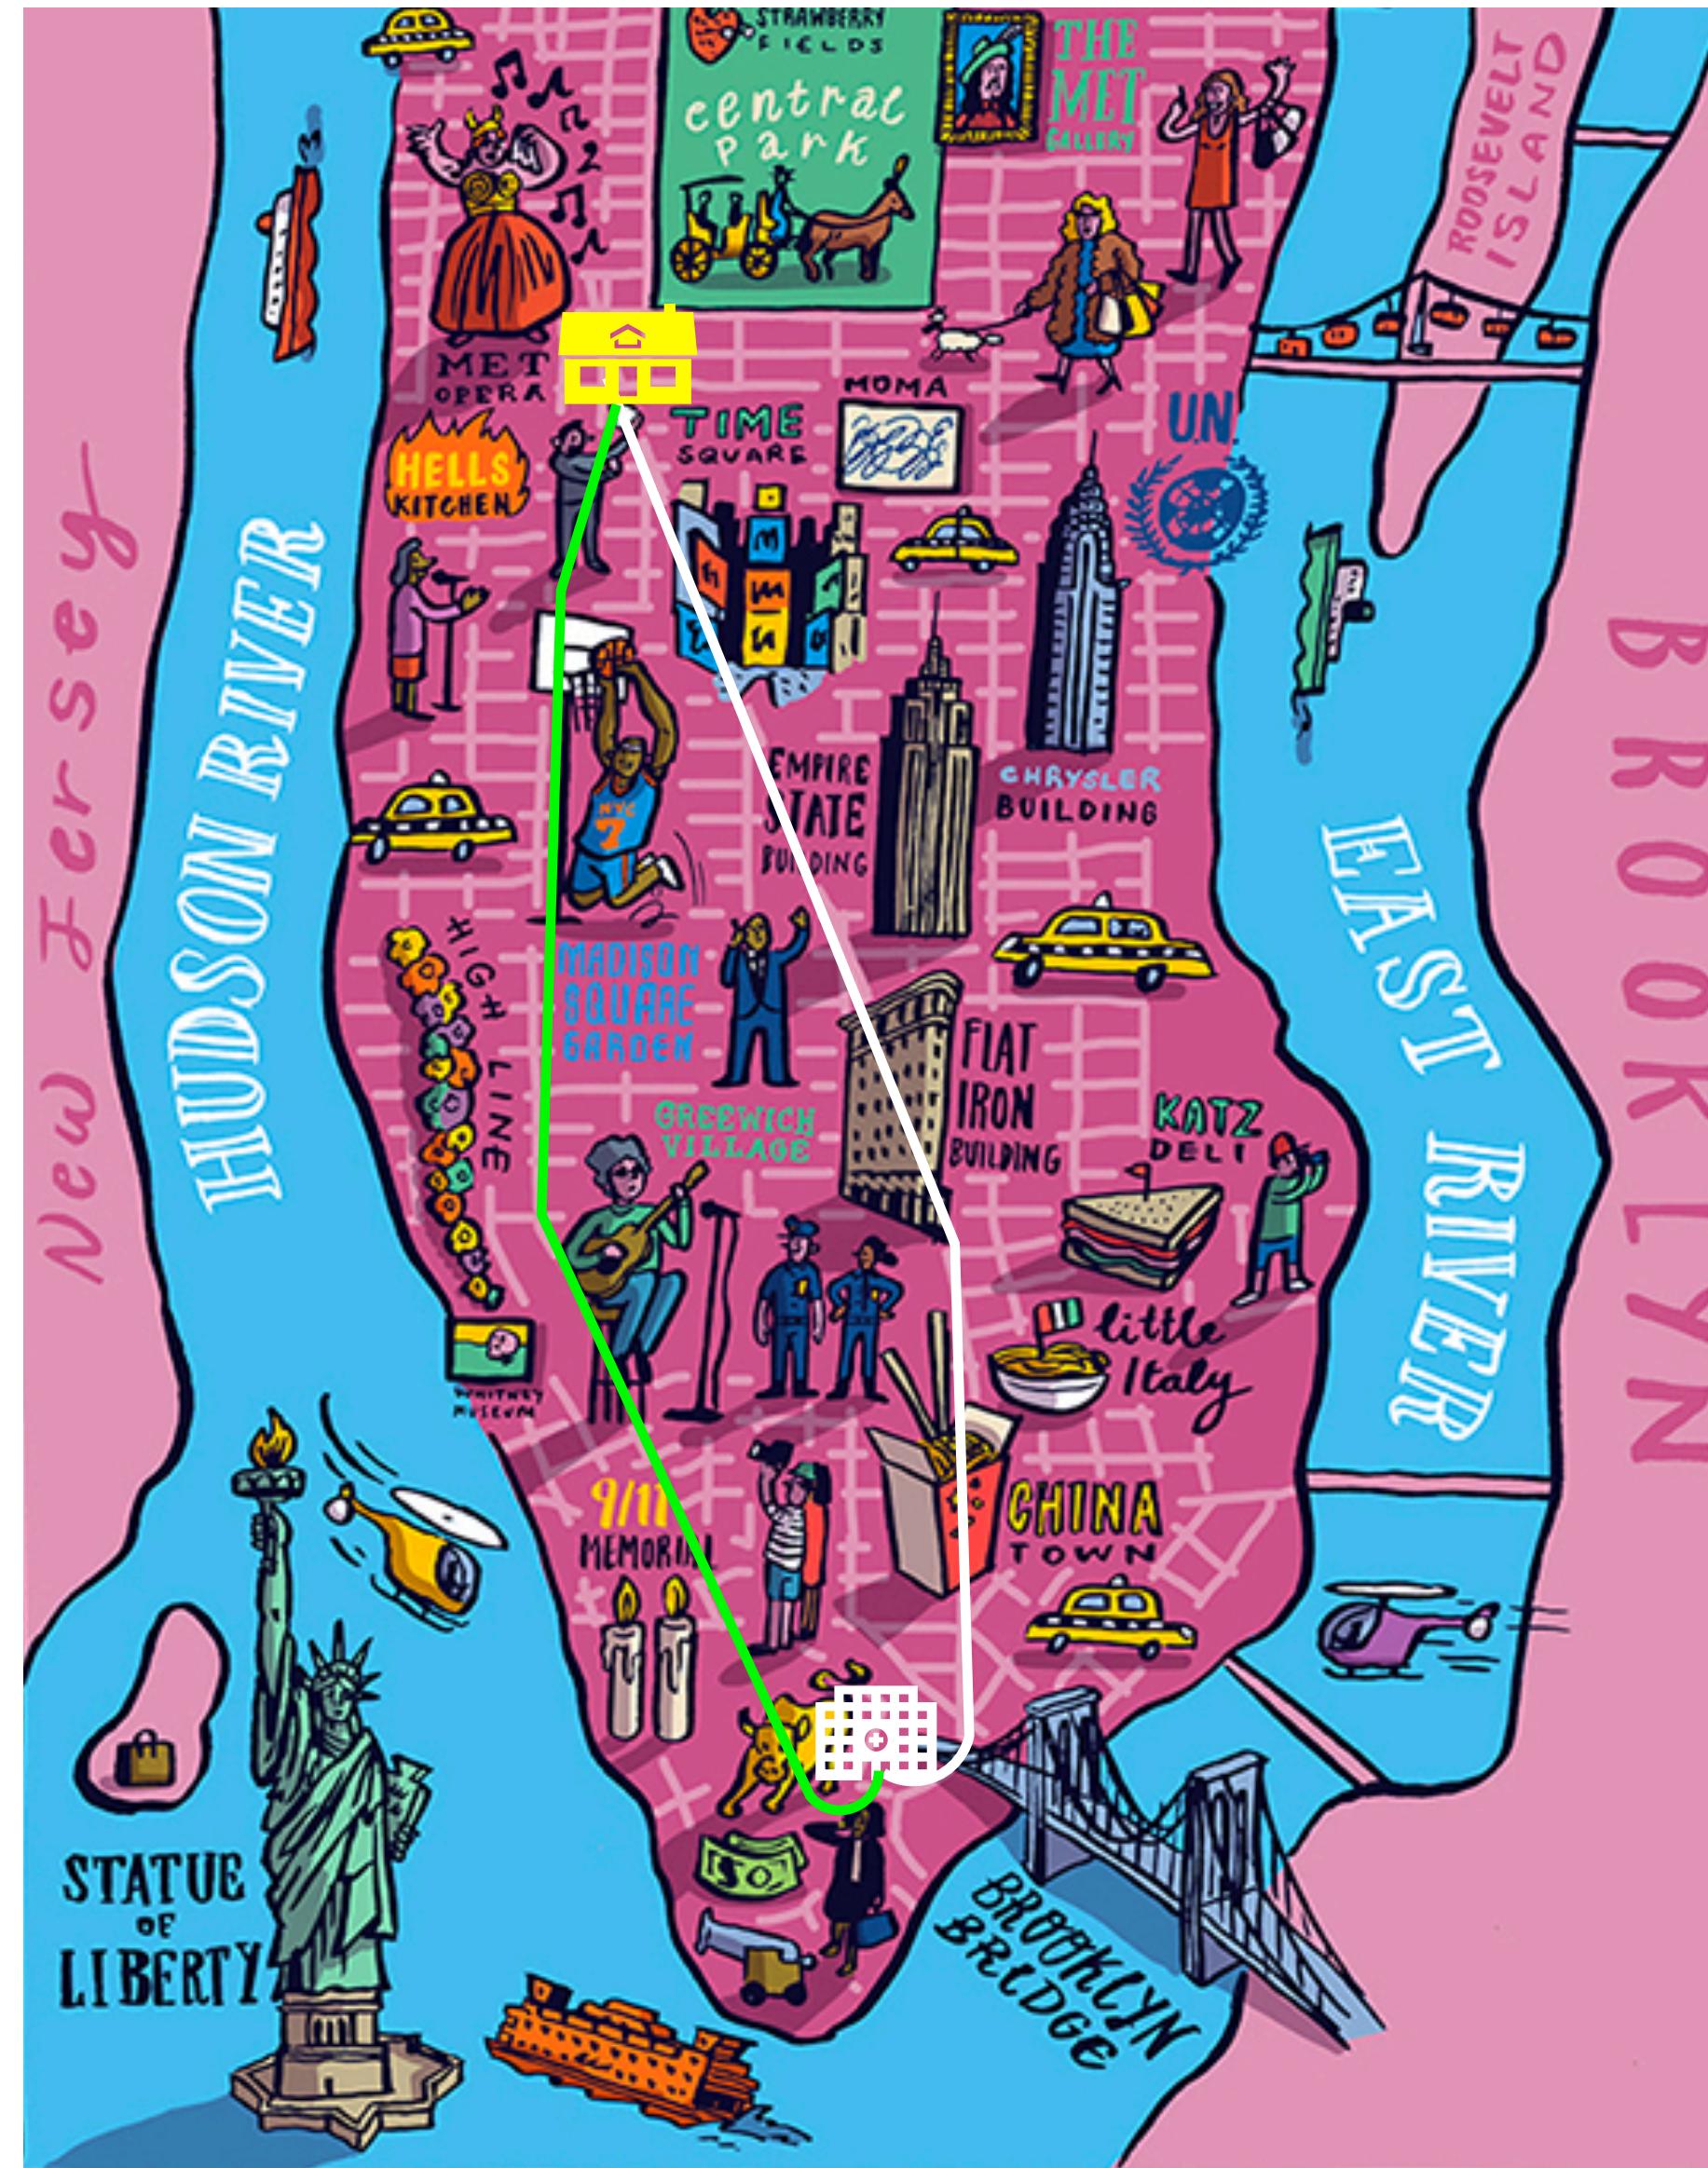
# Inference in Machine Learning



# Inference in Machine Learning



# Statistical Inference with Quantified Uncertainty



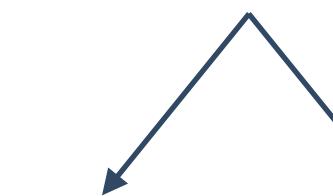
Your house Columbus Circle

NY-Presbyterian Hospital/Lower Manhattan

Route 1 → 25min ± 9min

Route 2 → 29min ± 2min

Uncertainties are important to make informed decisions



Frequentist approach  
(Bootstrapping)

Bayesian Inference  
(Includes prior knowledge)

# Bayesian Inference

Probability as a measure of *believability in an event*

$$P(A|B) = \frac{P(B|A) P(A)}{P(B)}$$

THE PROBABILITY OF "B" BEING TRUE GIVEN THAT "A" IS TRUE  
↓  
THE PROBABILITY OF "A" BEING TRUE

↑ THE PROBABILITY OF "A" BEING TRUE GIVEN THAT "B" IS TRUE  
THE PROBABILITY OF "B" BEING TRUE

A.I. Wiki

$$p(\theta|y) = \frac{p(y|\theta)p(\theta)}{p(y)}$$

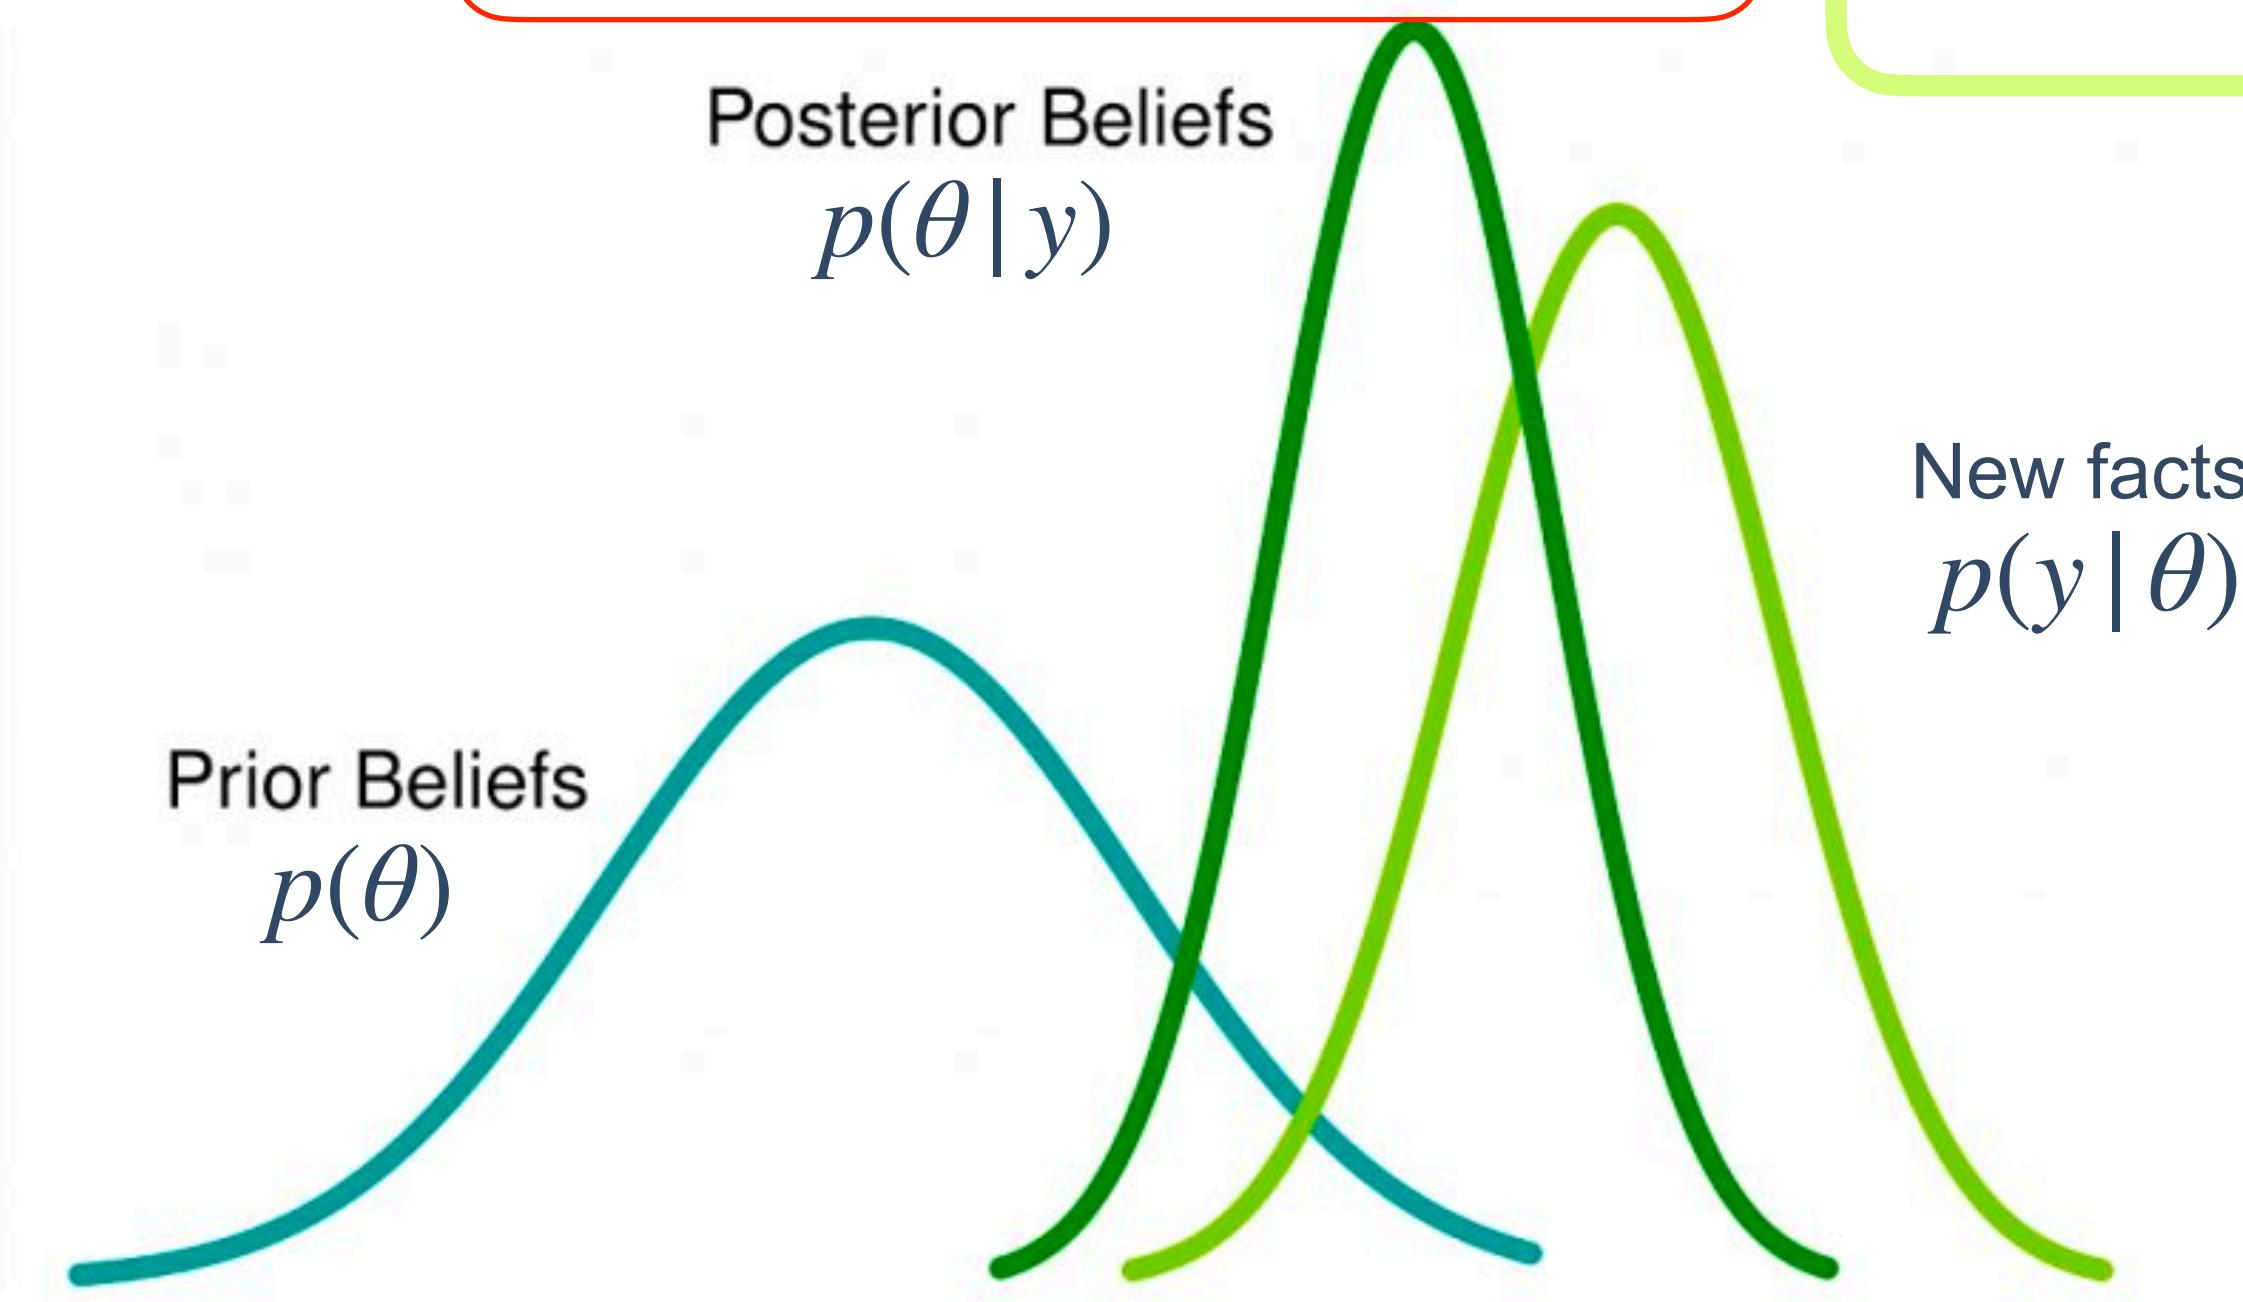
Likelihood → Prior  
Model → Data → Evidence

# Bayesian Inference

$$p(\theta | y) = \frac{p(y | \theta)p(\theta)}{p(y)}$$

Intractable

$$p(y) = \int_{\theta} p(y | \theta)p(\theta)d\theta$$



“When the facts change, I change my mind. What do you do, sir? “

John Maynard Keynes

# Bayesian Inference: Markov Chain Monte Carlo

$$p(\theta | y) \propto p(y | \theta)p(\theta)$$

$$p(\theta^0 | y) \propto p(y | \theta^0)p(\theta^0)$$

$$\begin{aligned} \theta &= (\mu, \sigma) \\ \mu^j &\sim N(\mu^{(j-1)}, \sigma) \end{aligned}$$

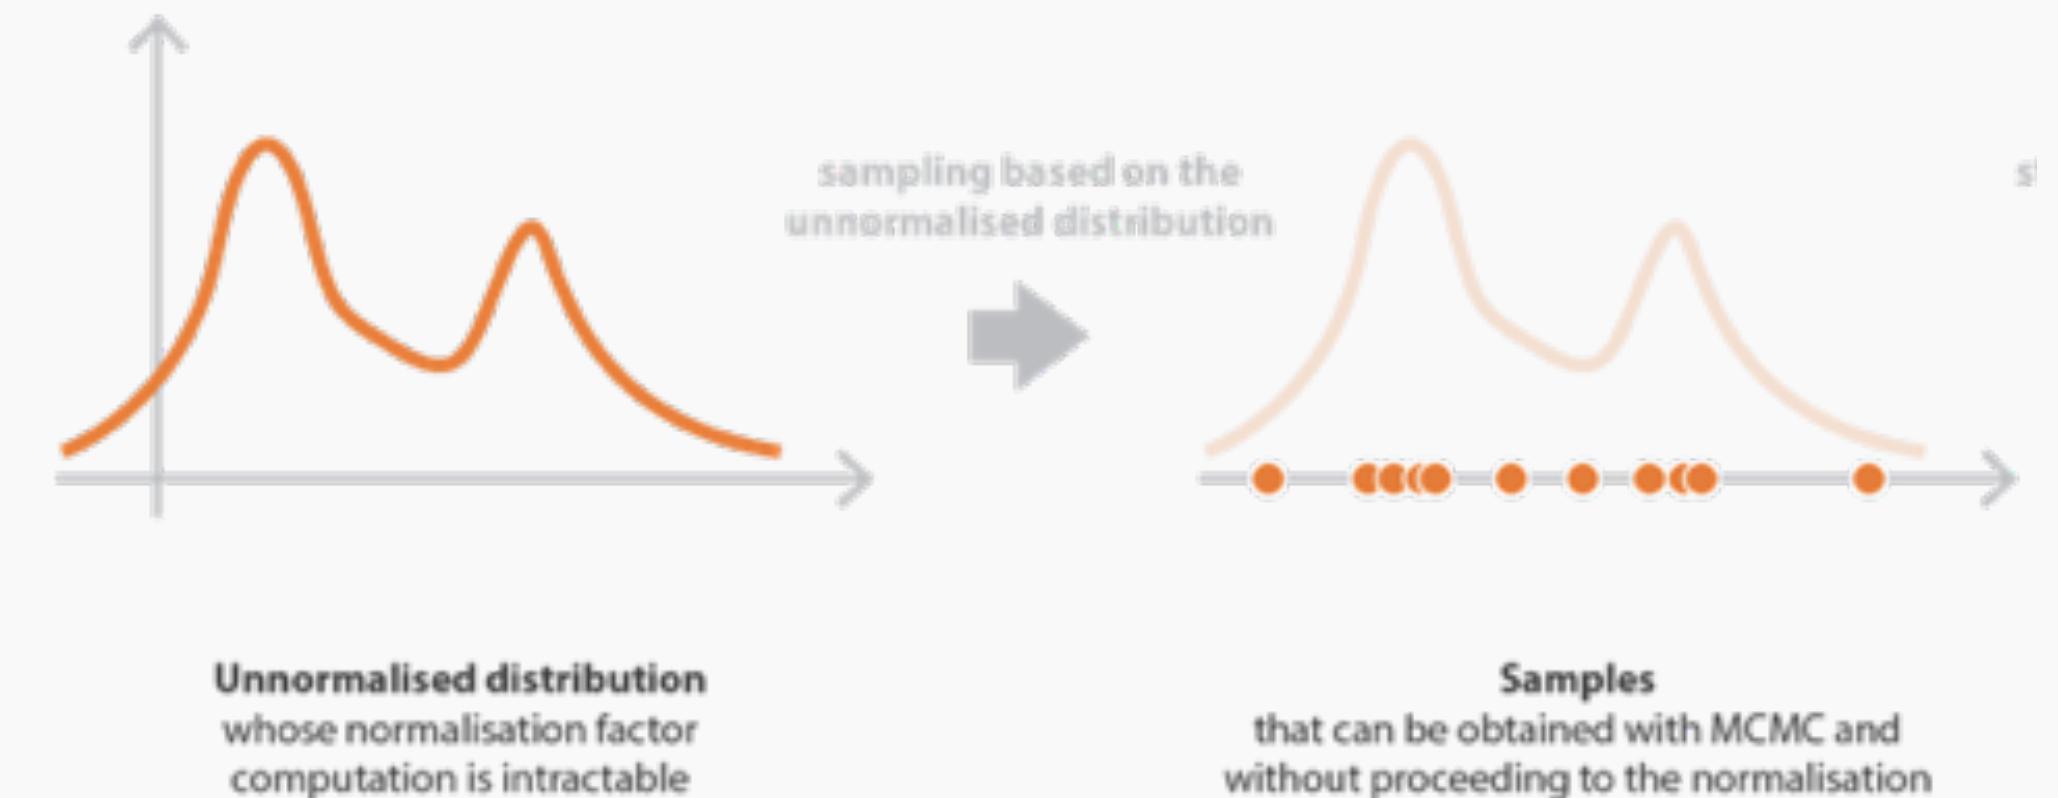
Step of random walk

$$\frac{p(\theta^j | y)}{p(\theta^{j-1} | y)} = \frac{p(y | \theta^j)p(\theta^j)}{p(y | \theta^{j-1})p(\theta^{j-1})}$$

accept

flip a coin

Histogram of my samples will be identical to the true posterior distribution... eventually

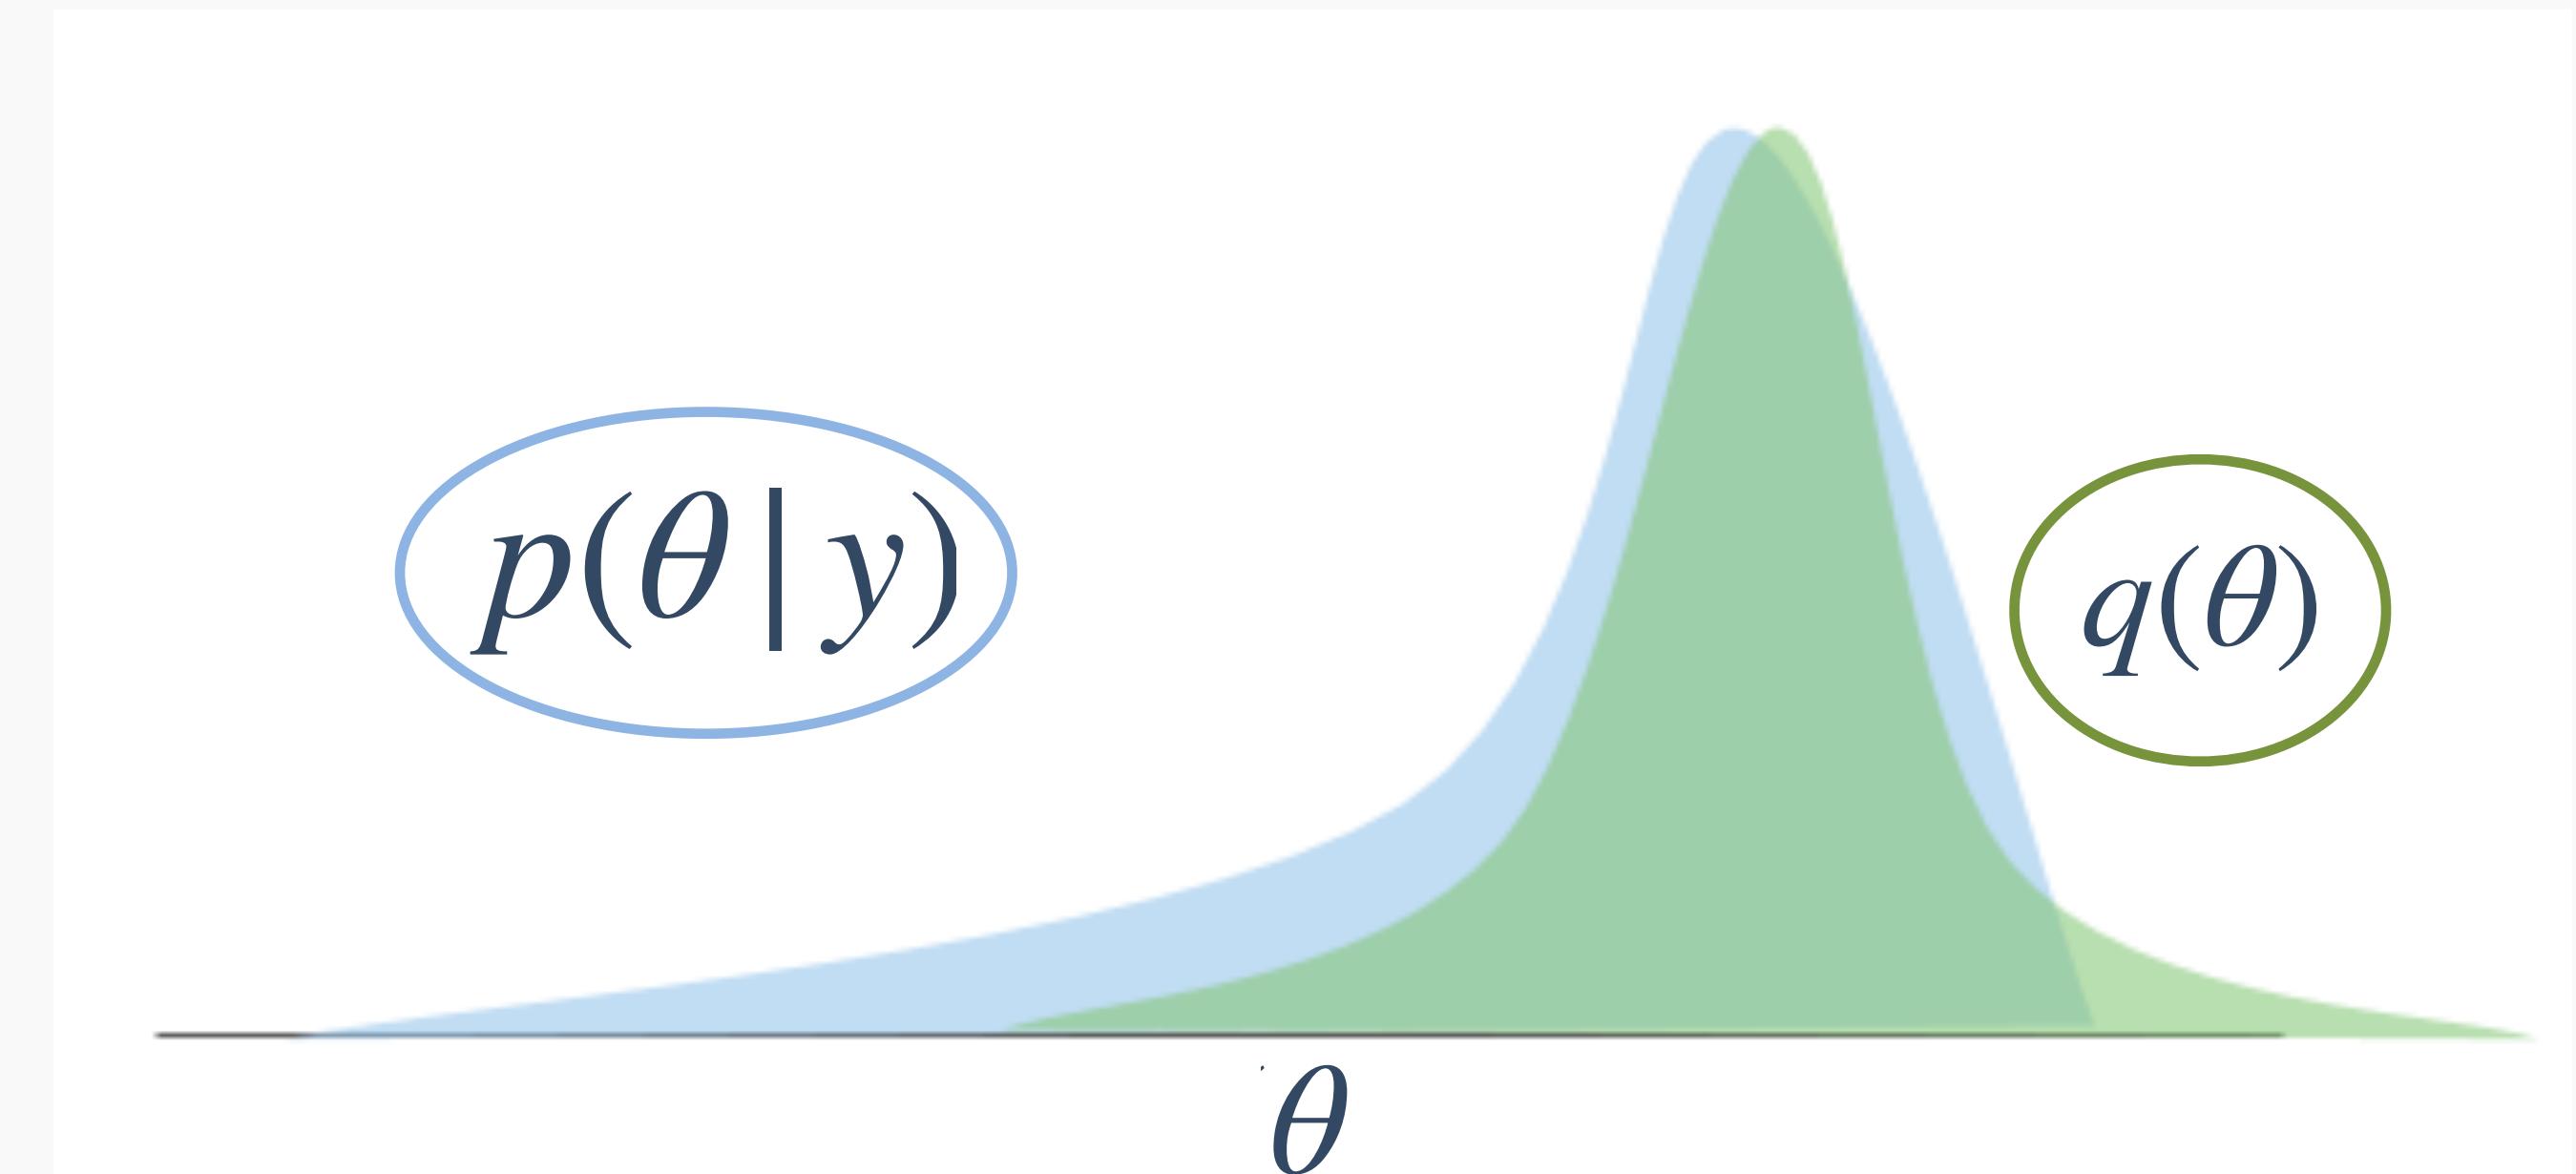


MCMC is eventually accurate, but not scalable to large models

# Approximate Bayesian Inference: Variational Inference

Optimization approach -> Q a family of “nice” distributions

$$p(\theta | y) = \frac{p(y | \theta) p(\theta)}{\int p(y, \theta) d\theta}$$



# Approximate Bayesian Inference: Variational Inference

$$p(\theta | y) = \frac{p(y | \theta) p(\theta)}{\int p(y, \theta) d\theta} = p(y) \leftarrow \text{evidence}$$

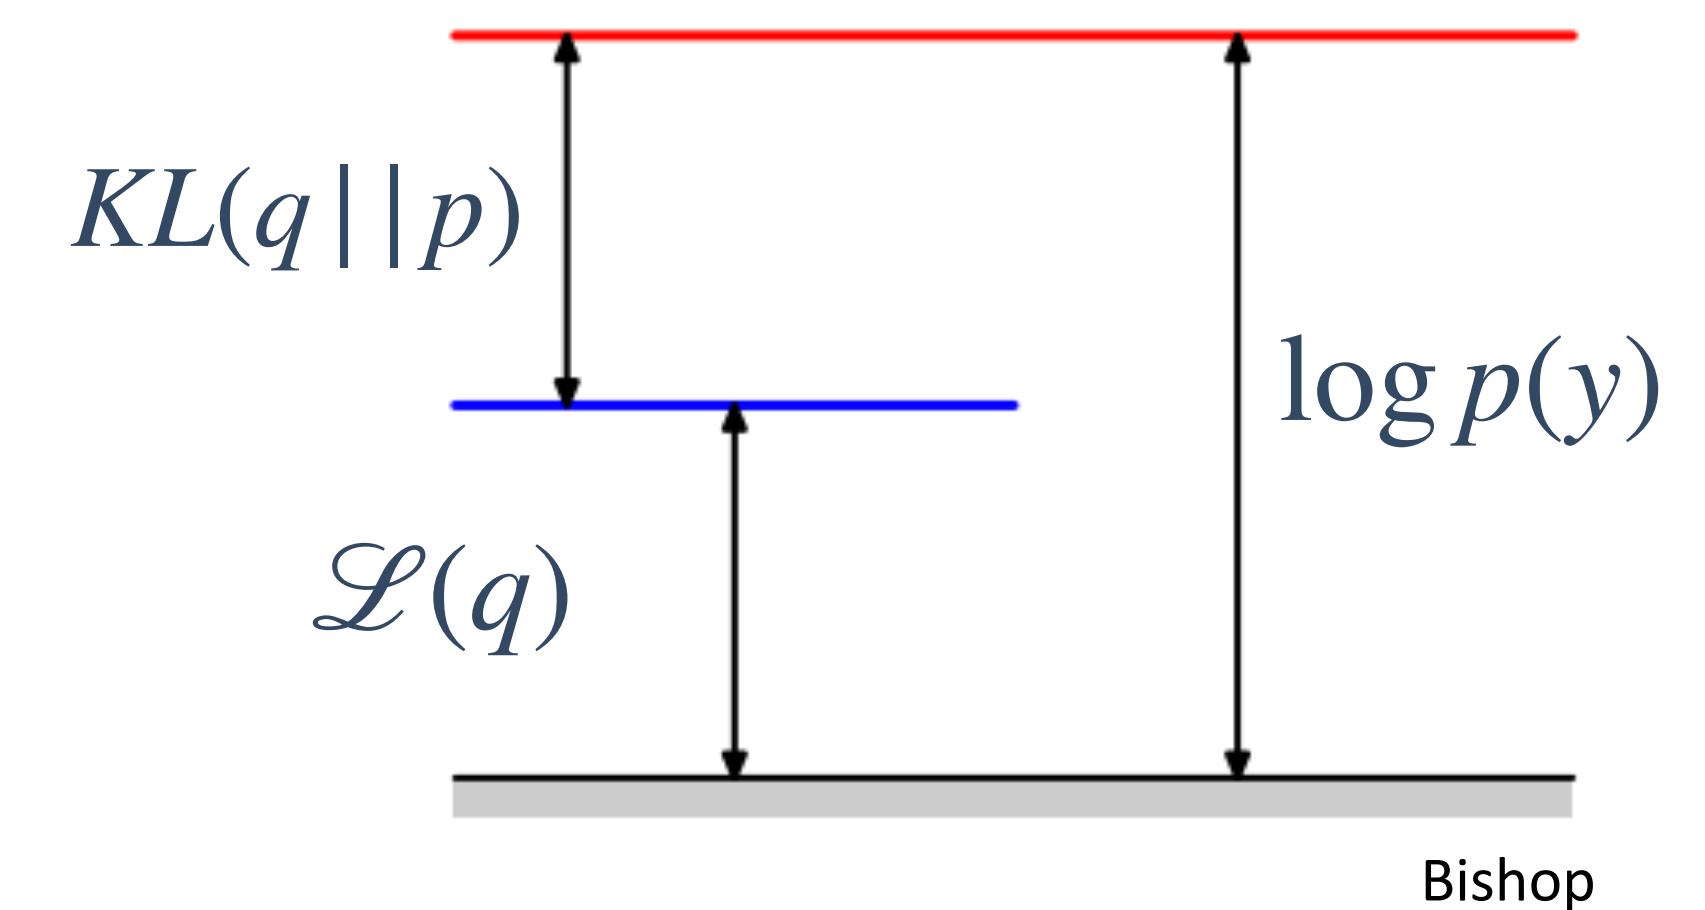
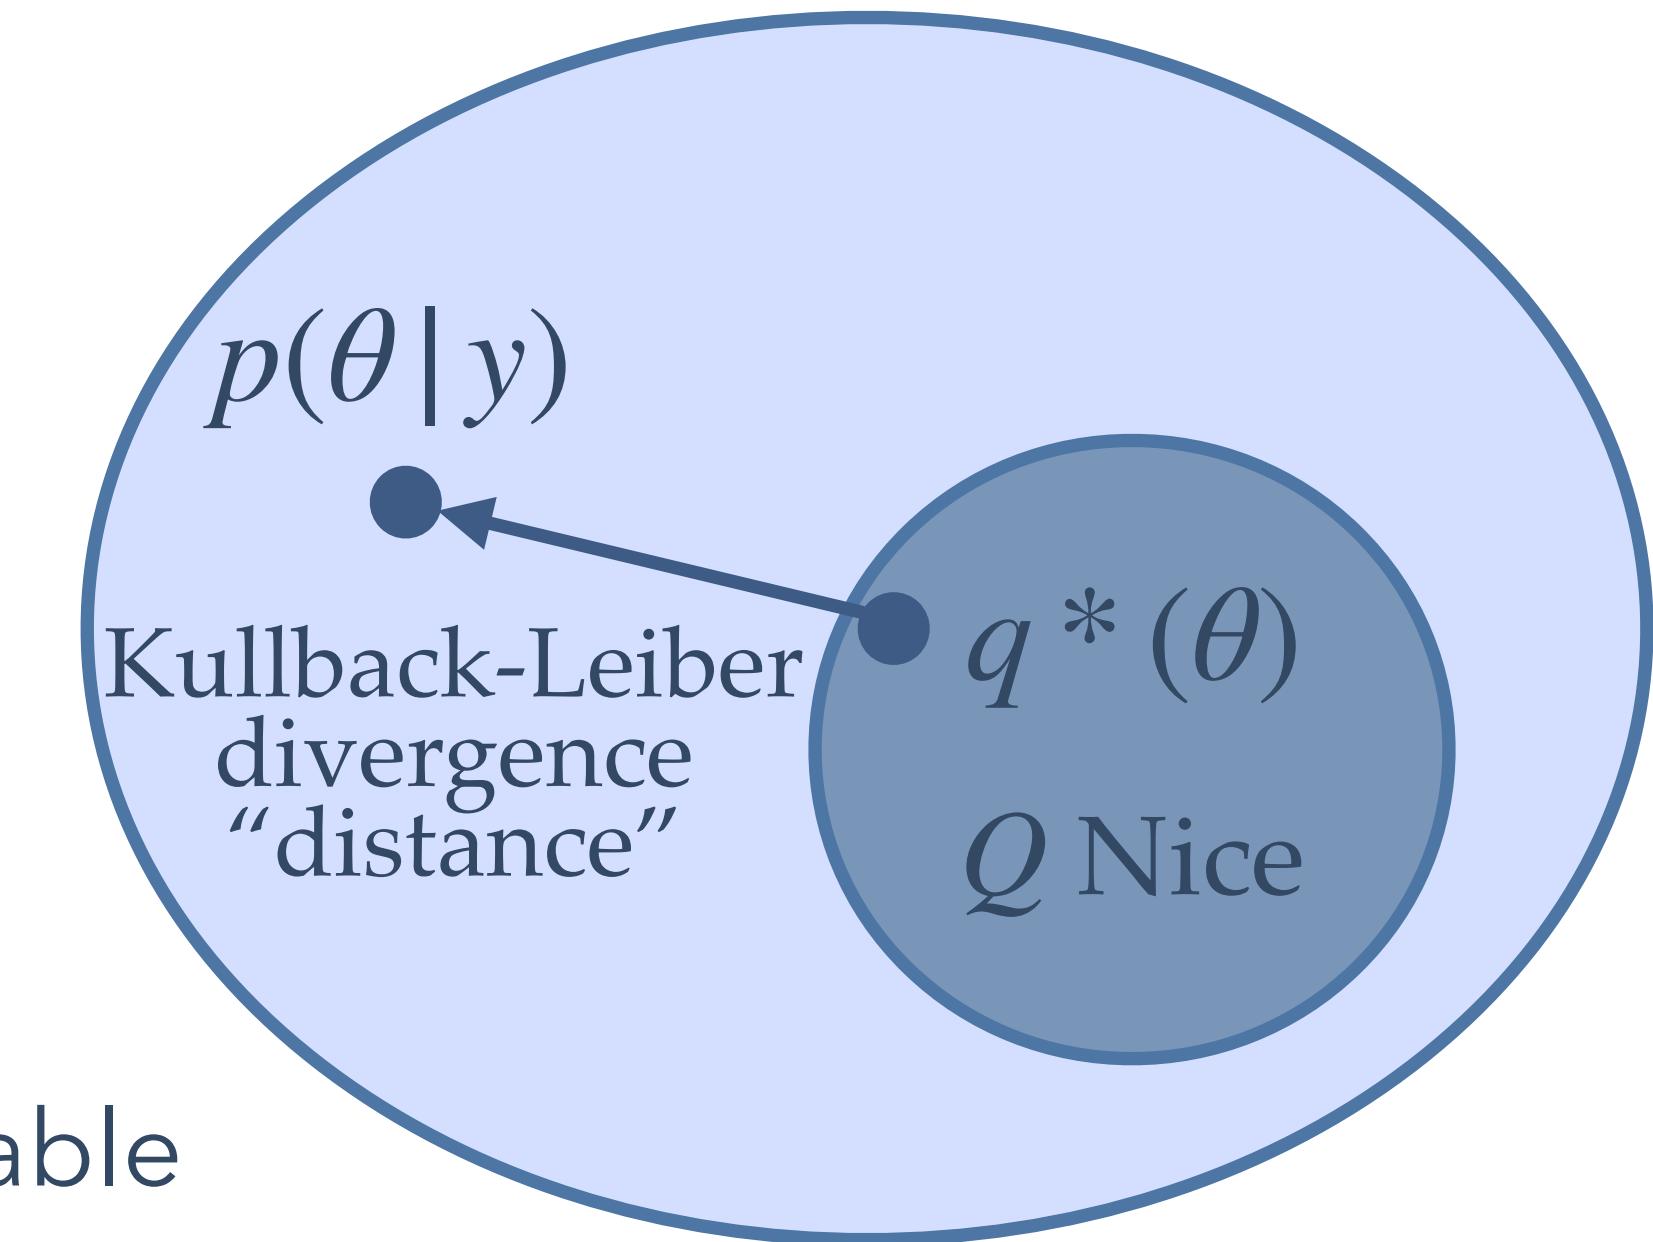
$$p(y) = \iiint_{\dots} p(y | \theta_1, \theta_2, \theta_3, \dots, \theta_m) d\theta_1 d\theta_2 d\theta_3 \dots d\theta_m$$

approximation

$$p(\theta | y) \approx q^*(\theta) \leftarrow \text{tractable family of distributions}$$

$$q^*(\theta) = \operatorname{argmin}_{q \in Q} KL(q(\cdot)) || p(\cdot | y)$$

intractable



# Approximate Bayesian Inference: Variational Inference

$$q^*(\theta) = \operatorname{argmin}_{q \in Q} KL(q(\cdot)) || p(\cdot | y)$$

$$\theta^* = \operatorname{argmin}_\theta KL[q(\omega | \theta) || p(\omega | y)]$$

$$KL(q(\cdot)) || p(\cdot | y) := \int q(\theta) \log \frac{q(\theta)}{p(\theta | y)} d\theta$$

$$p(\omega | y) = \frac{p(y | \omega)p(\omega)}{p(y)}$$

$$\theta^* = \operatorname{argmin}_\theta \int (q(\omega | \theta)) \log \frac{q(\omega | \theta)p(y)}{p(\omega)p(y | \omega)} d\omega$$

$$= \operatorname{argmin}_\theta \int q(\omega | \theta) \log \frac{q(\omega | \theta)}{p(\omega)p(y | \omega)} d\omega + \operatorname{argmin}_\theta [\log p(y) \int q(\omega | \theta) d\omega]$$

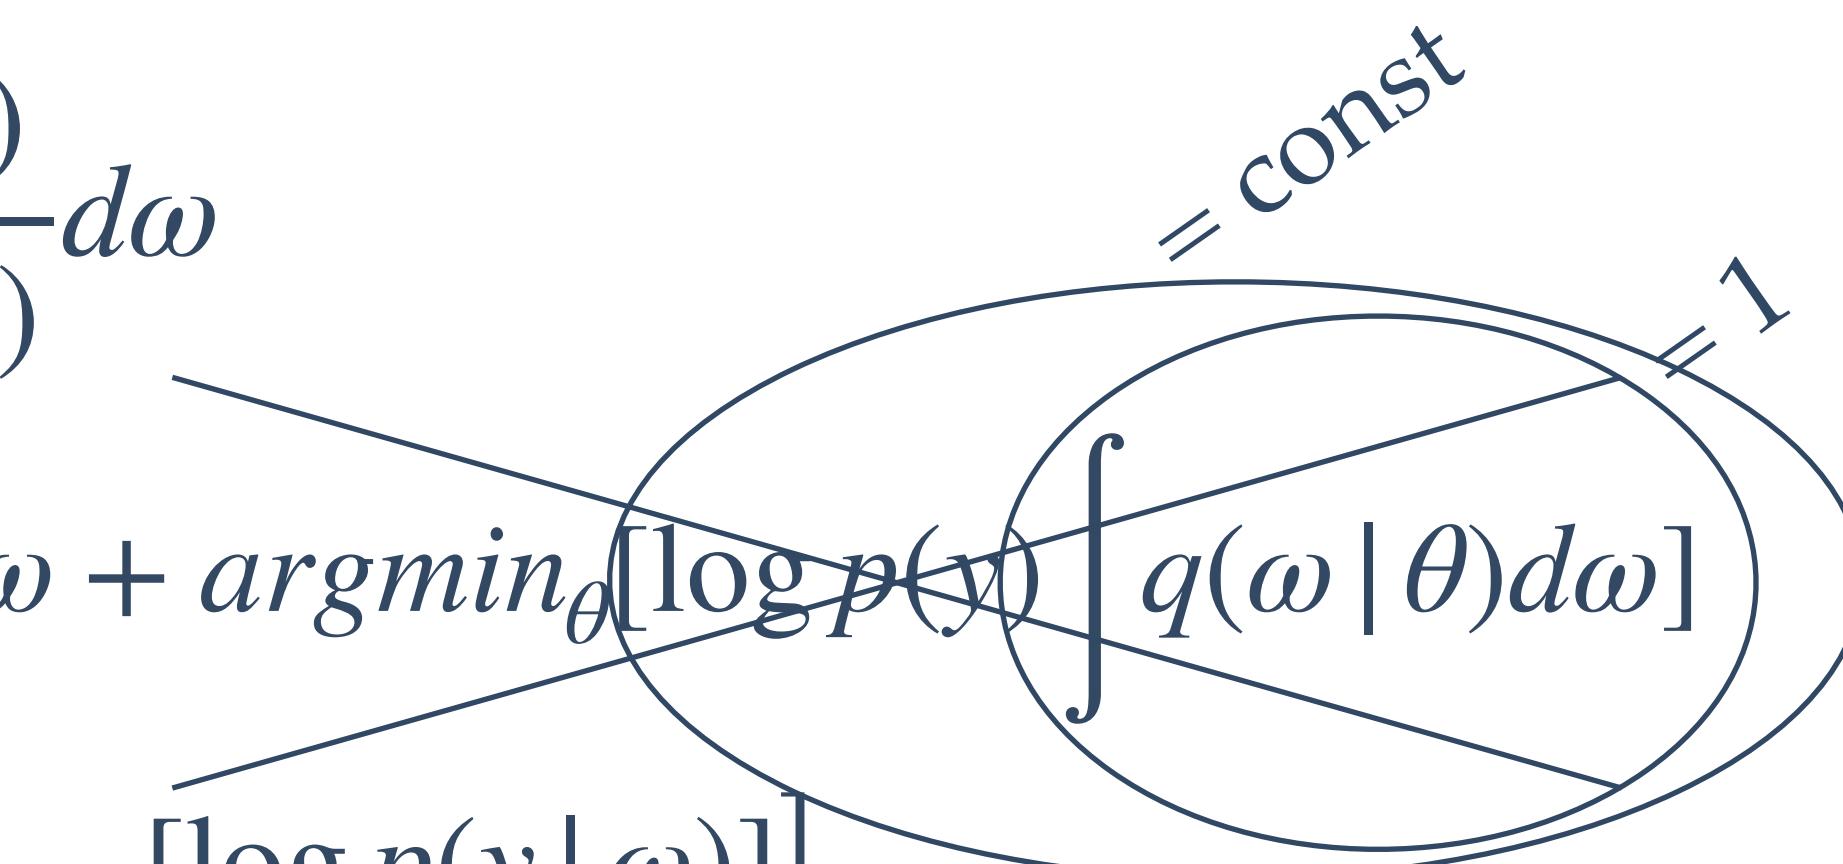
$$\theta^* = \operatorname{argmin}_\theta [KL[q(\omega | \theta) || p(\omega)] - \mathbb{E}_{q(\omega | \theta)}[\log p(y | \omega)]]$$

$$\mathcal{L}(y | \theta) = KL[q(\omega | \theta) || p(\omega)] - \mathbb{E}_{q(\omega | \theta)}[\log p(y | \omega)] \quad \xleftarrow{\hspace{1cm}} \text{Loss function}$$

Prior                      Likelihood

we need to optimize for  $q$

we can, instead, find the parameters  $\theta$  of a distribution on the weights  $q(\omega | \theta)$



# Approximate Bayesian Inference: Bayes by Backprop

$$\mathcal{L}(y | \theta) = KL[q(\omega | \theta) || p(\omega)] - \mathbb{E}_{q(\omega | \theta)}[\log p(y | \omega)] \quad \leftarrow \begin{array}{l} \text{Prior} \\ \text{Likelihood} \end{array} \quad \begin{array}{l} \text{Loss function} \\ \downarrow \text{Back propagate} \end{array}$$
$$\nabla_{\theta} \mathcal{L}(y | \theta)$$

# Approximate Bayesian Inference: Bayes by Backprop

$$\mathcal{L}(y|\theta) = KL[q(\omega|\theta) || p(\omega)] - \mathbb{E}_{q(\omega|\theta)}[\log p(y|\omega)] \quad \leftarrow \begin{array}{l} \text{Prior} \\ \text{Likelihood} \end{array} \quad \begin{array}{l} \text{Loss function} \\ \downarrow \text{Back propagate} \end{array}$$
$$\nabla_{\theta} \mathcal{L}(y|\theta) \rightarrow \nabla_{\theta} \mathcal{L}(\omega, \theta)$$

Let's define the *pointwise* (sampled) loss for a given  $\omega$  and  $\theta$ :

$$\mathcal{L}(\omega, \theta) = \log q(\omega|\theta) - \log p(\omega) - \log p(y|\omega)$$

Notice that if you take the expectation under  $q(\omega|\theta)$ :

$$\mathbb{E}_{q(\omega|\theta)}[\mathcal{L}(\omega, \theta)] = \text{KL}[q(\omega|\theta) || p(\omega)] - \mathbb{E}_{q(\omega|\theta)}[\log p(y|\omega)] = \mathcal{L}(y|\theta)$$

# Approximate Bayesian Inference: Bayes by Backprop

$$\mathcal{L}(y|\theta) = KL[q(\omega|\theta)||p(\omega)] - \mathbb{E}_{q(\omega|\theta)}[\log p(y|\omega)] \quad \leftarrow \begin{array}{l} \text{Prior} \\ \text{Likelihood} \end{array} \quad \begin{array}{l} \text{Loss function} \\ \downarrow \text{Back propagate} \end{array}$$
$$q(\theta) = N(\widehat{\mu, \sigma}) \quad \omega \text{ are samples of this distribution}$$
$$\omega = \mu + \sigma \odot \epsilon, \quad q(\epsilon), \quad q(\epsilon)d\epsilon = q(\omega|\theta)d\omega$$

Weight Uncertainty in Neural Networks,  
Blundell et al. 1015( <https://arxiv.org/pdf/1505.05424.pdf>)

Suppose you can write  $\omega = t(\epsilon, \theta)$ , with  $\epsilon \sim q(\epsilon)$  (independent of  $\theta$ ), so:

$$q(\omega|\theta)d\omega = q(\epsilon)d\epsilon$$

and thus,

$$\int f(\omega, \theta)q(\omega|\theta)d\omega = \int f(t(\epsilon, \theta), \theta)q(\epsilon)d\epsilon = \mathbb{E}_{q(\epsilon)}[f(t(\epsilon, \theta), \theta)]$$

# Approximate Bayesian Inference: Bayes by Backprop

$$\mathcal{L}(y|\theta) = KL[q(\omega|\theta)||p(\omega)] - \mathbb{E}_{q(\omega|\theta)}[\log p(y|\omega)] \quad \leftarrow \begin{array}{l} \text{Prior} \\ \text{Likelihood} \end{array} \quad \begin{array}{l} \text{Loss function} \\ \downarrow \text{Back propagate} \end{array}$$

$$q(\theta) = N(\widehat{\mu, \sigma})$$

$\omega$  are samples of this distribution

$$\omega = \mu + \sigma \odot \epsilon, \quad q(\epsilon), \quad q(\epsilon)d\epsilon = q(\omega|\theta)d\omega$$

$$\frac{\partial}{\partial \theta} \mathbb{E}_{q(\omega|\theta)}[f(\omega, \theta)] = \mathbb{E}_{q(\epsilon)} \left[ \frac{\partial f(\omega, \theta)}{\partial \omega} \frac{\partial \omega}{\partial \theta} + \frac{\partial f(\omega, \theta)}{\partial \theta} \right]$$

$$\frac{\partial}{\partial \theta} \mathcal{L}(y|\theta) \longrightarrow \nabla_\mu \mathcal{L}(\omega, \theta), \quad \nabla_\sigma \mathcal{L}(\omega, \theta)$$

Weight Uncertainty in Neural Networks,  
Blundell et al. 1015( <https://arxiv.org/pdf/1505.05424.pdf>)

$$\begin{cases} \frac{\partial}{\partial \mu} \mathbb{E}_{q(\omega|\theta)}[\mathcal{L}(\omega, \theta(\mu, \sigma))] = \mathbb{E}_{q(\epsilon)} \left[ \frac{\partial \mathcal{L}(\omega, \theta)}{\partial \theta} \frac{\partial \omega}{\partial \mu} + \frac{\partial \mathcal{L}(\omega, \theta)}{\partial \mu} \right] \\ \frac{\partial}{\partial \sigma} \mathbb{E}_{q(\omega|\theta)}[\mathcal{L}(\omega, \theta(\mu, \sigma))] = \mathbb{E}_{q(\epsilon)} \left[ \frac{\partial \mathcal{L}(\omega, \theta)}{\partial \theta} \frac{\partial \omega}{\partial \sigma} + \frac{\partial \mathcal{L}(\omega, \theta)}{\partial \sigma} \right] \end{cases}$$

this is the same for  
both  $\mu$  and  $\sigma$

update de variational  
parameters

$$\begin{cases} \mu = \mu - \alpha \nabla_\mu \mathcal{L} \\ \sigma = \sigma - \alpha \nabla_\sigma \mathcal{L} \end{cases} \quad \begin{array}{l} \leftarrow \text{Your standard} \\ \text{gradient descent} \end{array}$$

# Approximate Bayesian Inference: Flipout

The problem is that to make this computationally possible, we sample a single  $\epsilon$  per mini-batch and, therefore, sharing the same weight perturbation.



$$\omega = \mu + \sigma \odot \epsilon$$

This introduces correlations between the gradients of each mini-batch



Flipout decorrelates them by simply adding a flip-the-coin effect:

$$q(\epsilon) \longrightarrow \epsilon \xrightarrow{\text{flip the coin}} + \text{ or } - \longrightarrow \pm \epsilon \longrightarrow \omega = \mu \pm \sigma \odot \epsilon$$

For more details see Flipout: Efficient Pseudo-Independent Weight Perturbations on Mini-Batches  
Wen et al. (<https://arxiv.org/pdf/1803.04386.pdf>)

# Approximate Bayesian Inference: Variational Inference

$$p(\theta|y) = \frac{p(y|\theta)p(\theta)}{\int p(y,\theta)d\theta}$$

$$q^*(\theta) = \operatorname{argmin}_{q \in Q} KL(q(\cdot)) || p(\cdot | y)$$

Kullback-Leibler divergence:

$$KL(q(\cdot)) || p(\cdot | y) := \int q(\theta) \log \frac{q(\theta)}{p(\theta | y)} d\theta$$

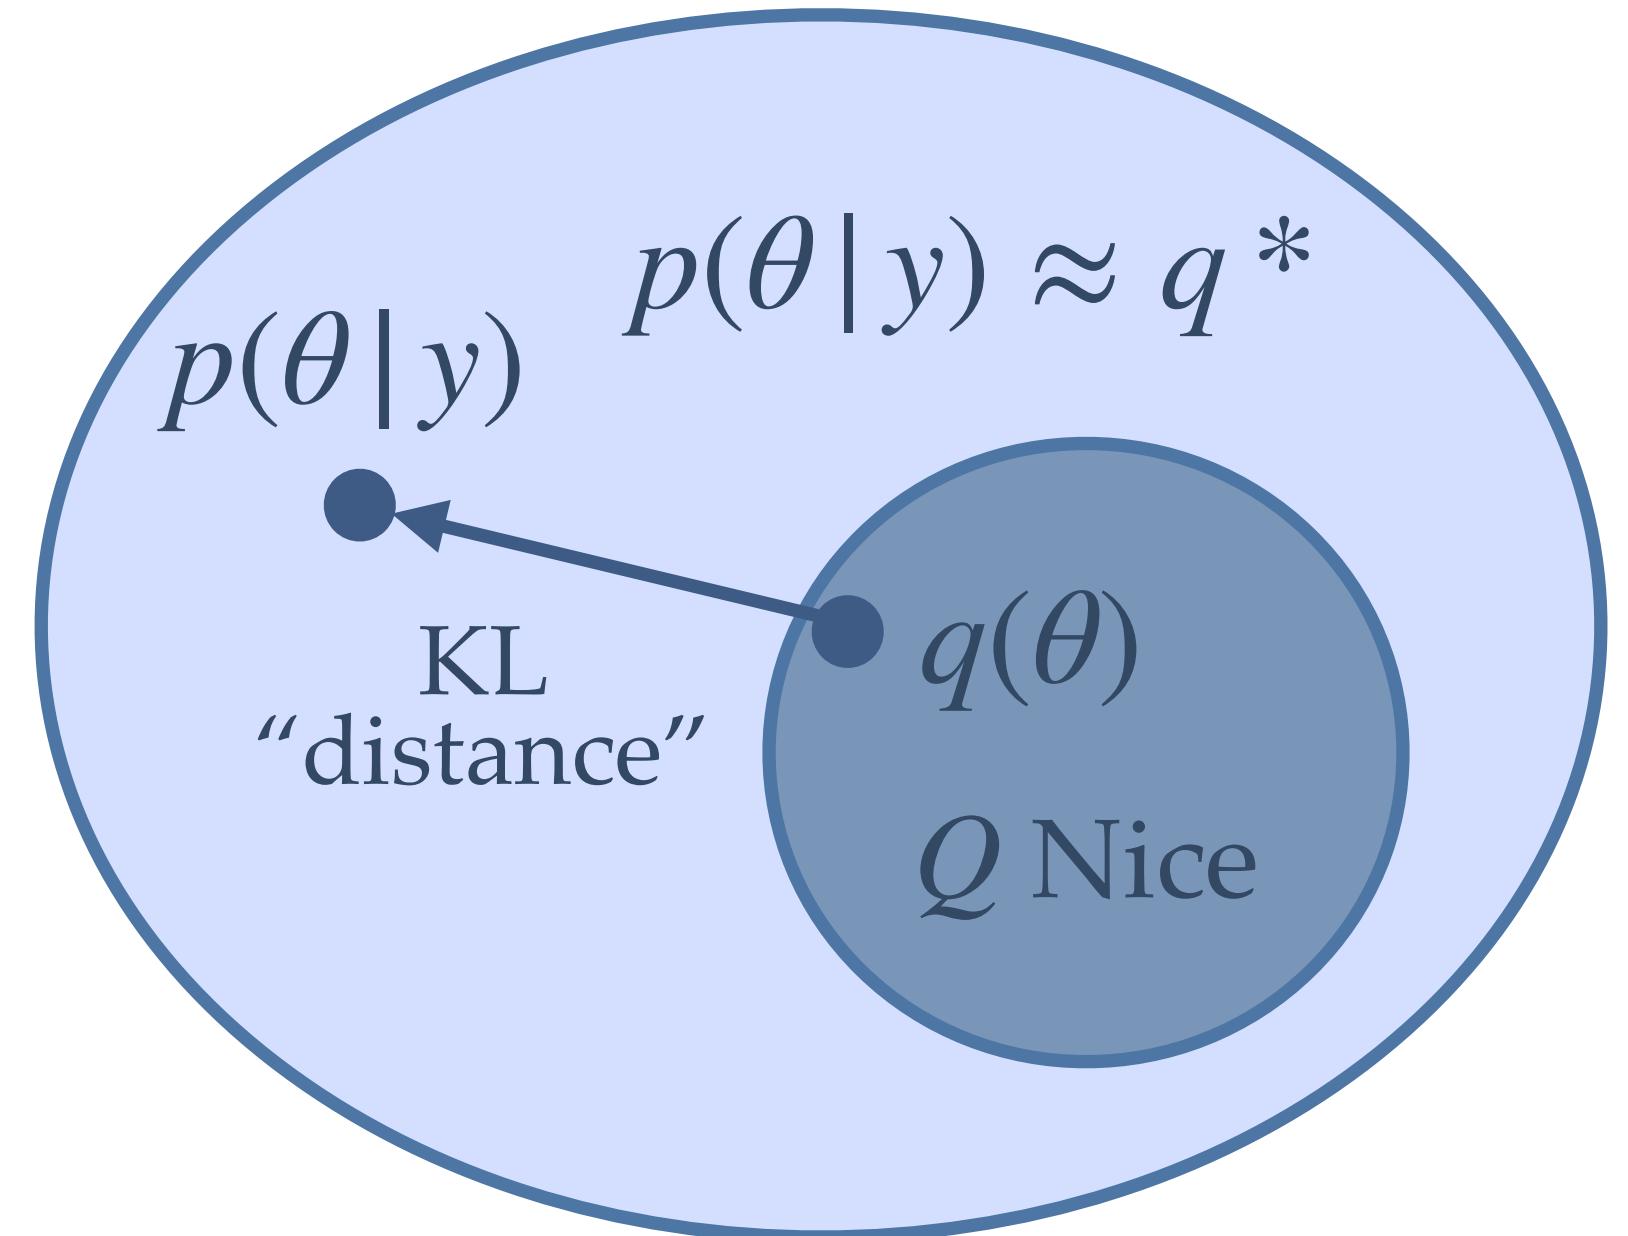
$$= \int q(\theta) \log \frac{q(\theta)p(y)}{p(\theta, y)} d\theta = \int q(\theta) \left[ \log p(y) + \log \frac{q(\theta)}{p(\theta, y)} \right] d\theta$$

$$= \log p(y) \int q(\theta) d\theta - \int q(\theta) \log \frac{p(\theta, y)}{q(\theta)} d\theta = \log p(y) - \int q(\theta) \log \frac{p(\theta, y)}{q(\theta)} d\theta$$

Evidence lower bound (ELBO)

one can prove (check Bishop)

$$\downarrow \\ KL \geq 0 \implies \log p(y) \geq \text{ELBO} \rightarrow q^* = \operatorname{argmax}_{q \in Q} \text{ELBO}(q) = \operatorname{argmax}_{q \in Q} \int q(\theta) \log \frac{p(\theta, y)}{q(\theta)} d\theta$$



$$\int q(\theta) \log \frac{p(\theta, y)}{q(\theta)} d\theta$$

# Approximate Bayesian Inference: Mean Field Variational Bayes

$$q^* = \operatorname{argmax}_{q \in Q} \int q(\theta) \log \frac{p(\theta, y)}{q(\theta)} d\theta \quad \text{ELBO}$$

$$q(\theta) = \prod_i^m q_i(\theta_i)$$

We assume  $q$  factorizes with respect to  $\theta$ s

$$\text{ELBO} = \int q(\theta) \log \frac{p(\theta, y)}{q(\theta)} d\theta = \int q(\theta) \log p(\theta, y) d\theta - \int q(\theta) \log q(\theta) d\theta$$

$$q(\theta) \approx \prod_i q_i(\theta_i)$$

$$\equiv \int \prod_i q_i(\theta_i) \log \frac{p(\theta, y)}{\prod_i q_i(\theta_i)} d\theta = \int \prod_i q_i(\theta_i) [\log p(\theta, y) - \sum_i \log q_i(\theta_i)] d\theta$$

choose one  $q_j$

$$= \int q_j(\theta_j) \left[ \int \log p(\theta, y) \prod_{i \neq j} q_i(\theta_i) d\theta_i \right] d\theta_j - \int q_j(\theta_j) \log q_j(\theta_j) \left[ \int \prod_{i \neq j} q_i(\theta_i) d\theta_i \right] d\theta_j + \text{const.}$$

$$= \int q_j(\theta_j) \log \tilde{p}(\theta_j, y) d\theta_j - \int q_j(\theta_j) \log q_j(\theta_j) d\theta_j + \text{const.} = -KL(q_j(\theta_j) || \tilde{p}(\theta_j, y)) + \text{const.}$$

= const.

# Approximate Bayesian Inference: Mean Field Variational Bayes

Optimize for  $q_j$ :

$$\text{ELBO} = -KL(q_j(\theta_j) || \tilde{p}(\theta_j, y)) + \text{const.} = -\int q_j(\theta_j) \log \frac{q_j(\theta_j)}{\tilde{p}(\theta_j, y)} d\theta_j + \text{const.}$$

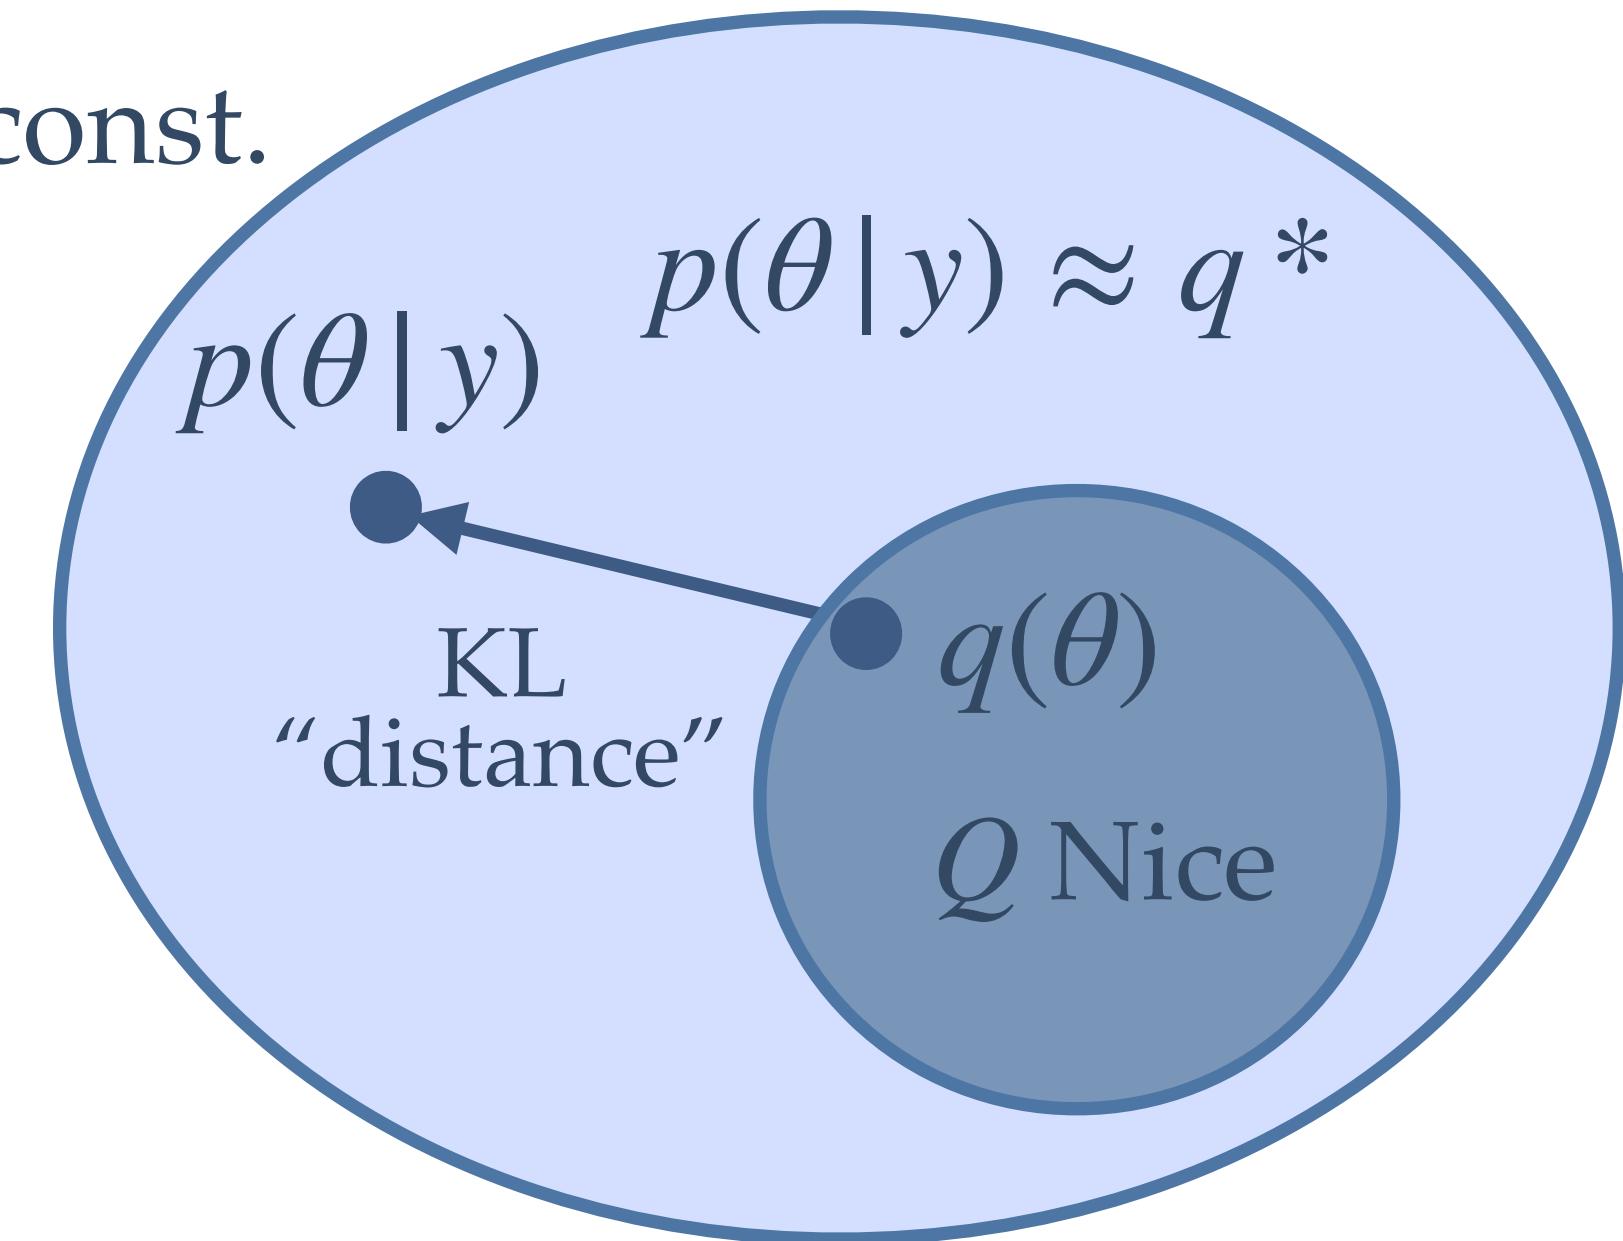
Maximizing the ELBO is equivalent to minimizing KL:

$$q^*_j = \operatorname{argmin}_{q(\theta_j)} KL(q_j(\theta_j) || \tilde{p}(\theta_j, y)) \quad \Rightarrow \quad q_j(\theta_j) = \tilde{p}(\theta_j, y)$$

$$\log q^*_j(\theta_j) = \log \tilde{p}(\theta_j, y) = \mathbb{E}_{i \neq j}[\log p(\theta, y)] + \text{const.}$$

$$q^*_j(\theta_j) = \frac{\exp(\mathbb{E}_{i \neq j}[\log p(\theta, y)])}{\int \exp(\mathbb{E}_{i \neq j}[\log p(\theta, y)]) d\theta_j}$$

set by normalizing  
 $q^*_j(\theta_j)$



This represents the conditions for the maximum of the ELBO, given the factorization assumption. The problem is that the expression for  $q^*_j(\theta_j)$  depends on expectations of all other  $q_i(\theta_i)$ . We need to iterate through them. Convergence is guaranteed because ELBO is convex with respect to each factor  $q_i(\theta_i)$ .

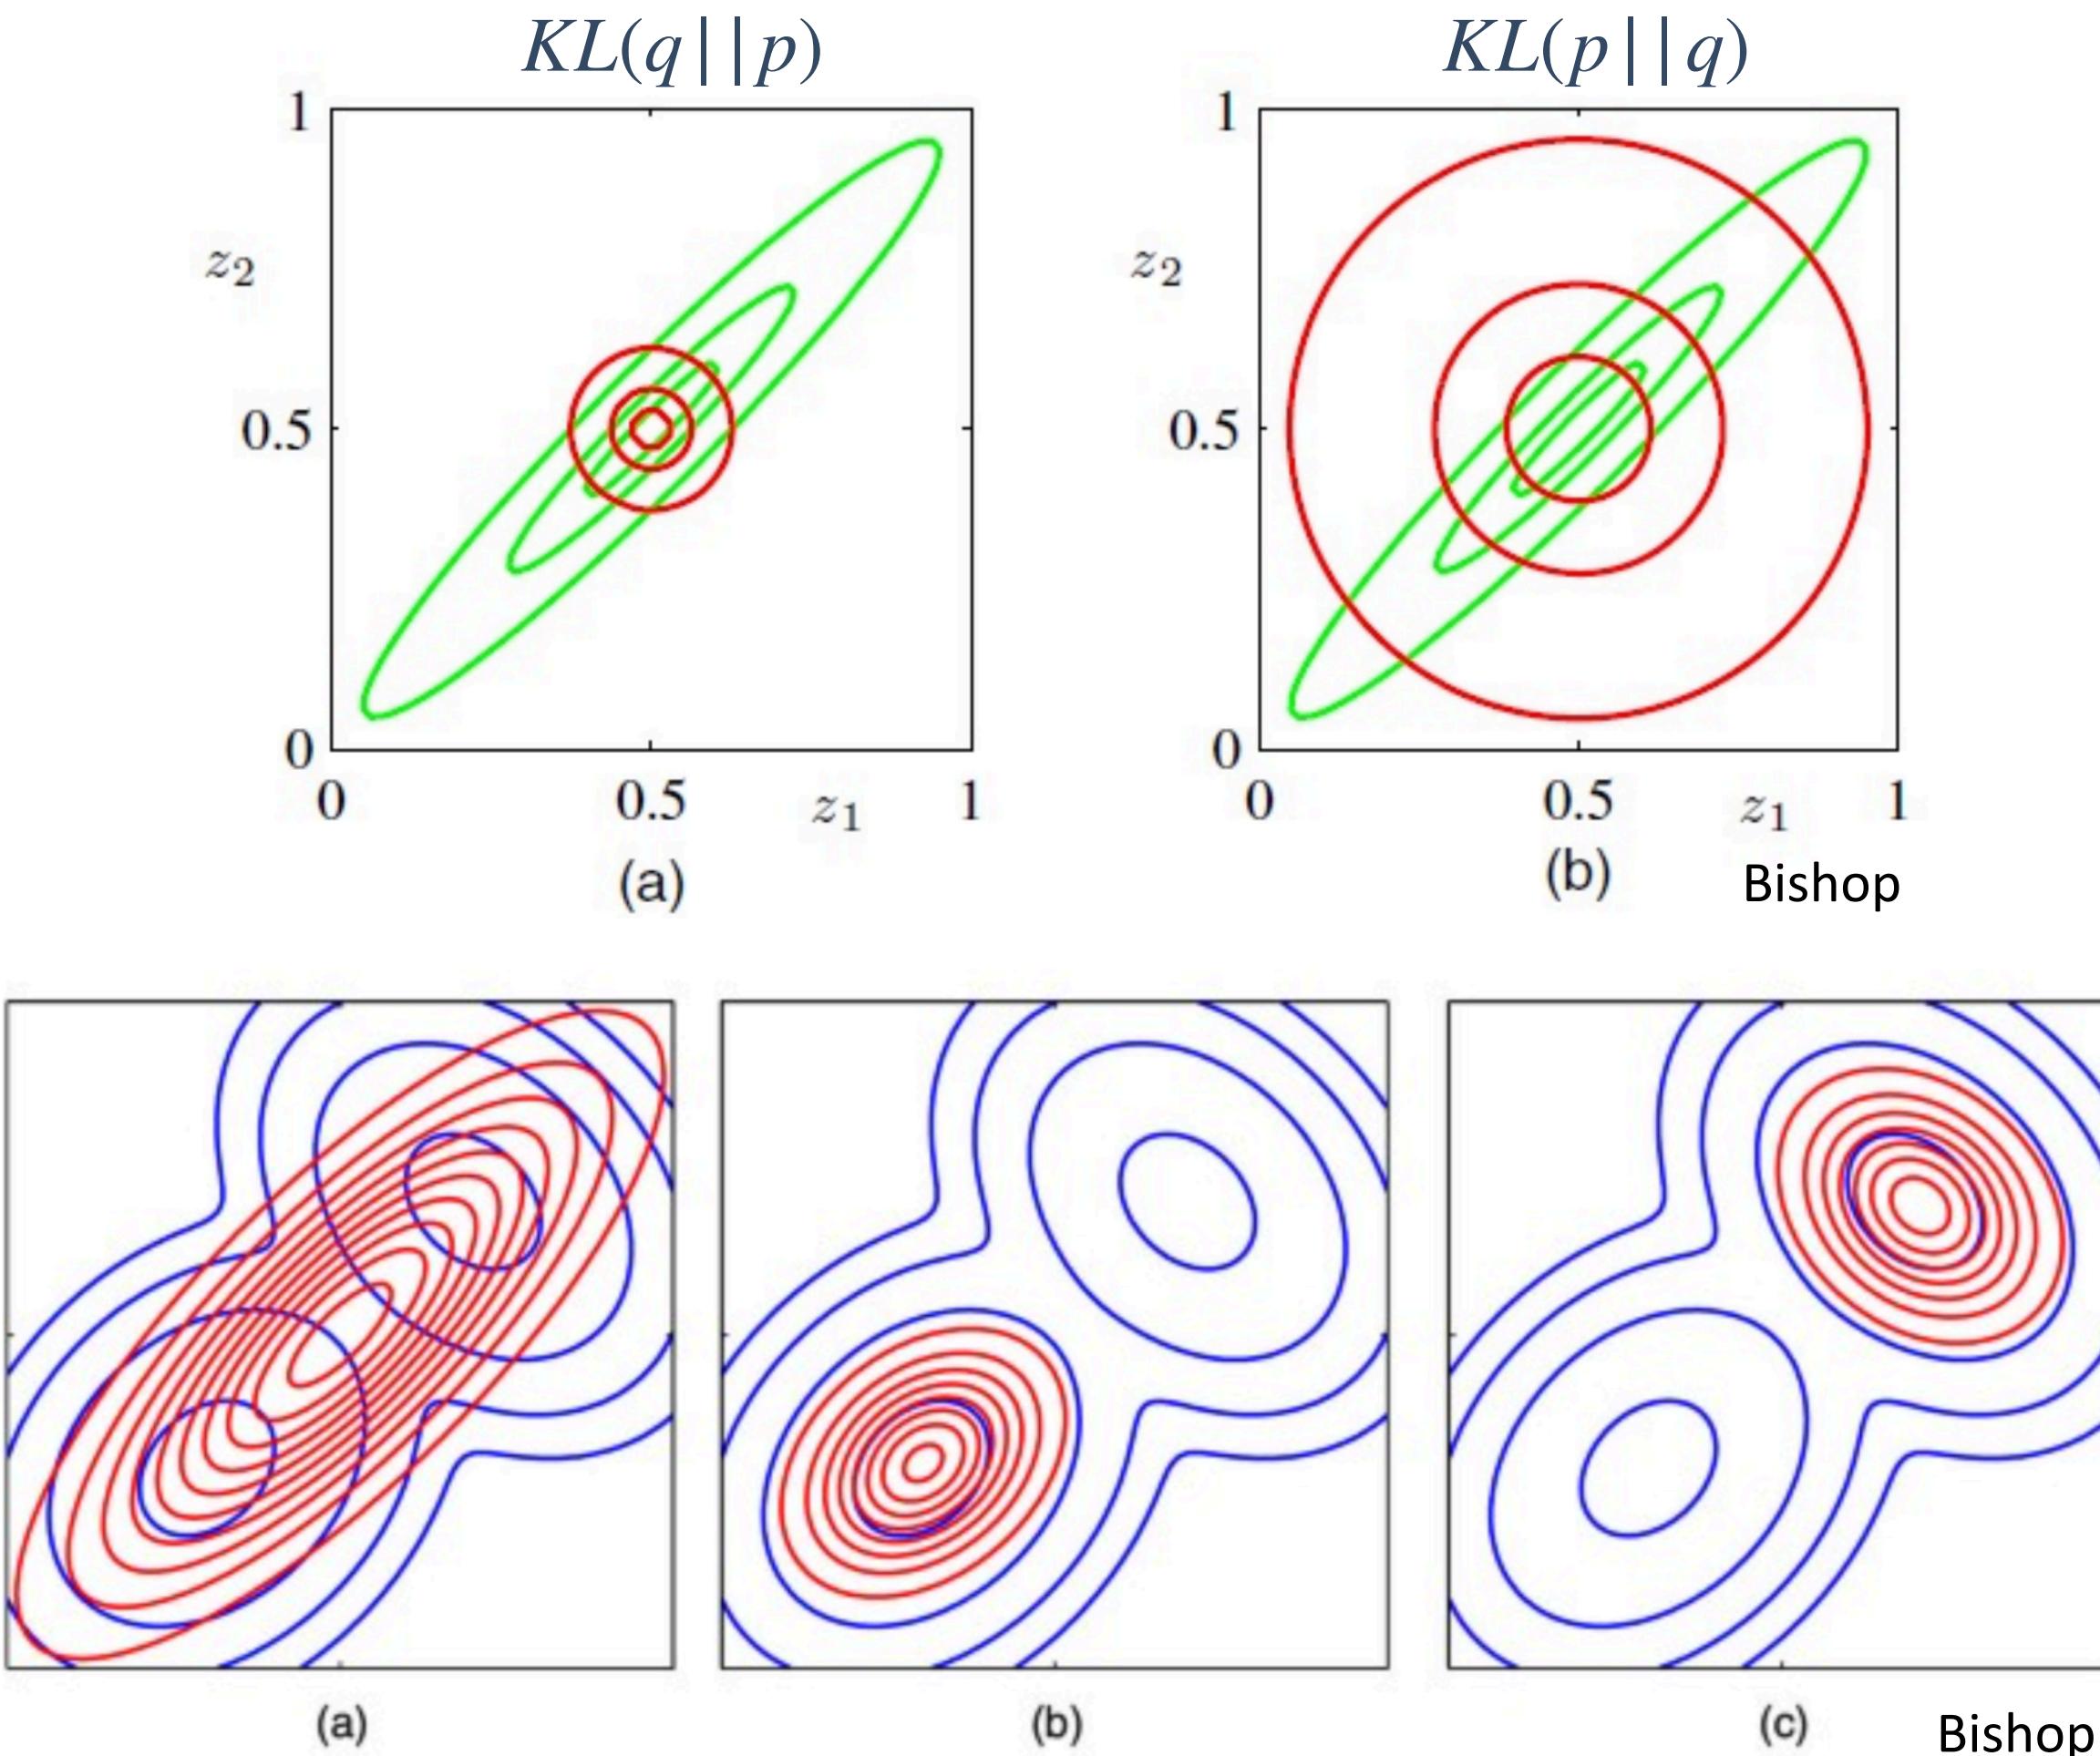
```
for  $j \in \{1, \dots, m\}$  do      until ELBO has converged
     $q_j(\theta_j) \propto \exp\{\mathbb{E}_{i \neq j}[\log(p(\theta_j | \theta_{i \neq j}, y))]\}$ 
    compute ELBO(q)
```

- Coordinate descent
- Stochastic variational inference (SVI) [Hoffman et al 2013]
- Automatic differentiation variational inference (ADVI) [Kucukelbir et al 2015, 2017]

# Approximate Bayesian Inference: Mean Field Variational Bayes

Green:  
1, 2, and 3 standard  
deviations for a correlated  
Gaussian distribution

Blue:  
Bimodal distribution given  
by a mixture of two  
Gaussians



Red:  
Same levels of an approximate  
distribution given by the product of  
two independent univariate  
Gaussians obtained by:  
a) minimizing  $KL(q \parallel p)$  divergence  
b) minimizing  $KL(p \parallel q)$  divergence

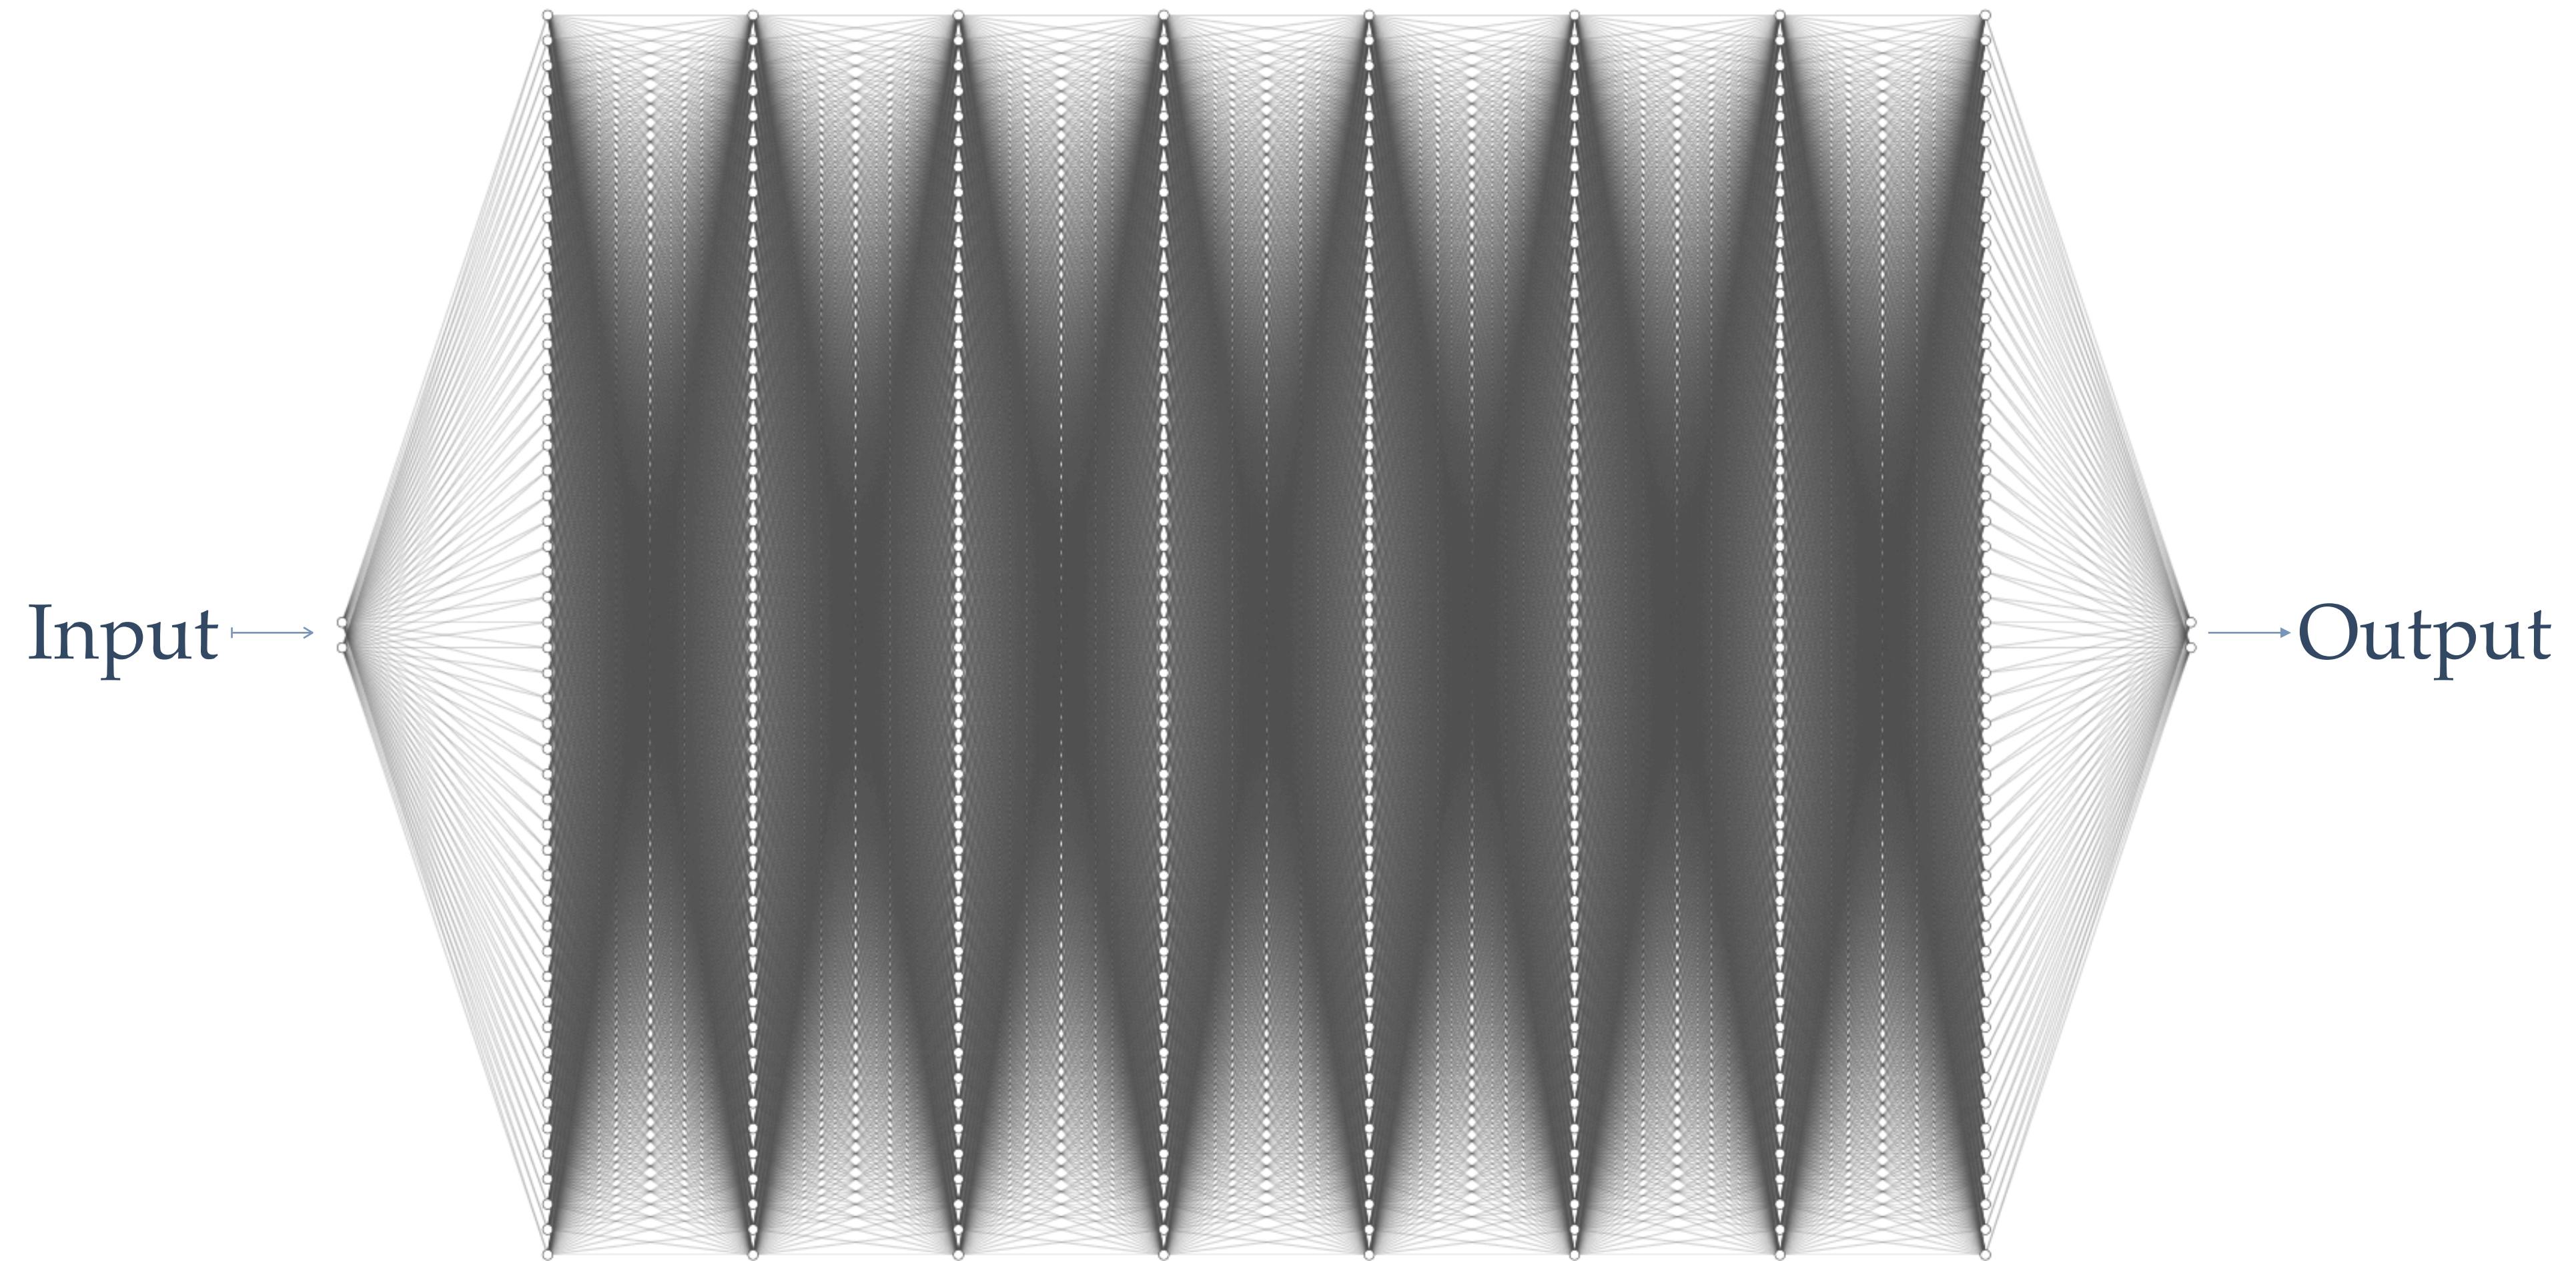
Red:  
a) Best approximation by a single  
Gaussian by minimizing the KL  
divergence  
b) same as (a) but numerically  
minimizing KL divergence  
c) same as in (b) but another local  
minimum of the KL divergence

# So far

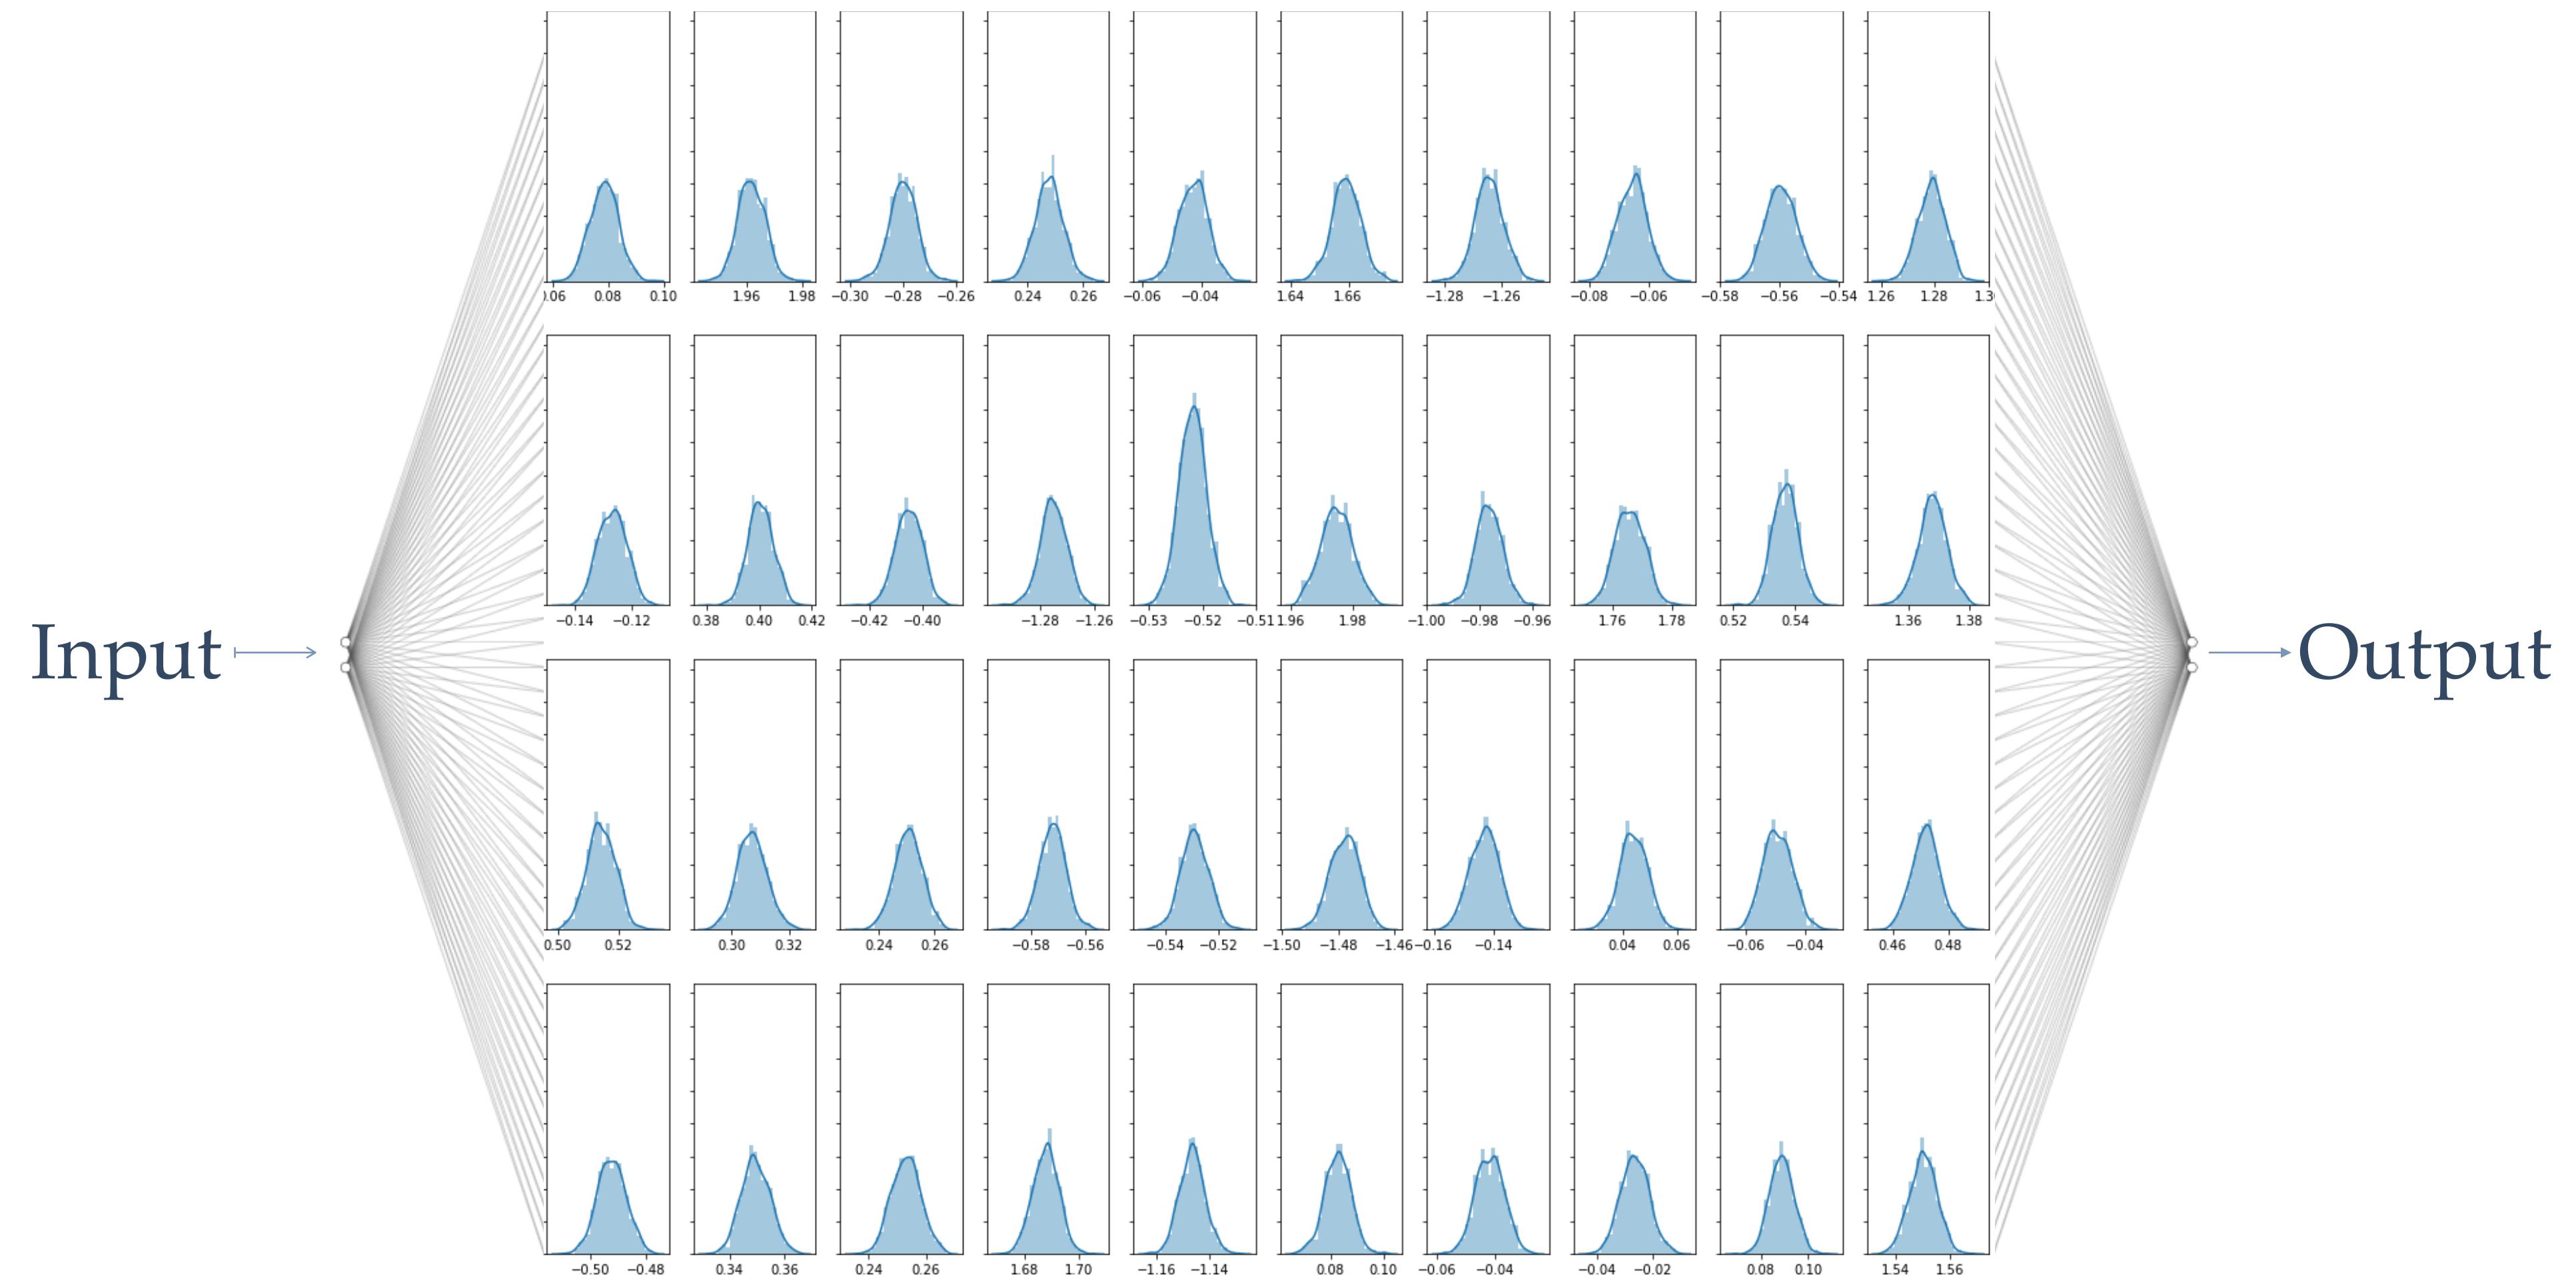
---

1. Bayesian Inference ✓
2. Markov Chain Monte Carlo ✓
3. Variational Inference ✓
4. Bayesian Neural Networks ✓
5. Bayes by backprop and flipout ✓
6. Mean Field Variational Bayes ✓
7. Application to Neural Networks
8. Drop Out as a Bayesian Approximation
9. Bootstrap for Inference
10. Measure performance of methods on uncertainty quantification

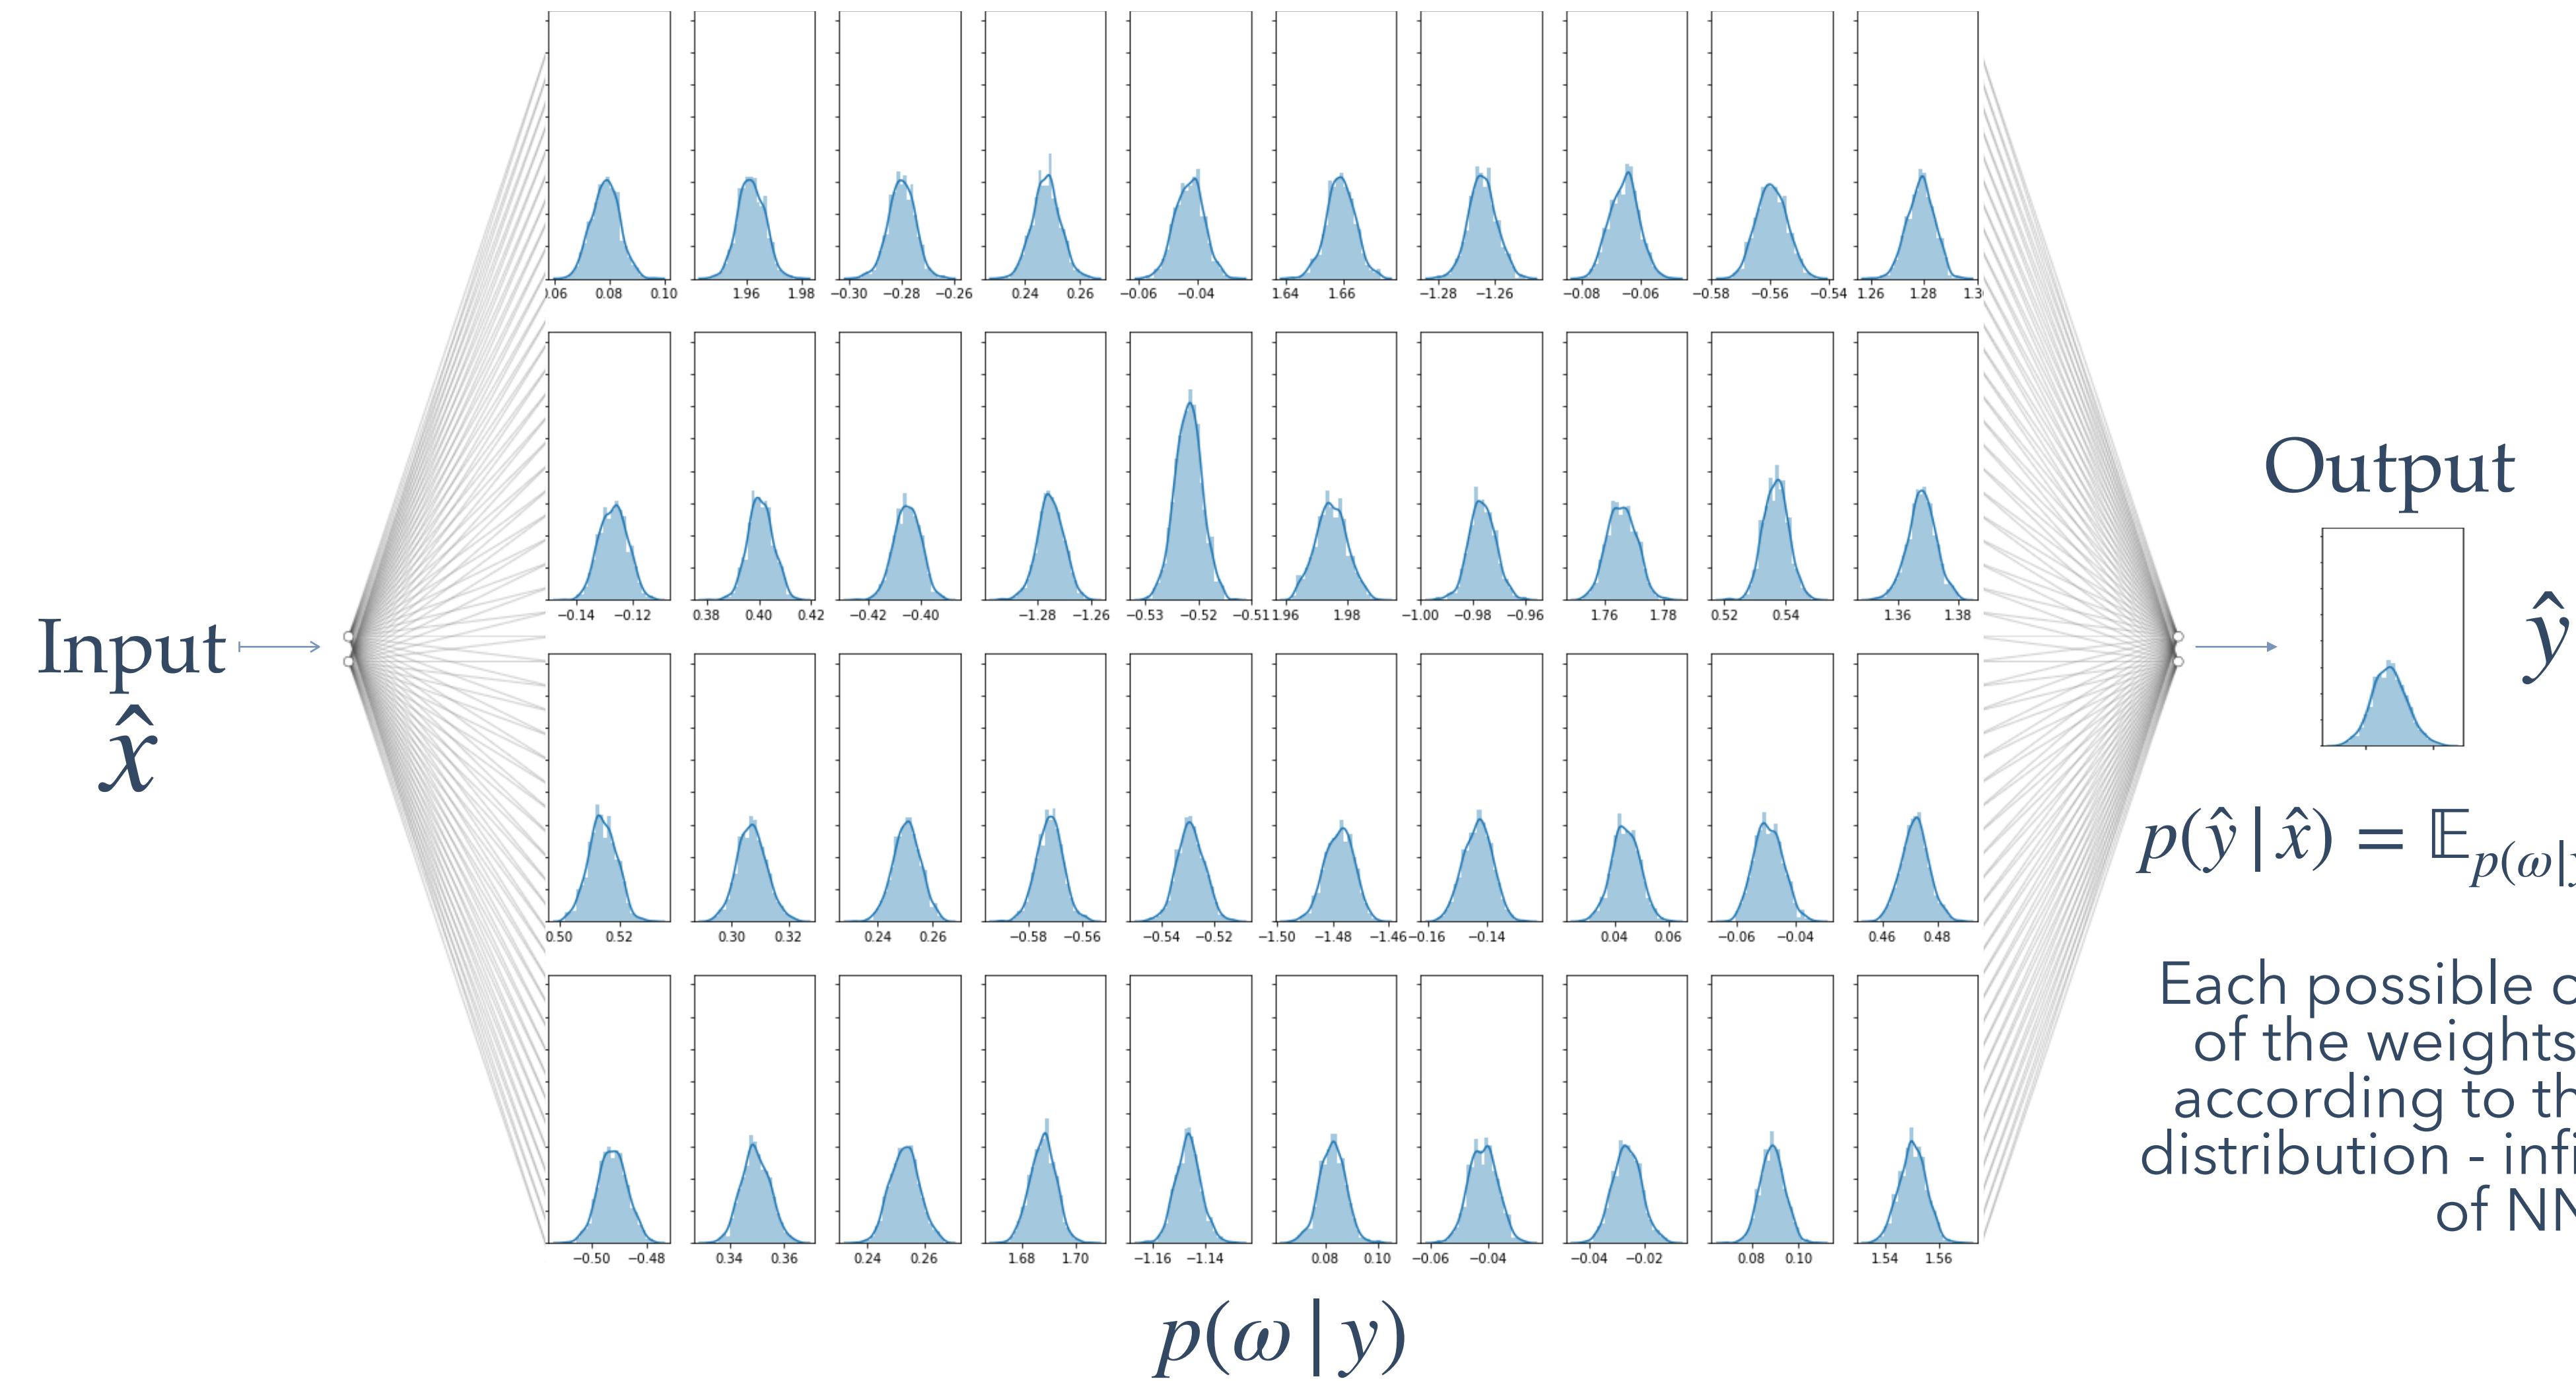
# Bayesian Inference for Neural Networks



# Bayesian Inference for Neural Networks



# Bayesian Inference for Neural Networks



BI for NN calculates the posterior distribution of the weights given the training data

# Bayesian Neural Networks

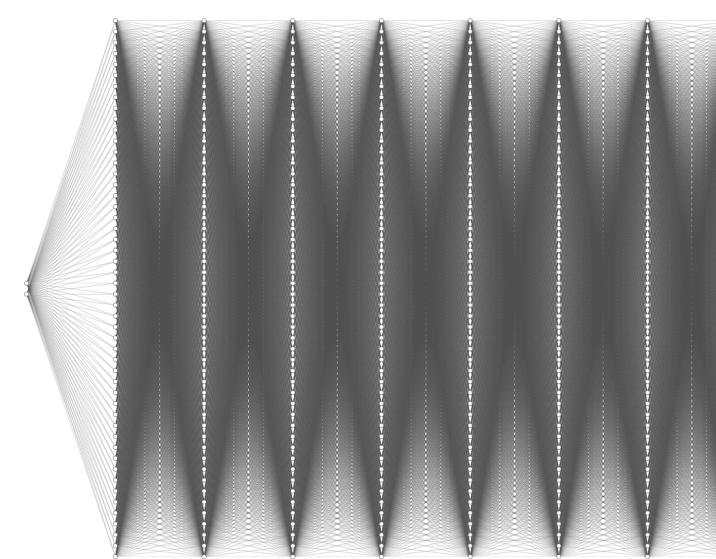
Priors  $p(\theta)$

$$\frac{p(\theta) p(y|\theta)}{p(y)} = p(\theta|y)$$

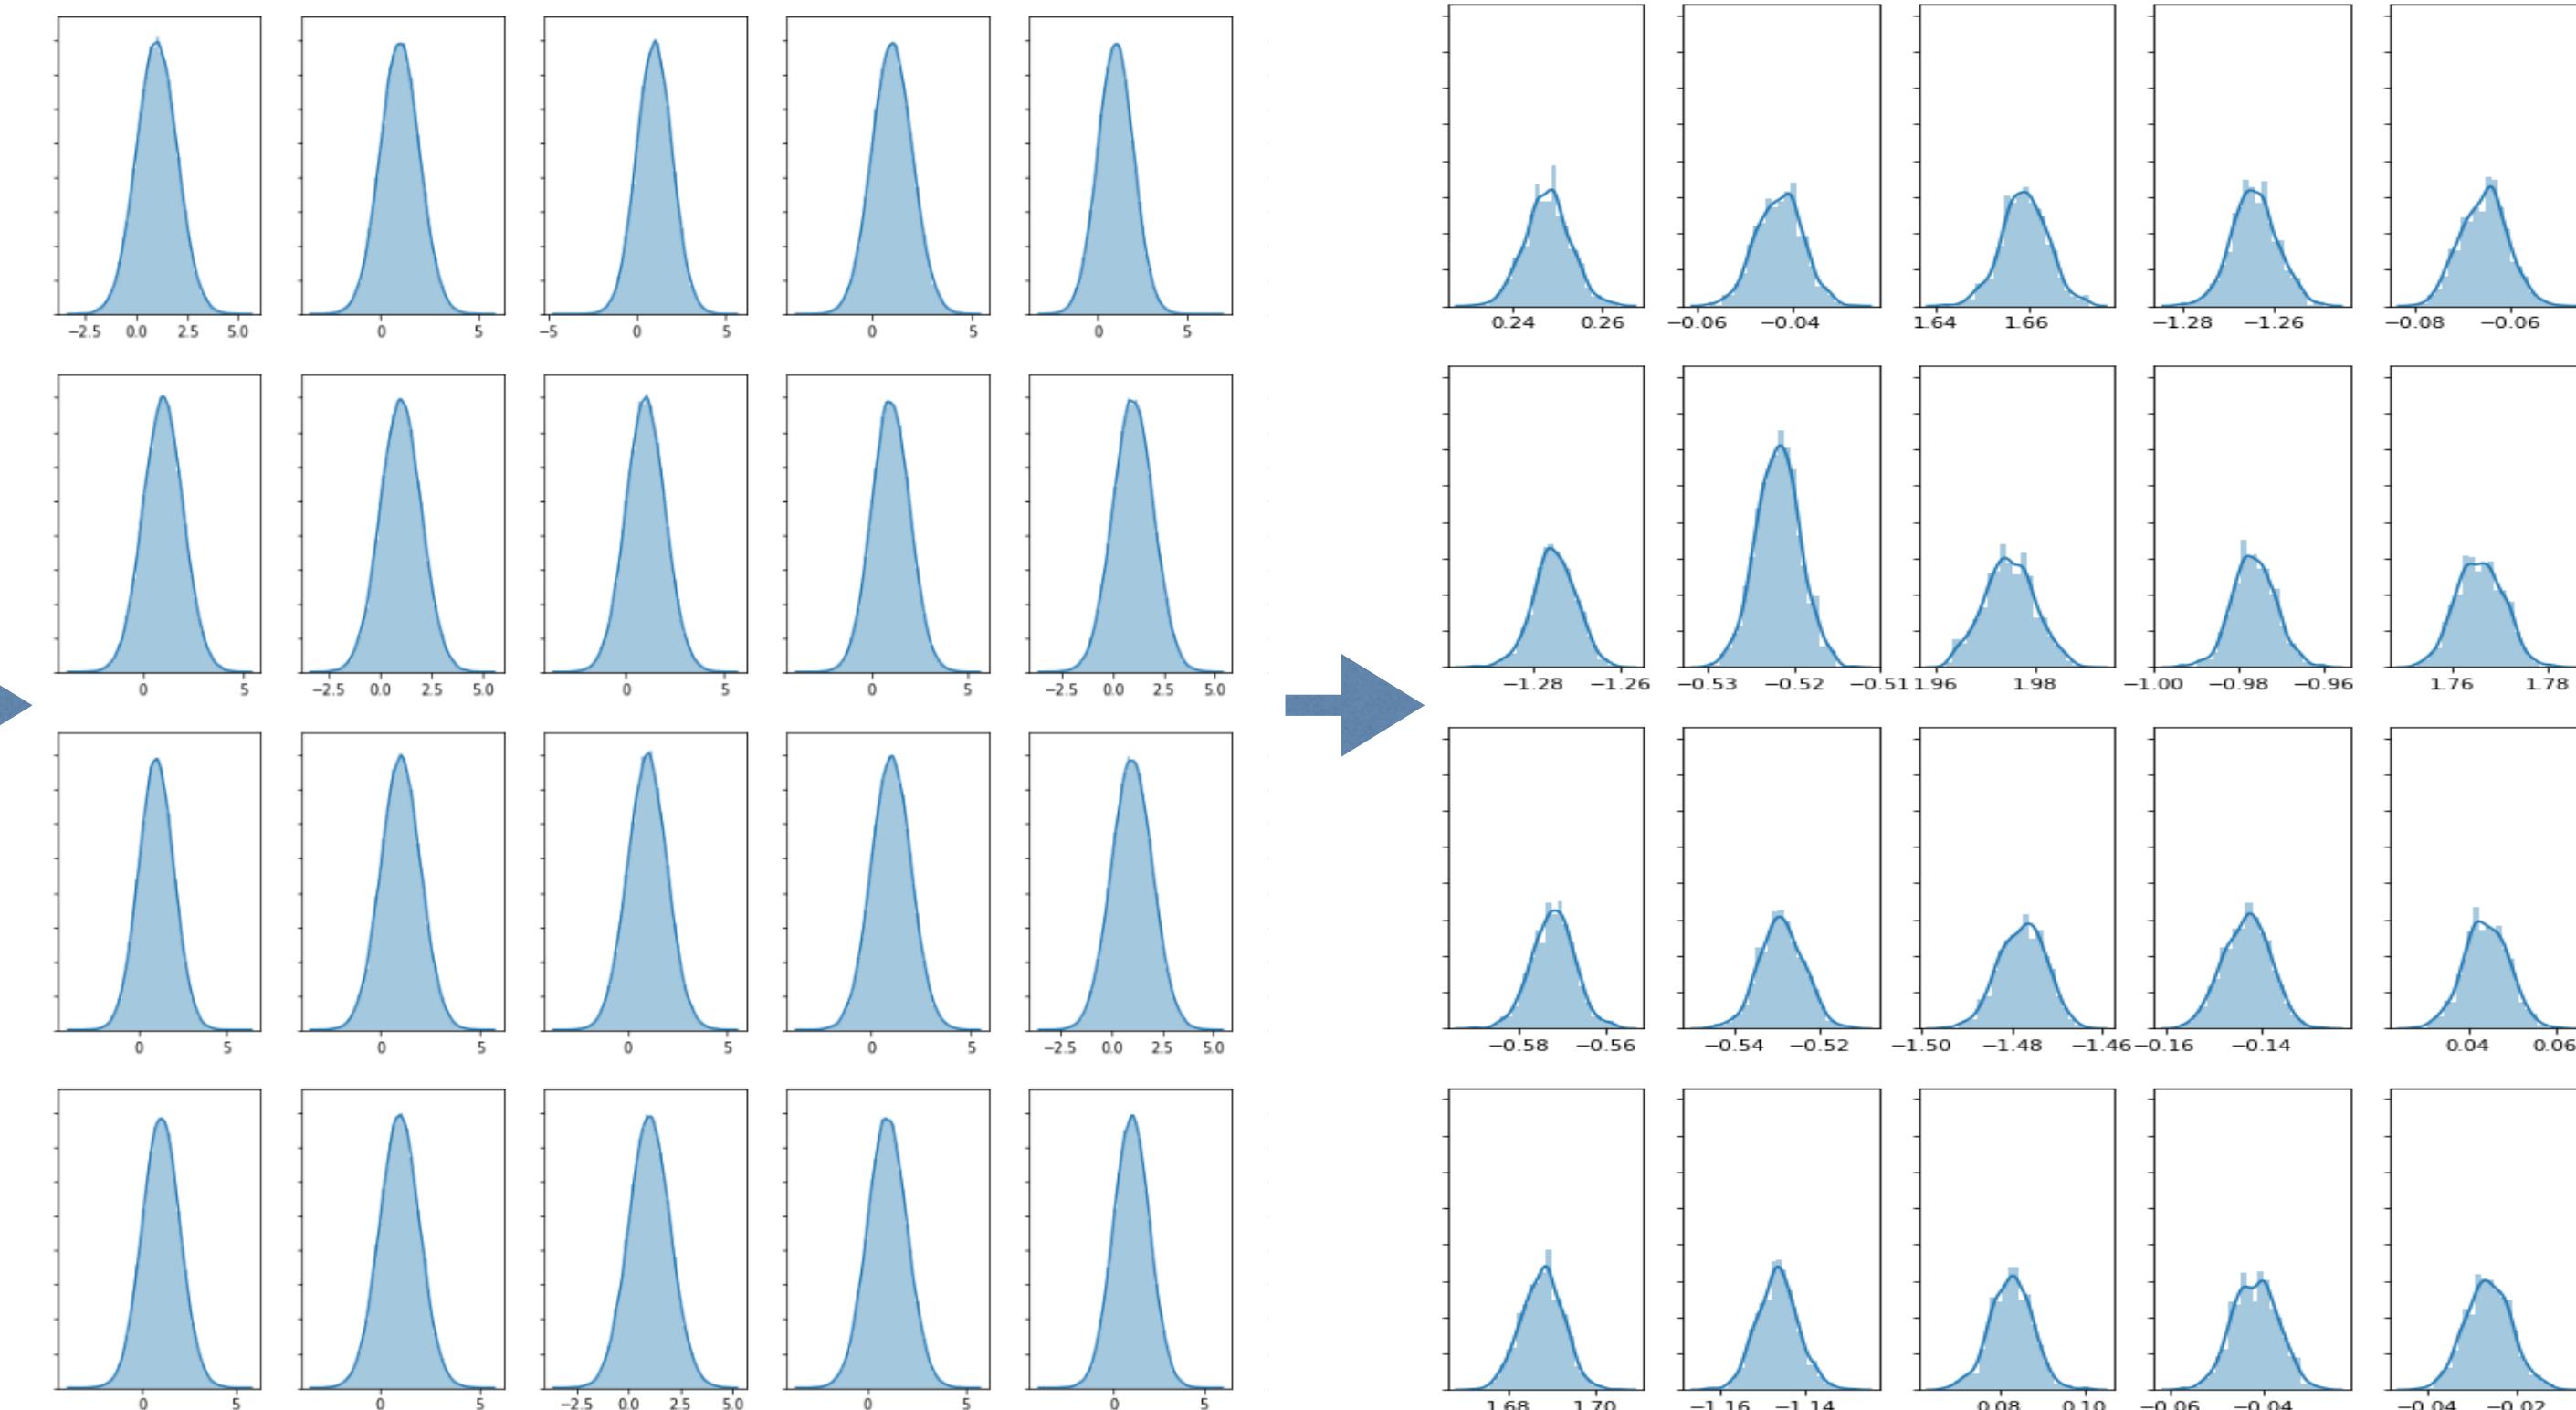
intractable

Means

FCNN



&  
Scale



# Bayesian Neural Networks

MCMC: Eventually accurate

$$p(\theta) \times p(y | \theta) \propto p(\theta | y)$$

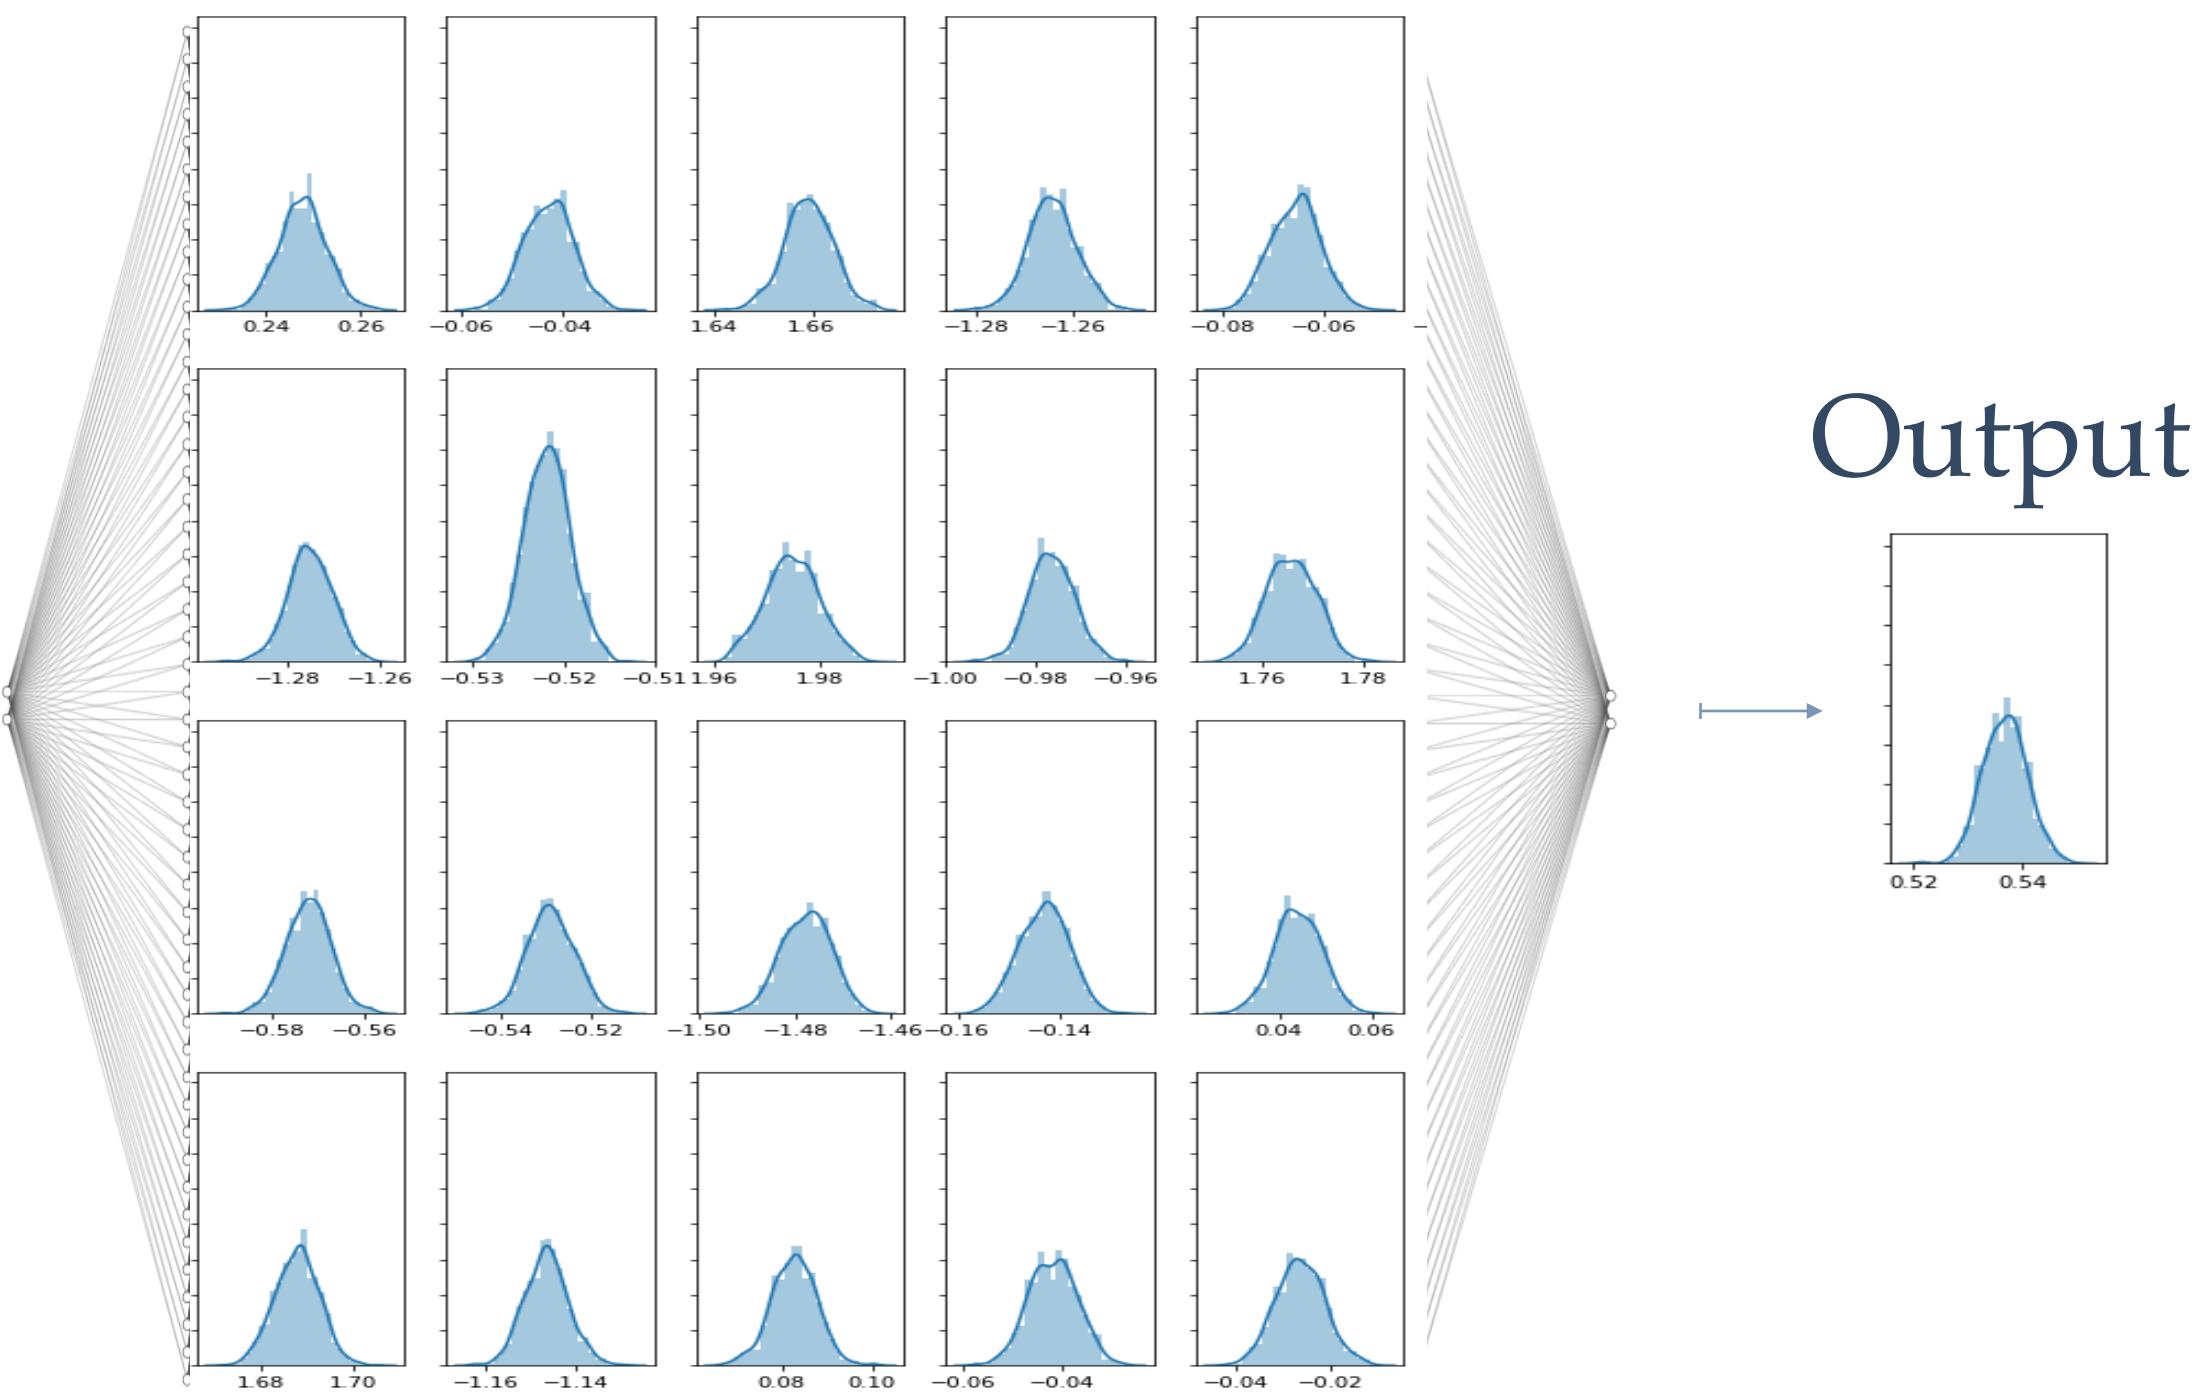
Only forward passes

$$\theta = (\mu, \sigma)$$

$$\mu^j \sim N(\mu^{(j-1)}, \sigma)$$

Step of random walk

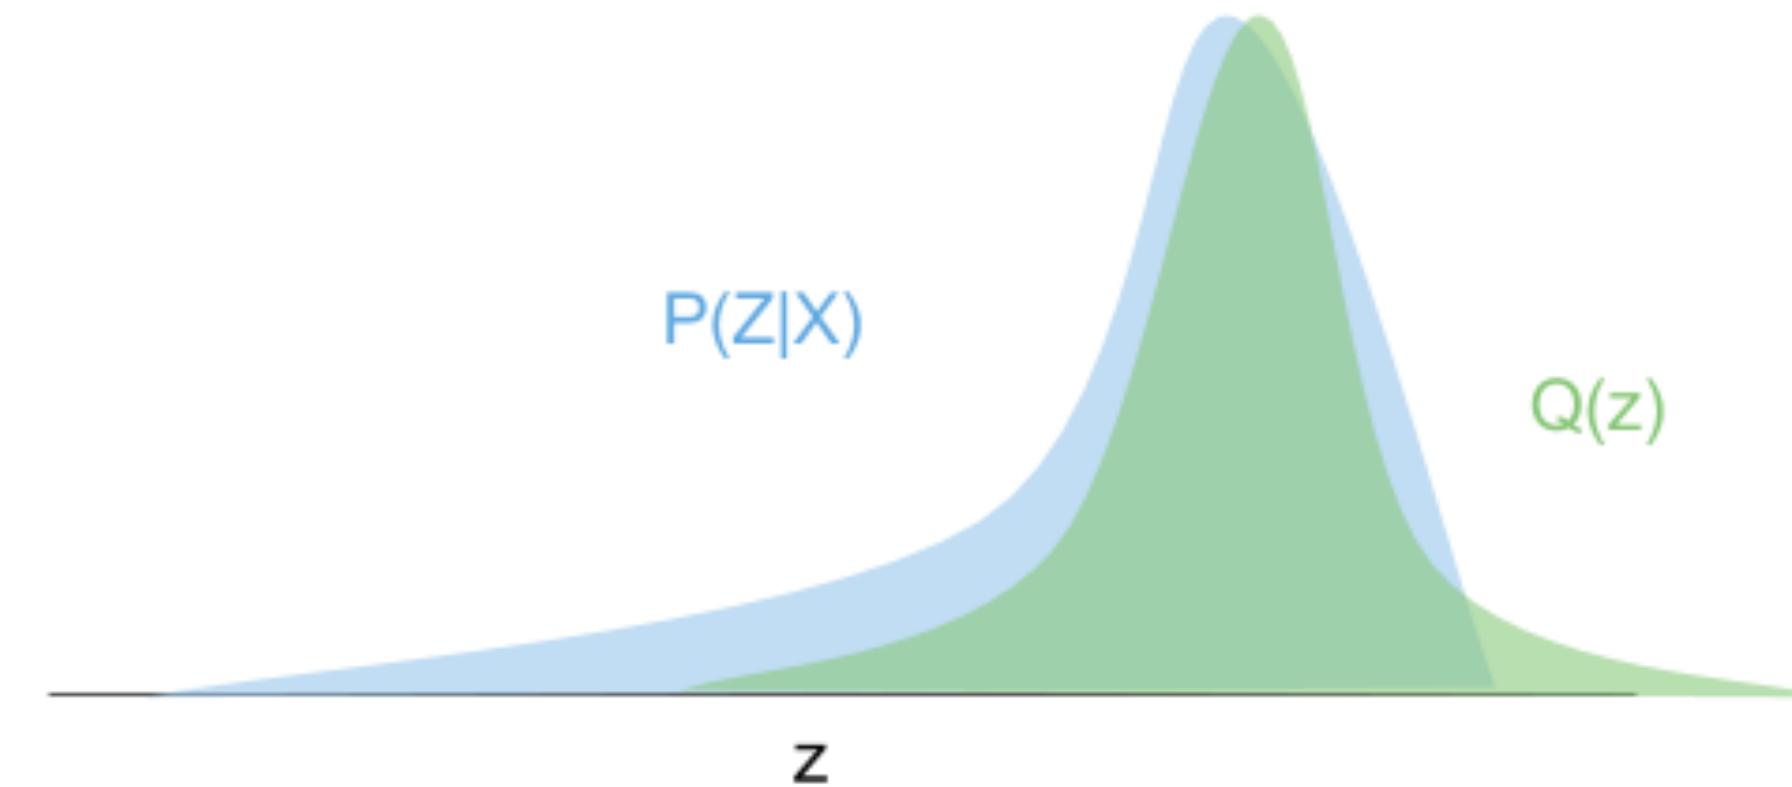
Input



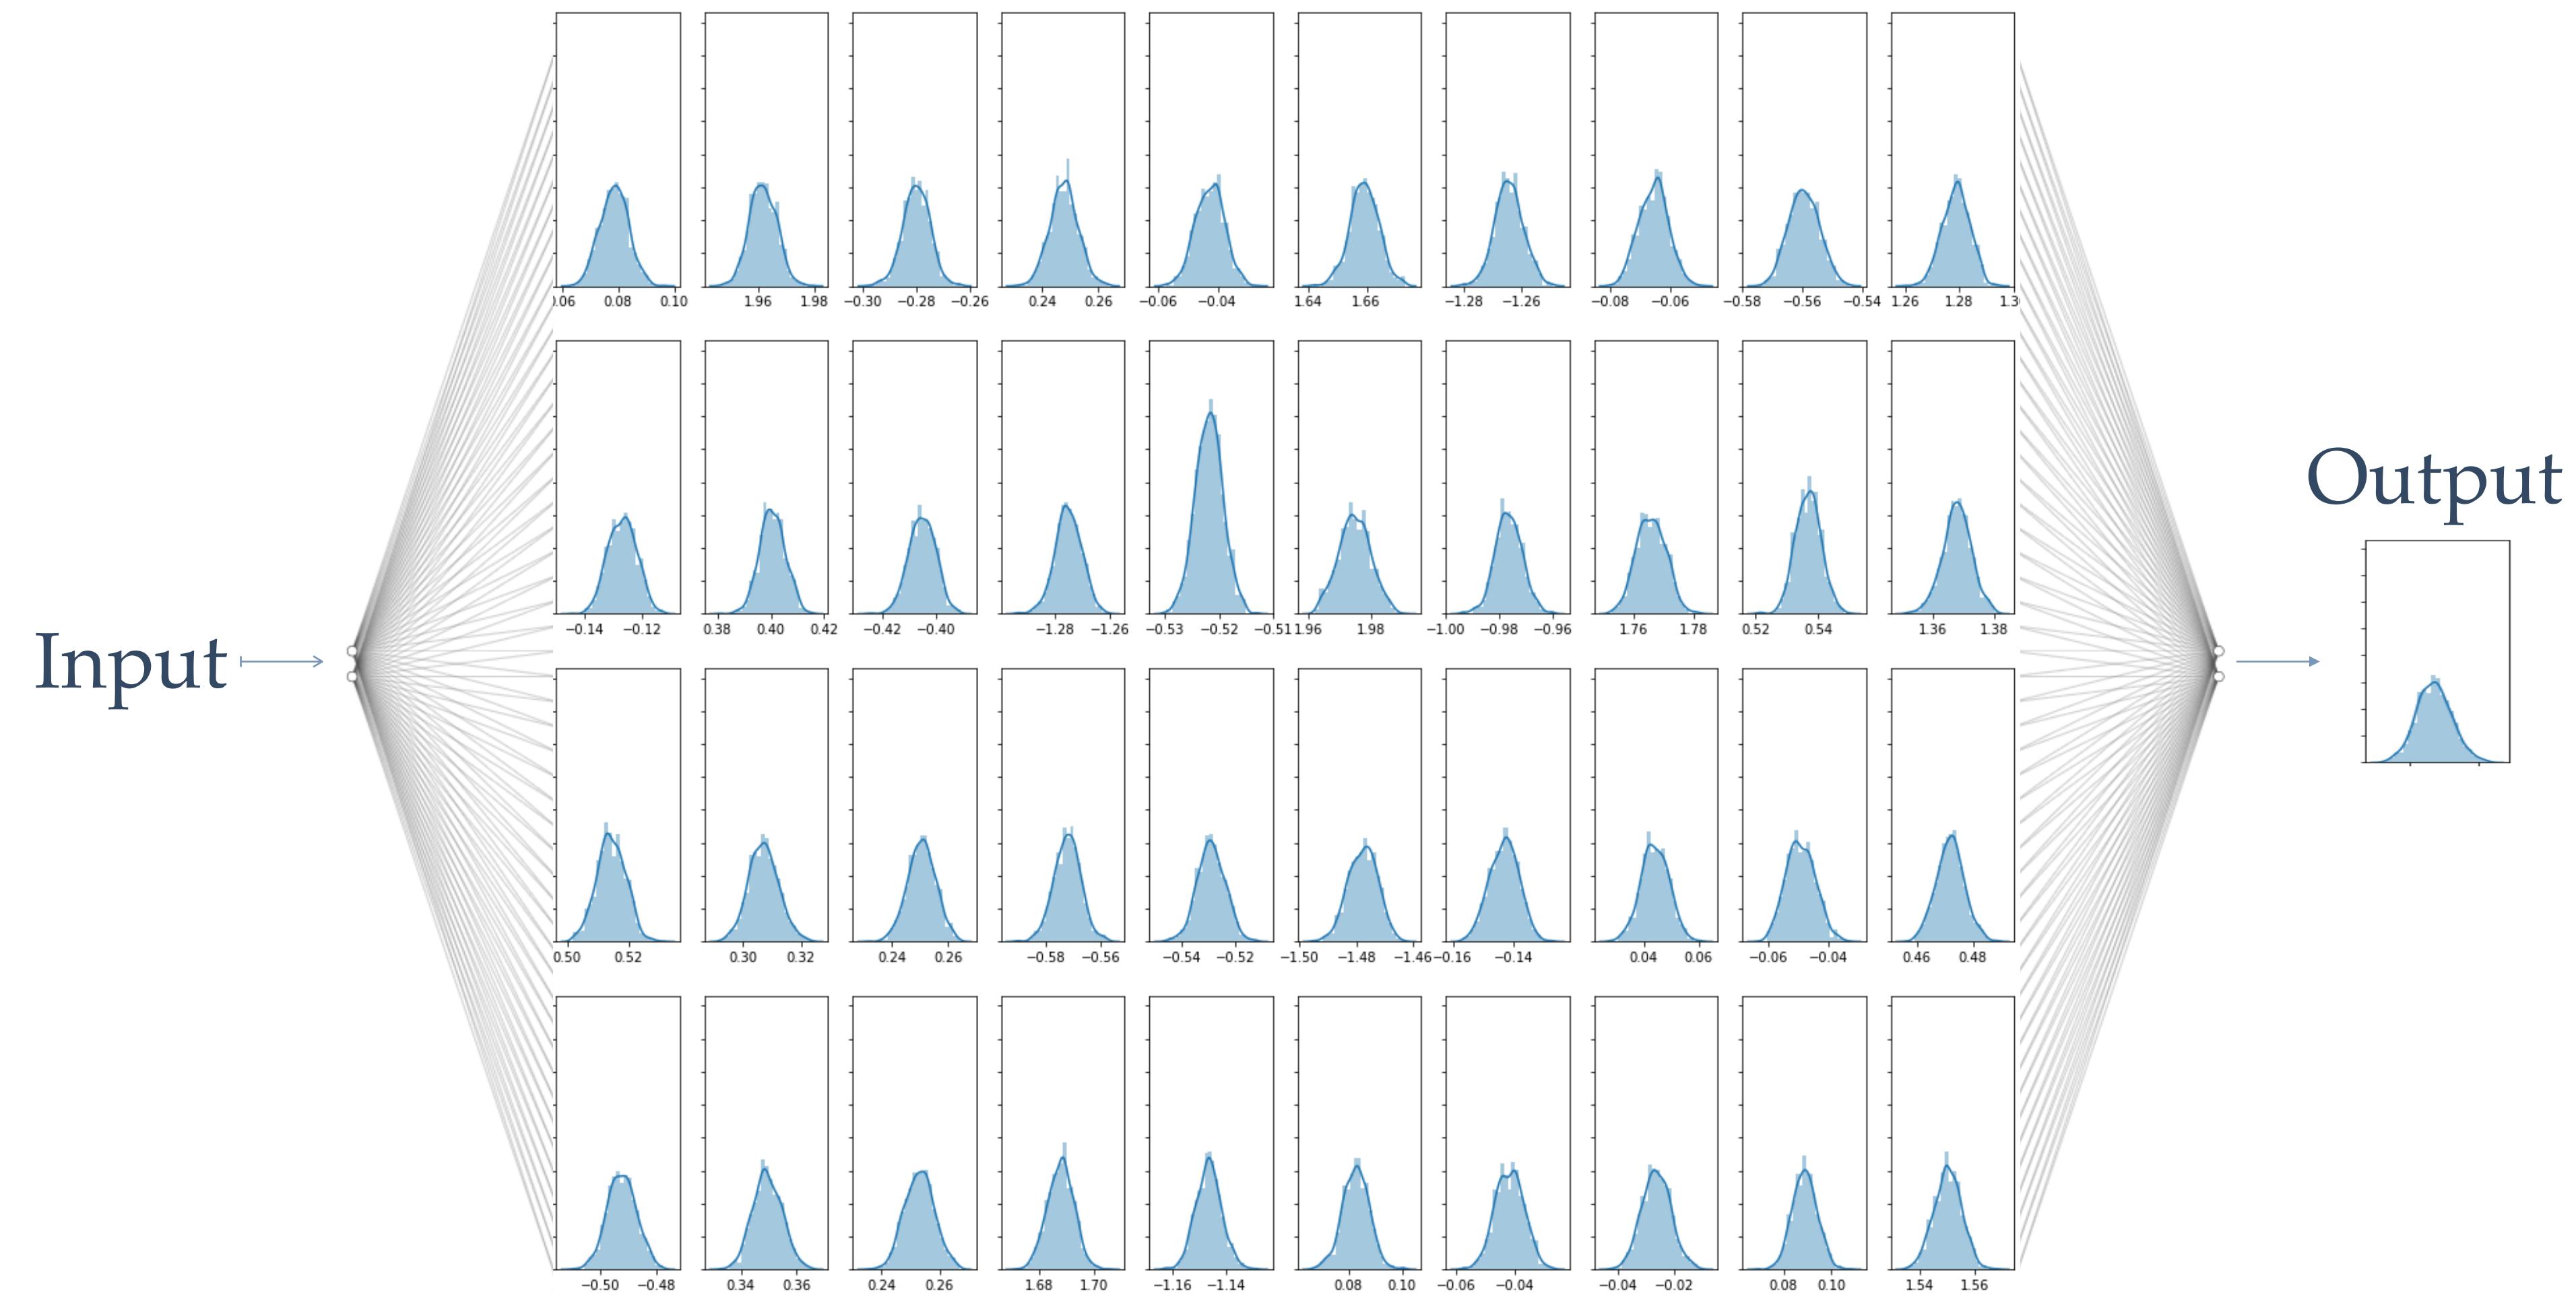
# Approximate Bayesian Inference: Variational Bayes

$$p(\theta | y) = \frac{p(y | \theta) p(\theta)}{\int p(y, \theta) d\theta} \approx q^*$$

Optimization approach →  $Q$  a family of “nice” distributions



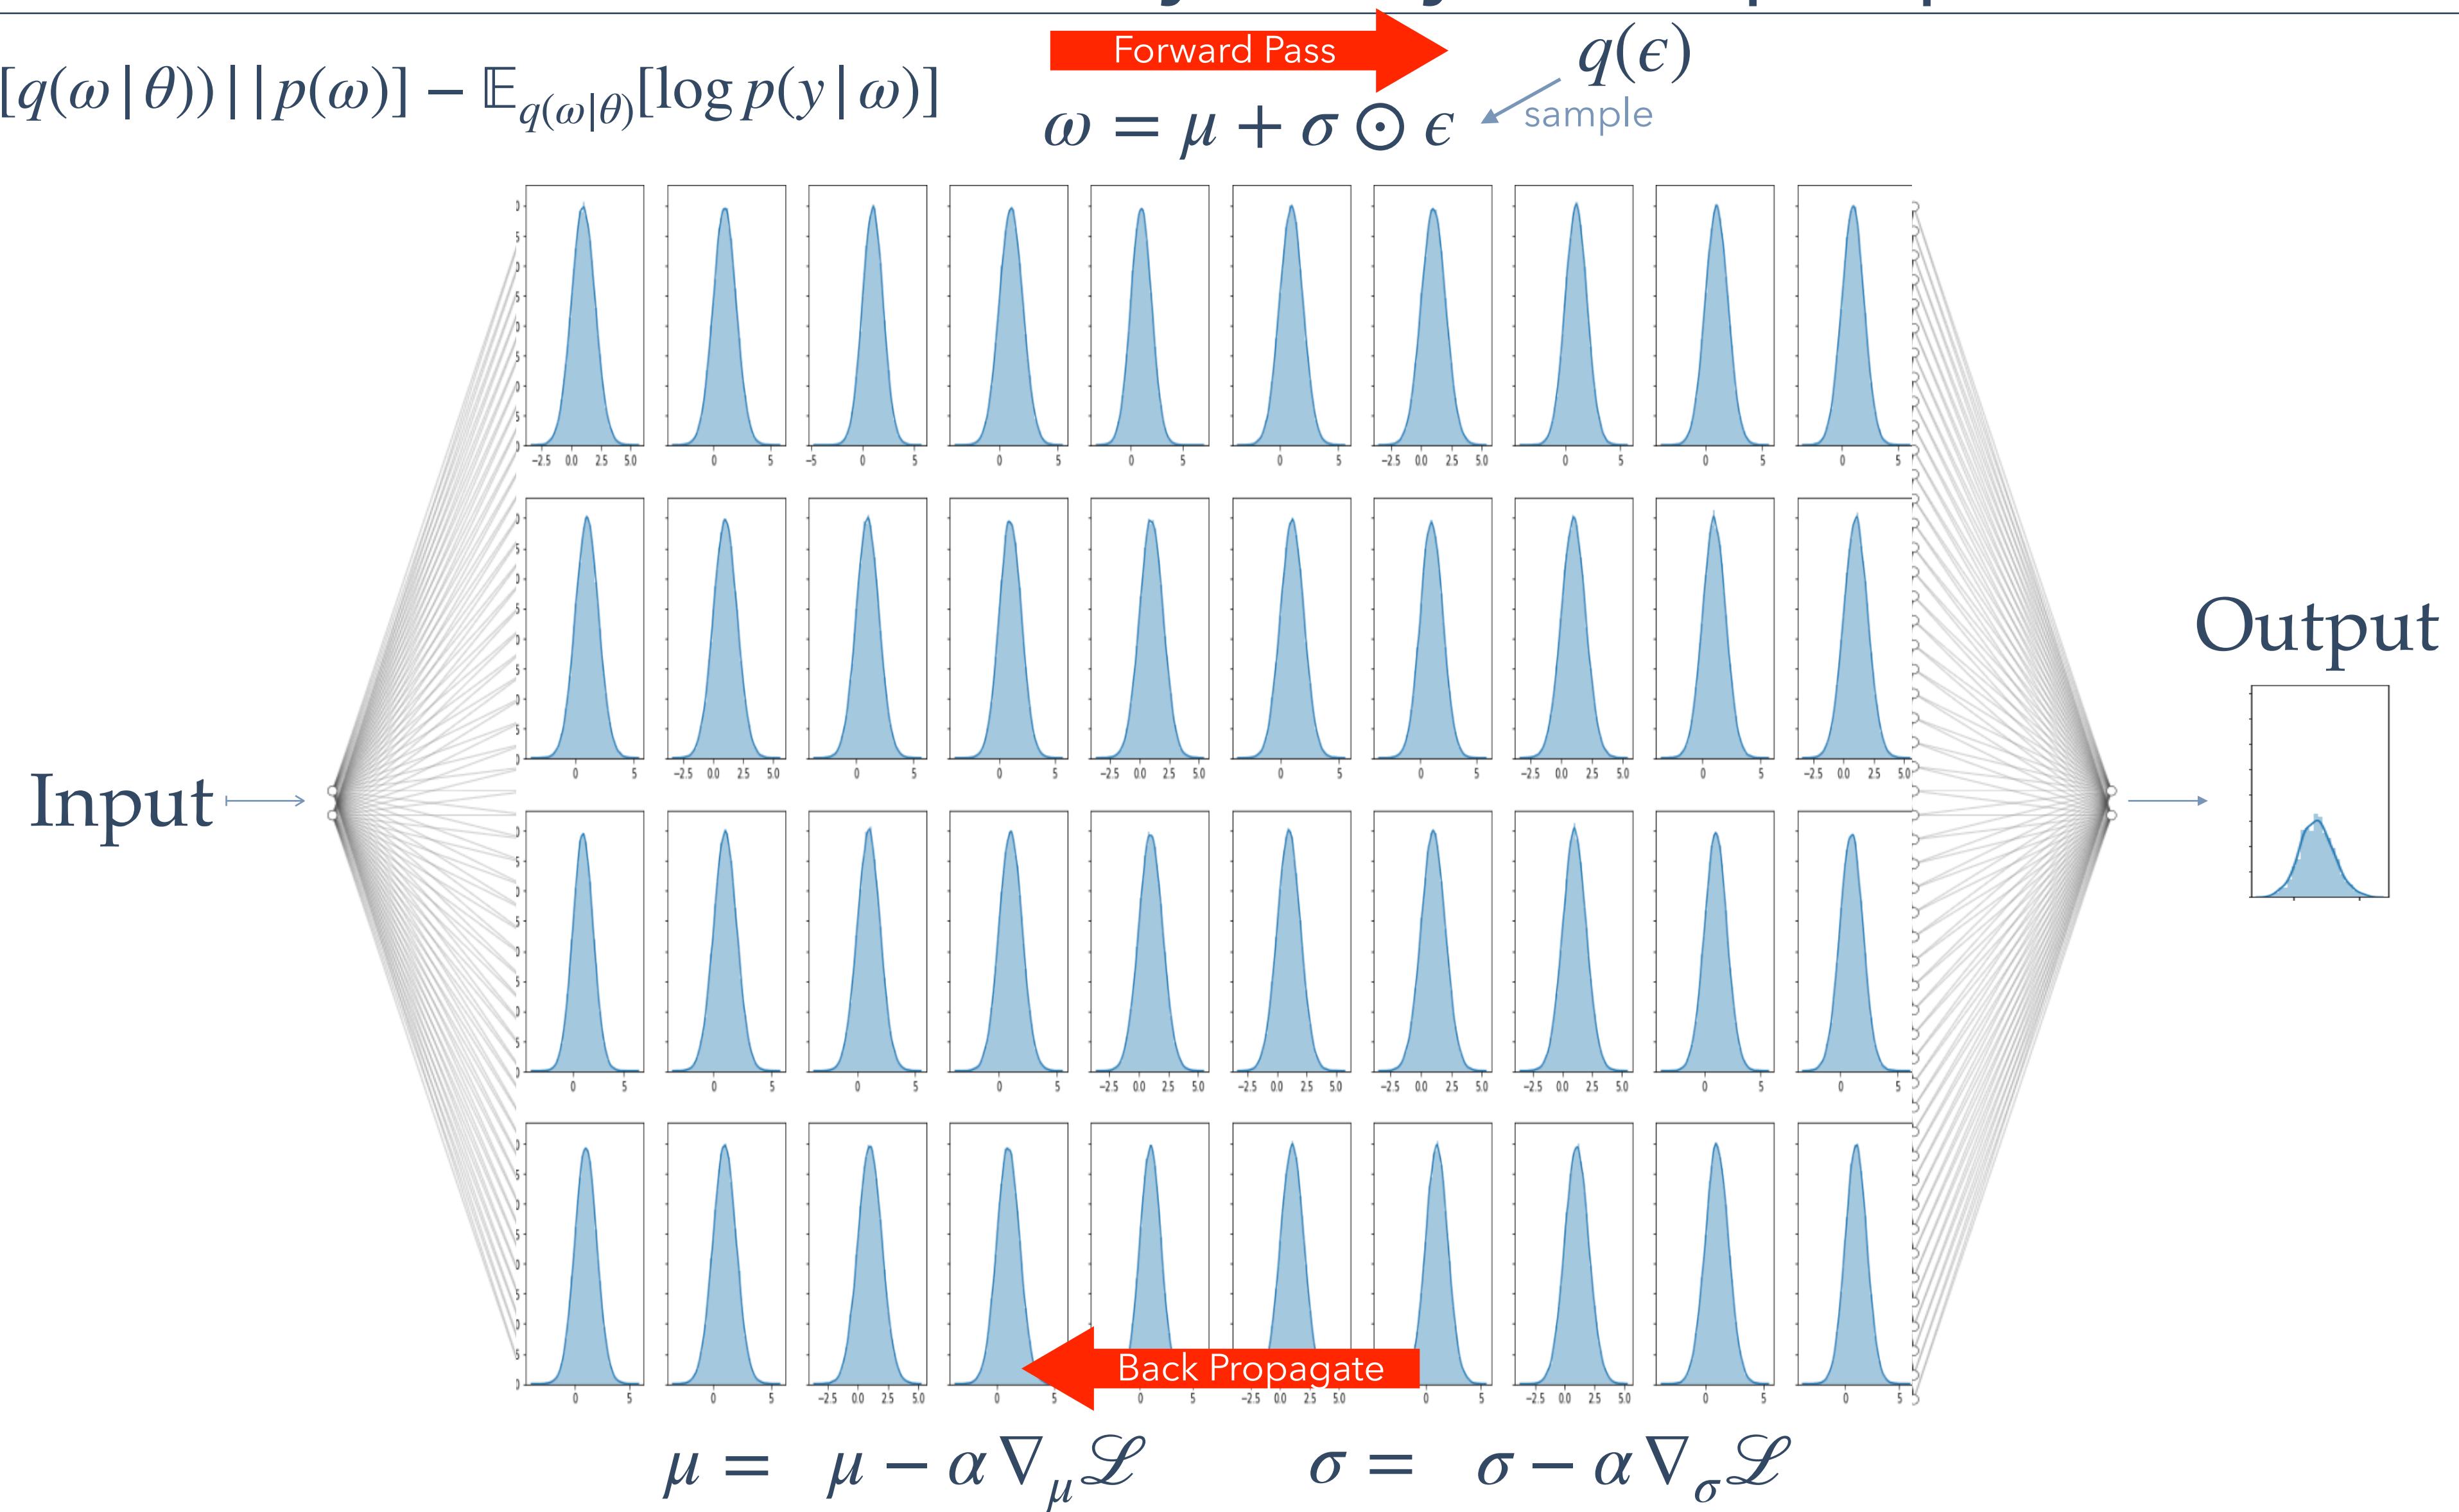
# Bayesian Neural Network



# Bayesian Neural Network: Bayes by Backprop

$$\mathcal{L}(y|\theta) = KL[q(\omega|\theta)||p(\omega)] - \mathbb{E}_{q(\omega|\theta)}[\log p(y|\omega)]$$

Forward Pass →  
 $\omega = \mu + \sigma \odot \epsilon$  ← sample  
 $q(\epsilon)$



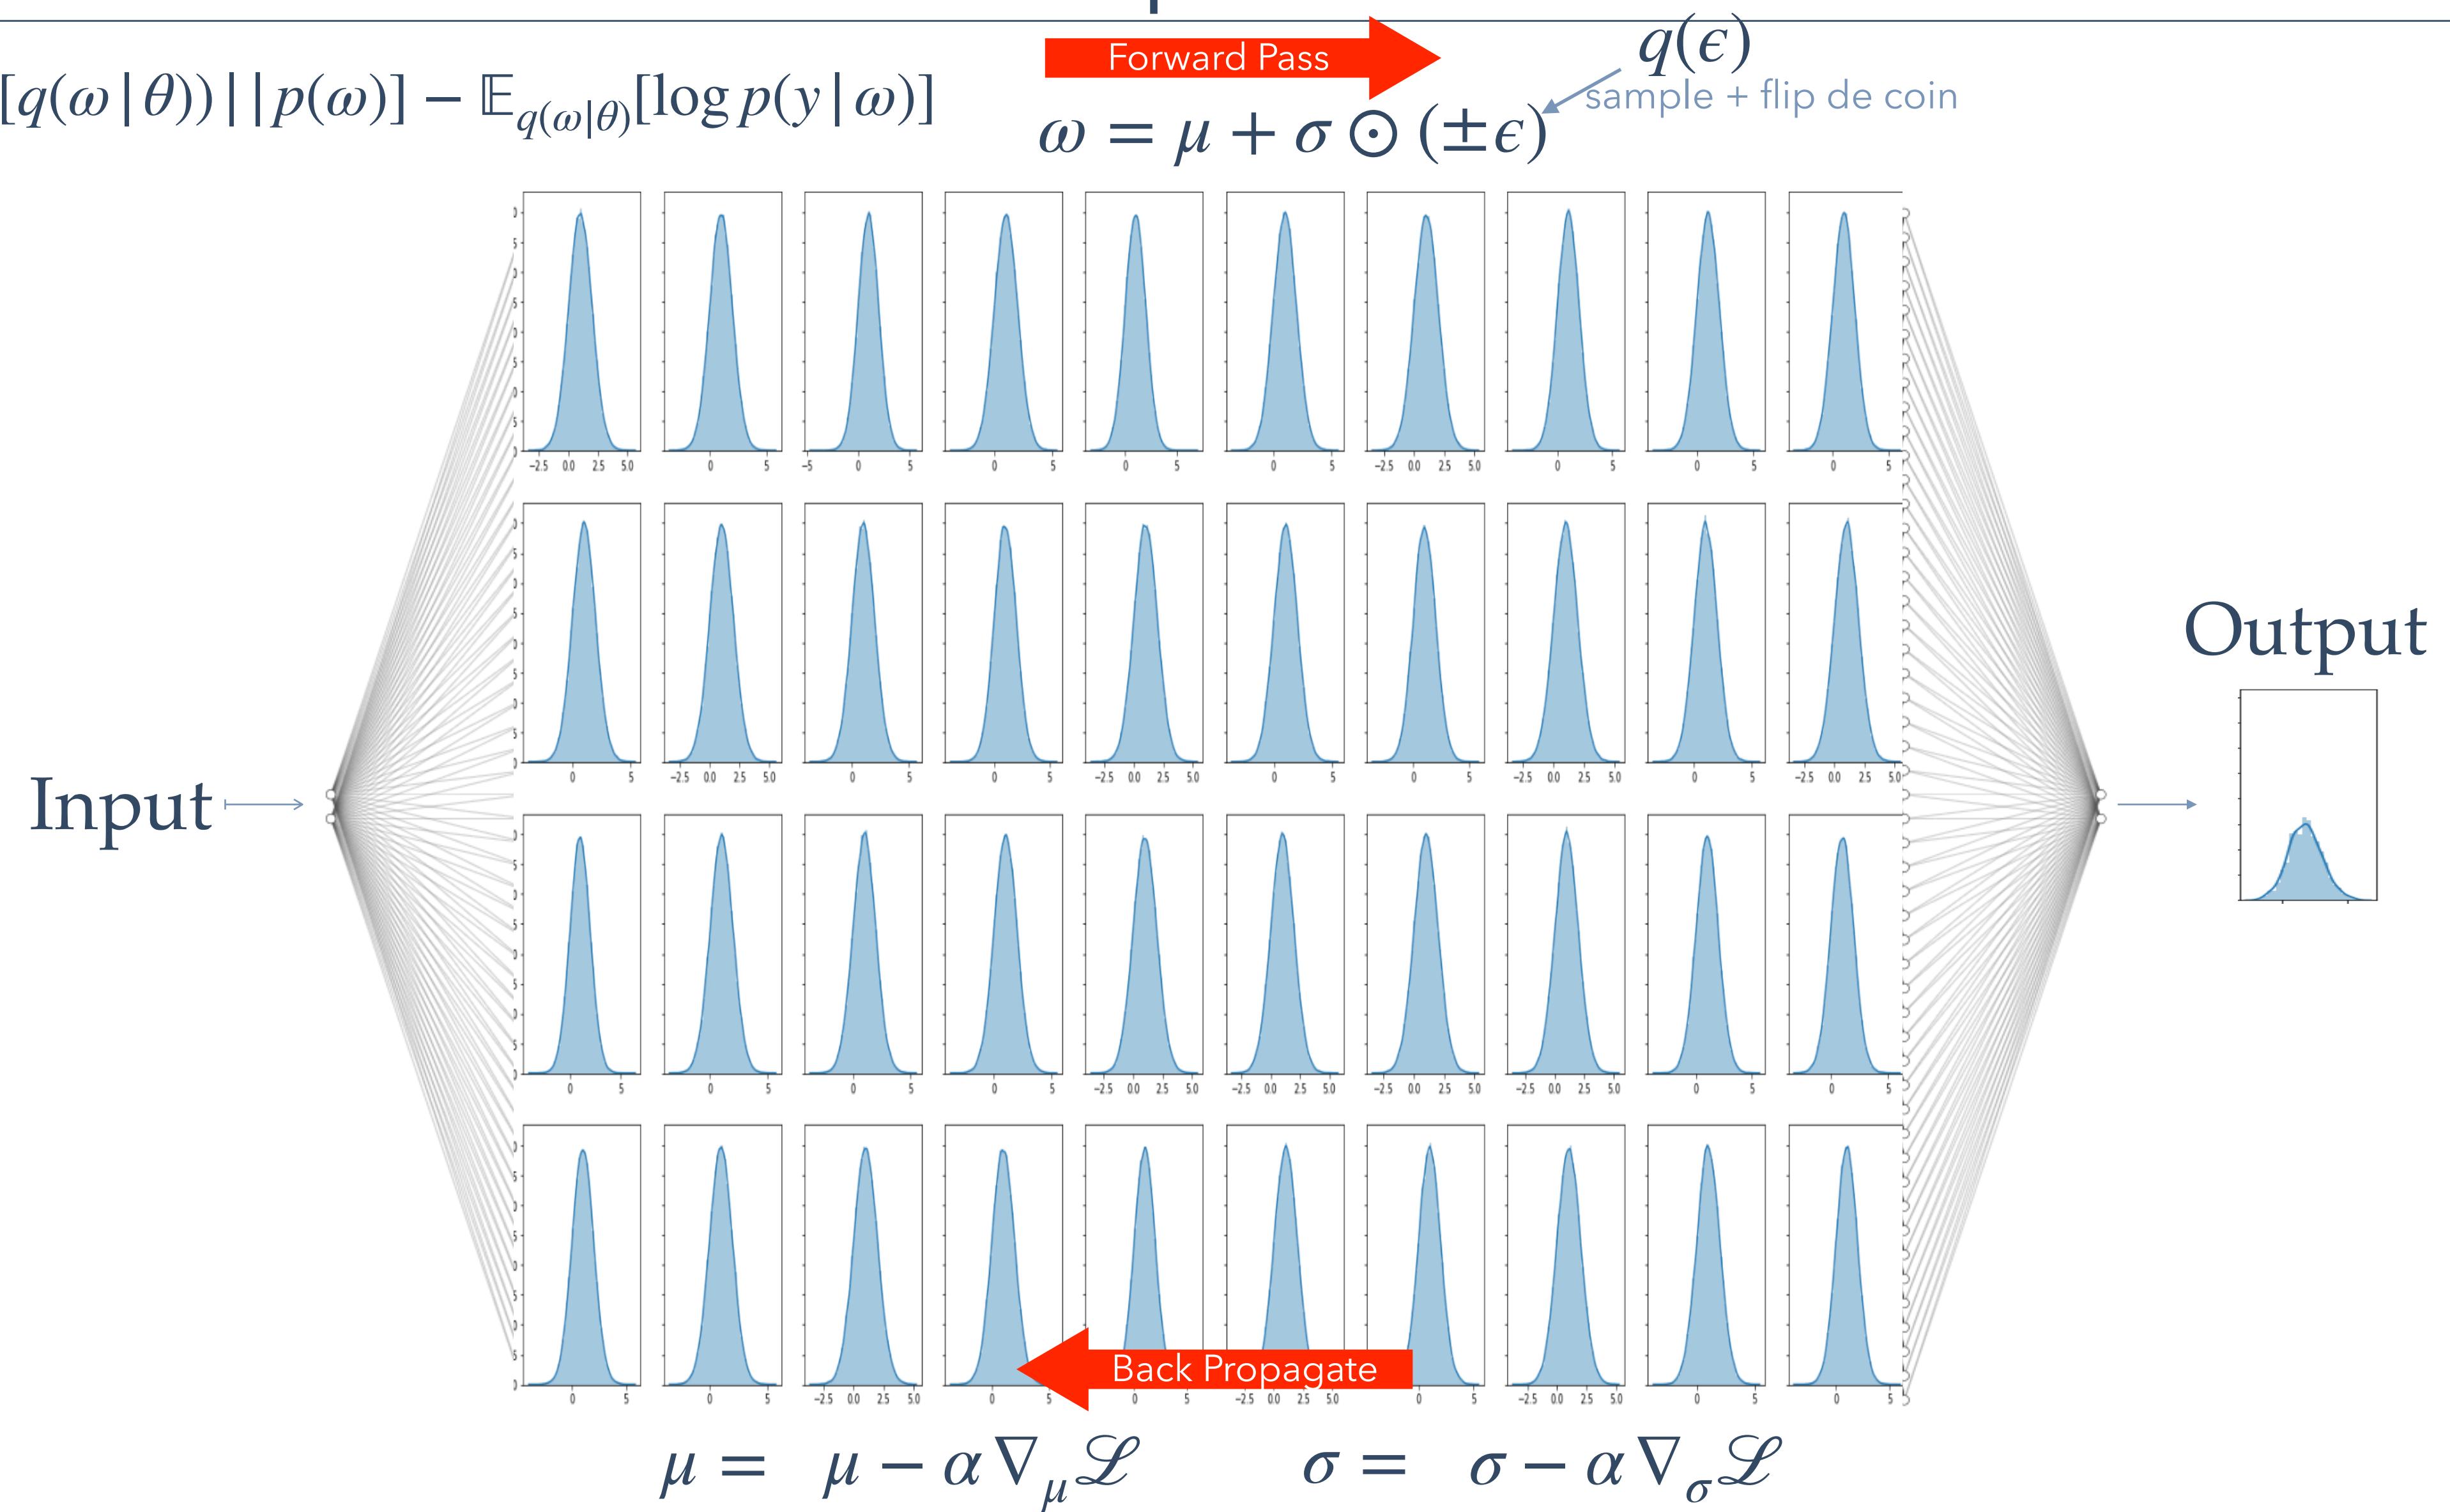
# Bayesian Neural Network: Flipout

$$\mathcal{L}(y|\theta) = KL[q(\omega|\theta)||p(\omega)] - \mathbb{E}_{q(\omega|\theta)}[\log p(y|\omega)]$$

Forward Pass →

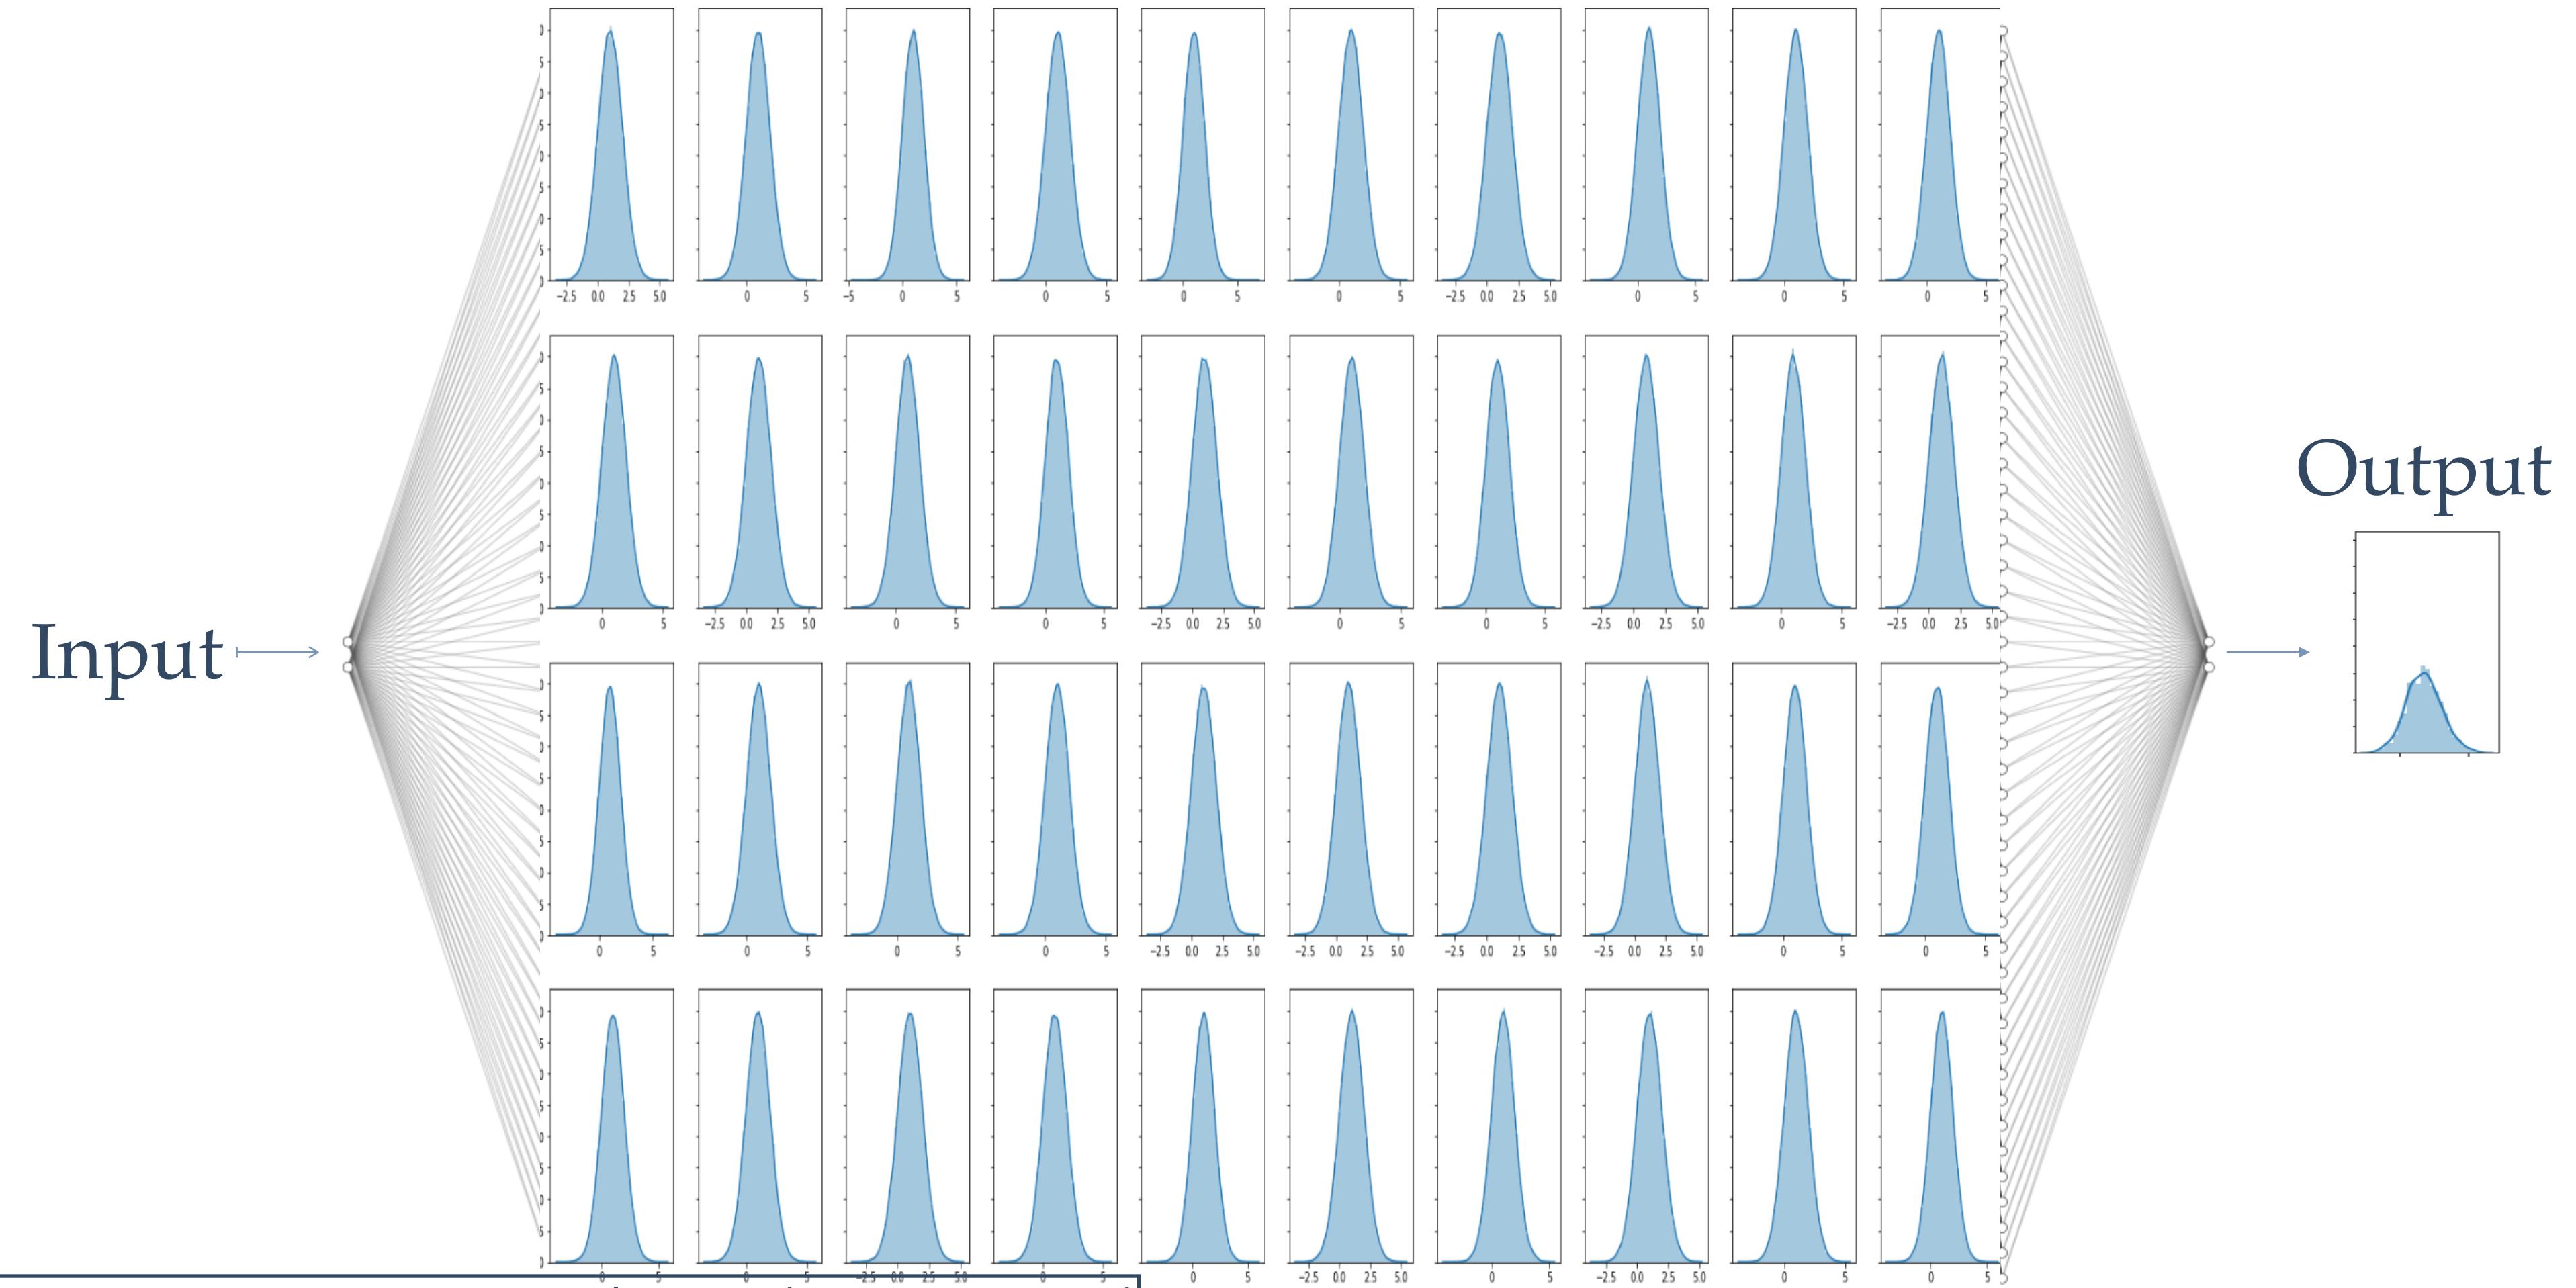
$$\omega = \mu + \sigma \odot (\pm \epsilon)$$

$q(\epsilon)$   
sample + flip de coin



# Bayesian Neural Network: MFVB

$$p(\theta | y) \approx q^* = \operatorname{argmax}_{q \in Q} \text{ELBO}$$



for  $j \in \{1, \dots, m\}$  do

until ELBO has converged

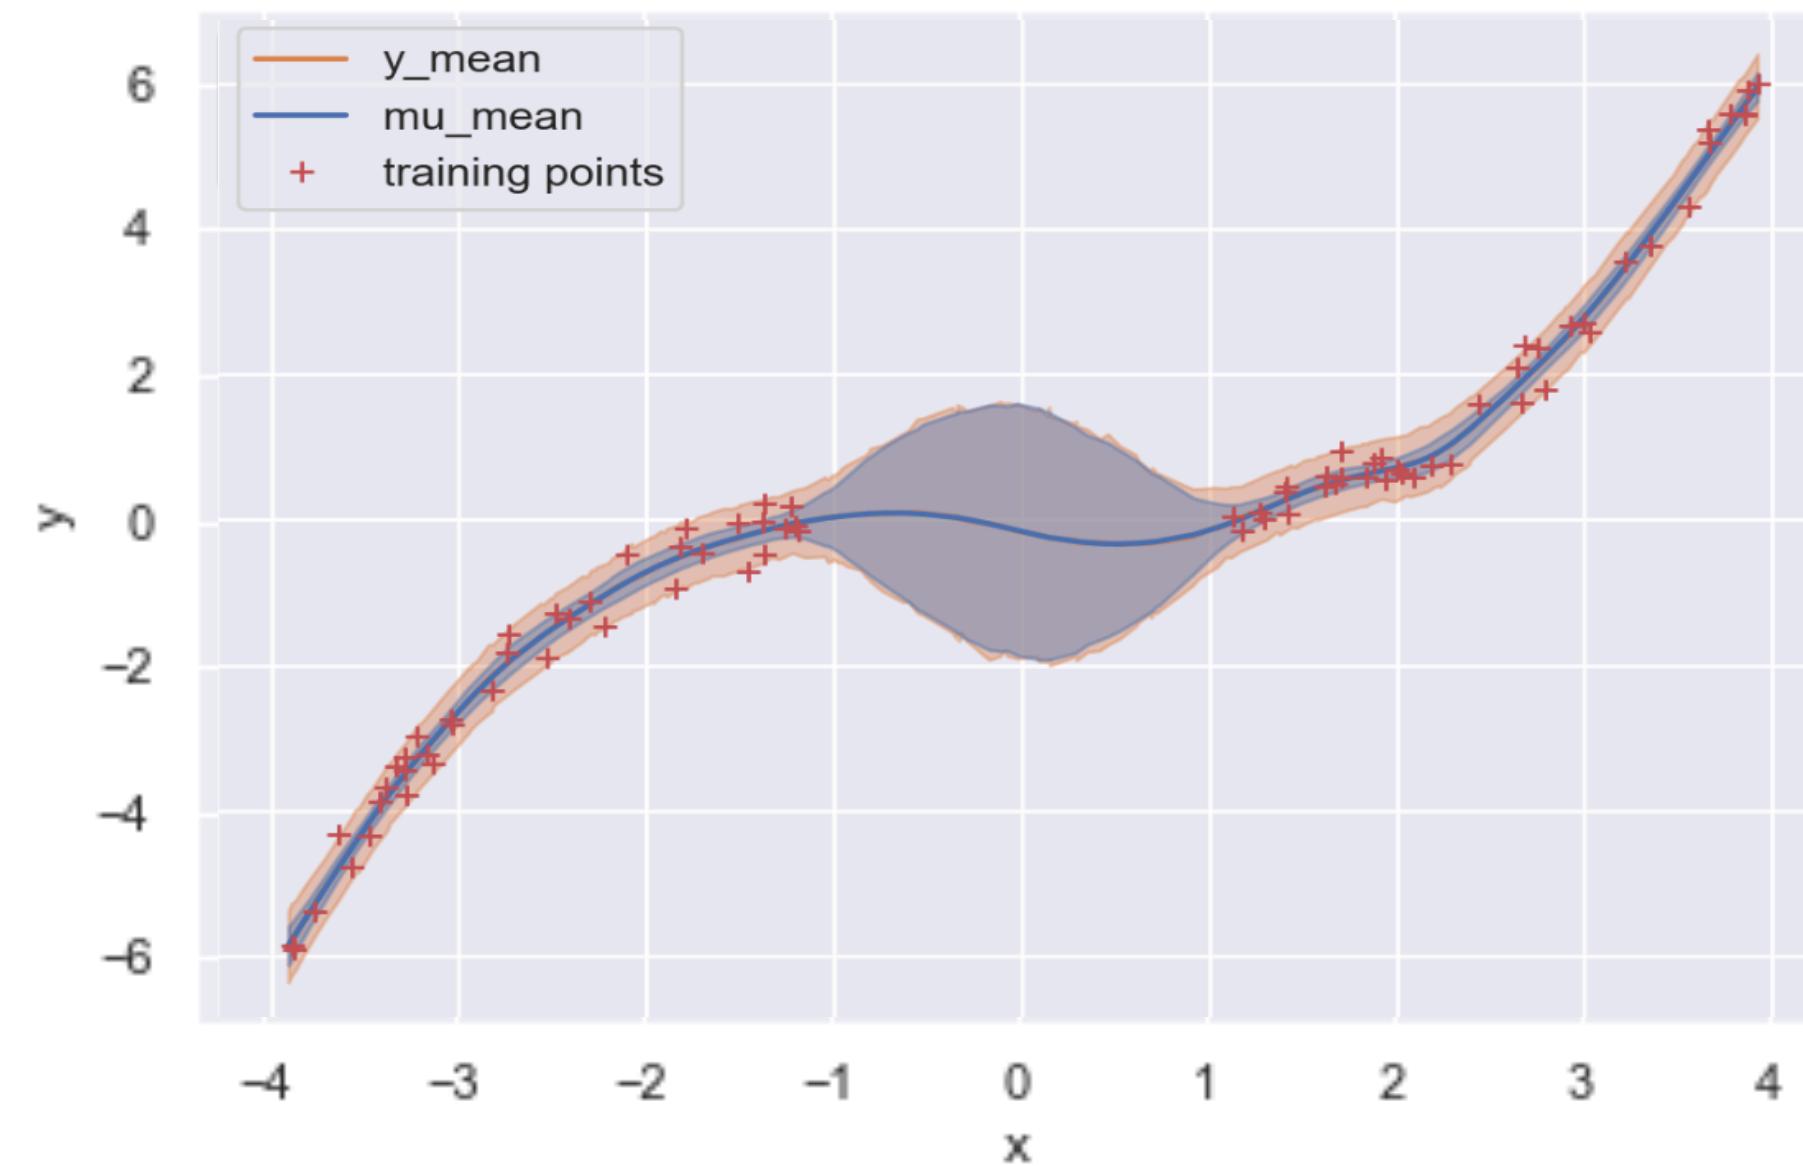
$$q_j(\theta_j) \propto \exp\{\mathbb{E}_{i \neq j}[\log(p(\theta_j | \theta_{i \neq j}, y))]\}$$

compute ELBO(q)

# Variational Bayesian Inference: The problem

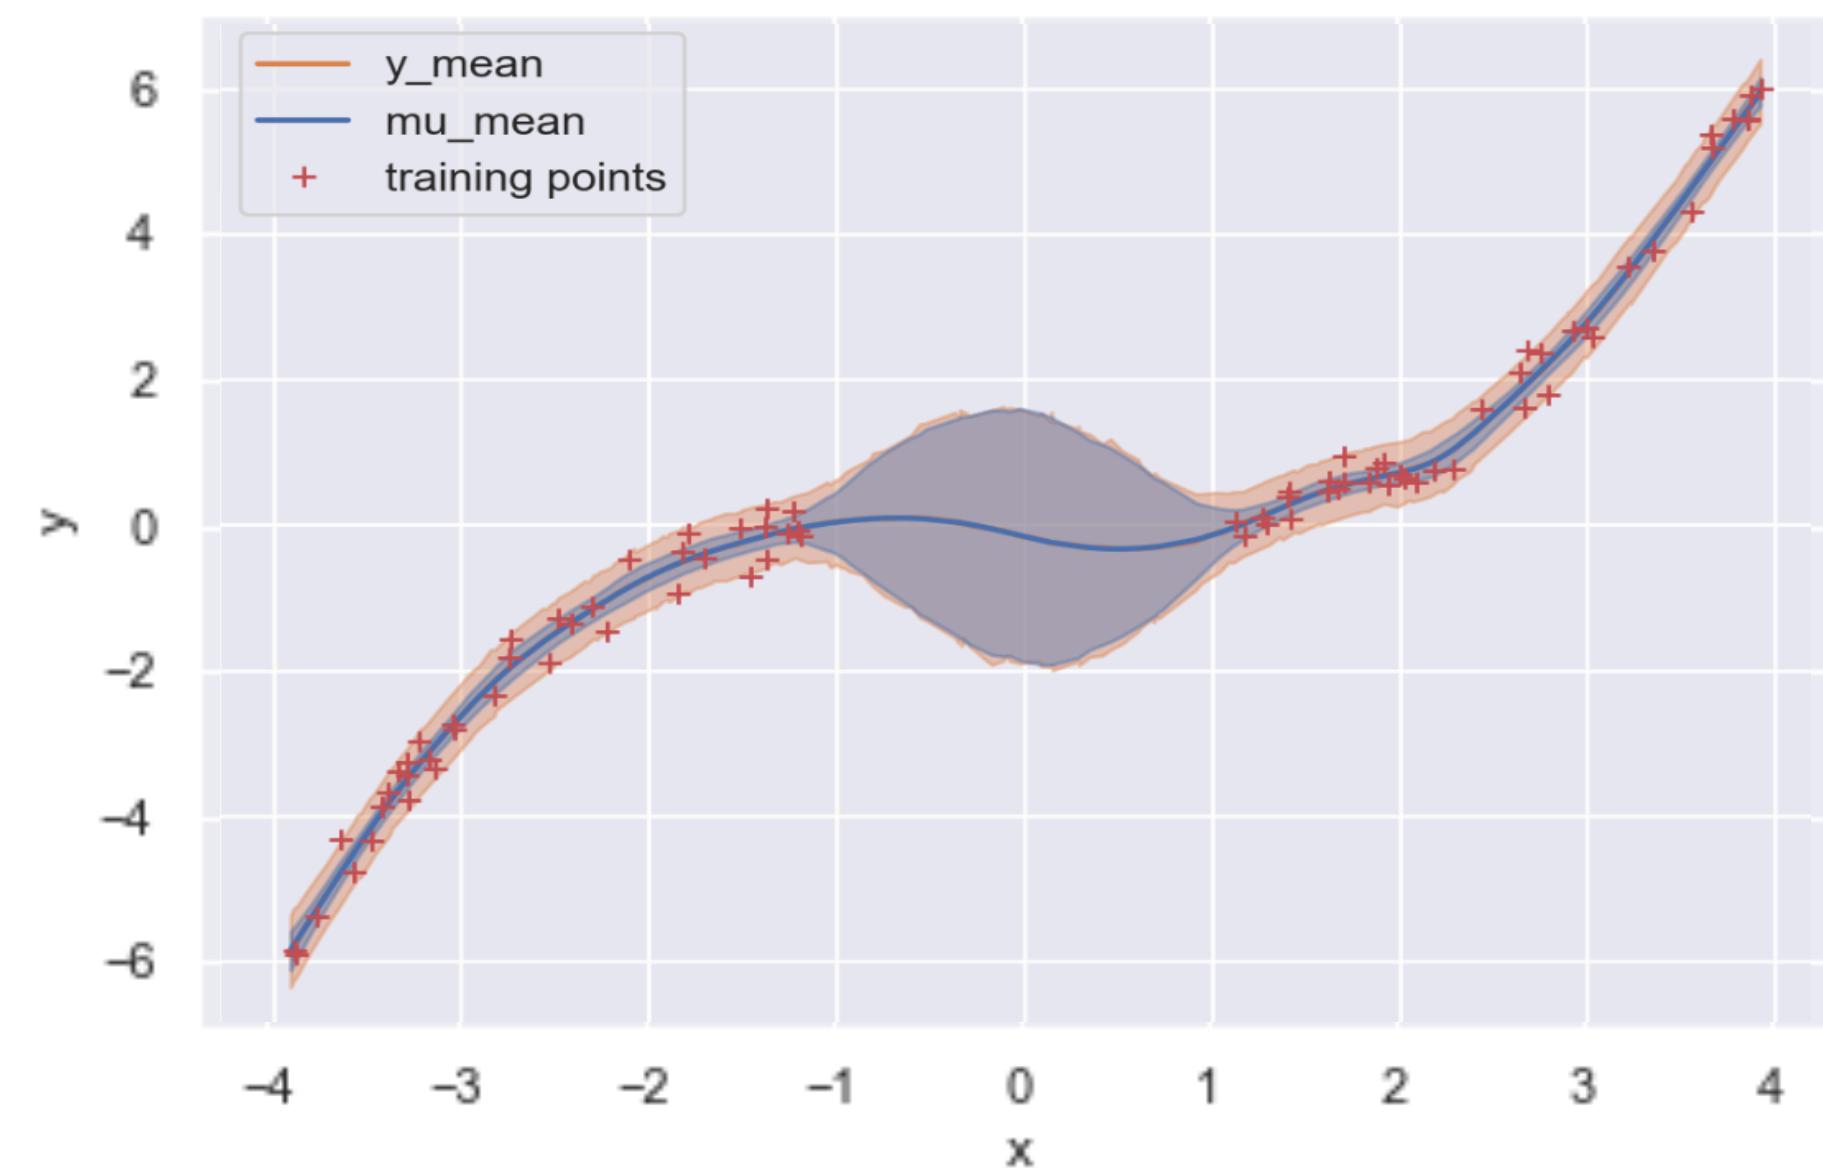


# Variational Bayesian Inference: The right solution (MCMC)

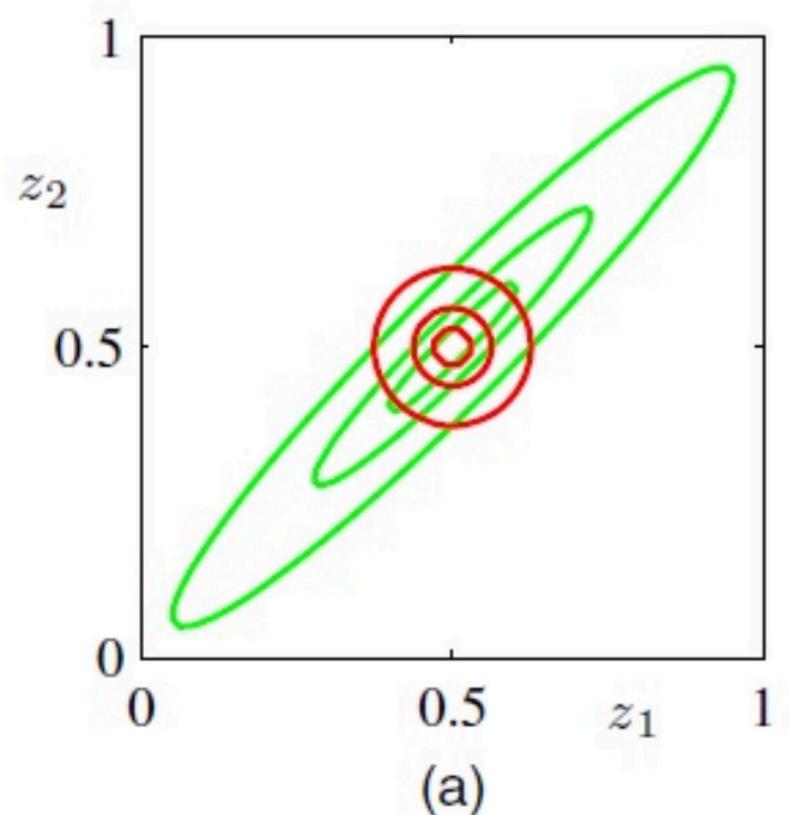
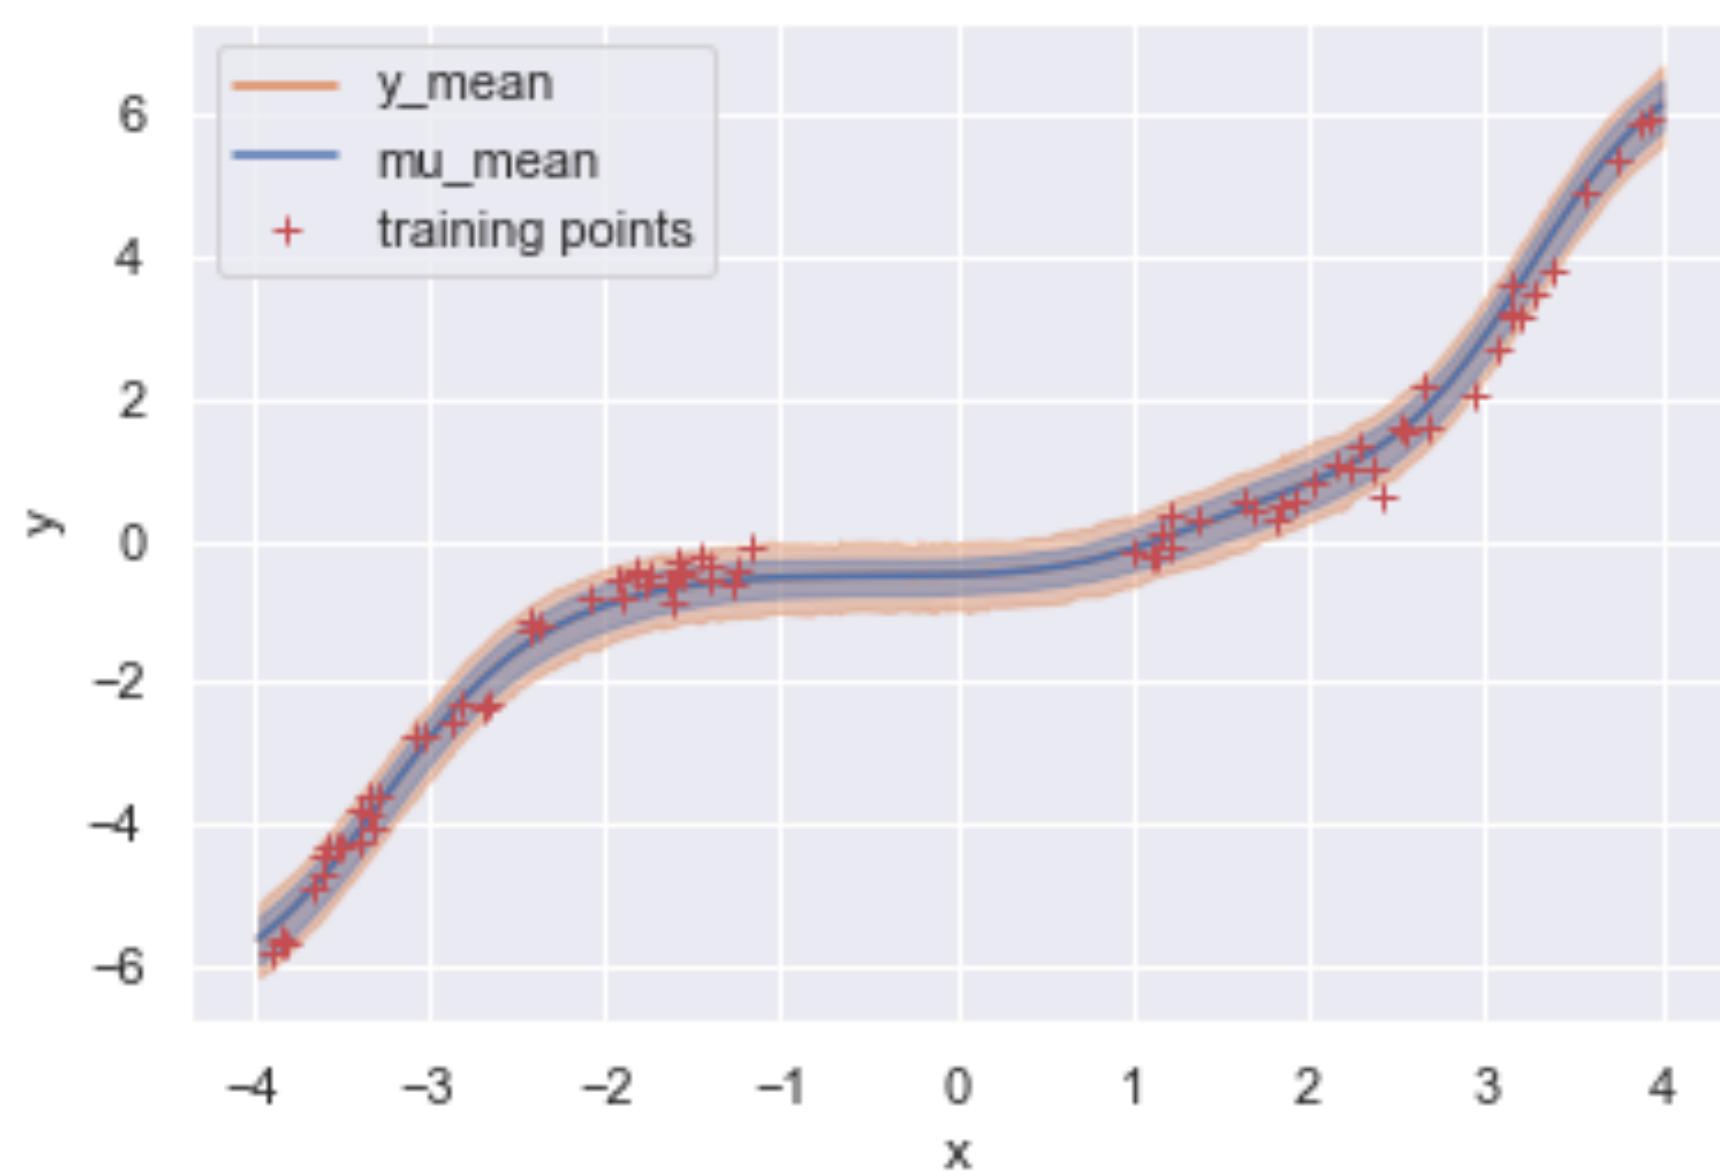


# Variational Bayesian Inference

MCMC



MFVB



# Dropout

---

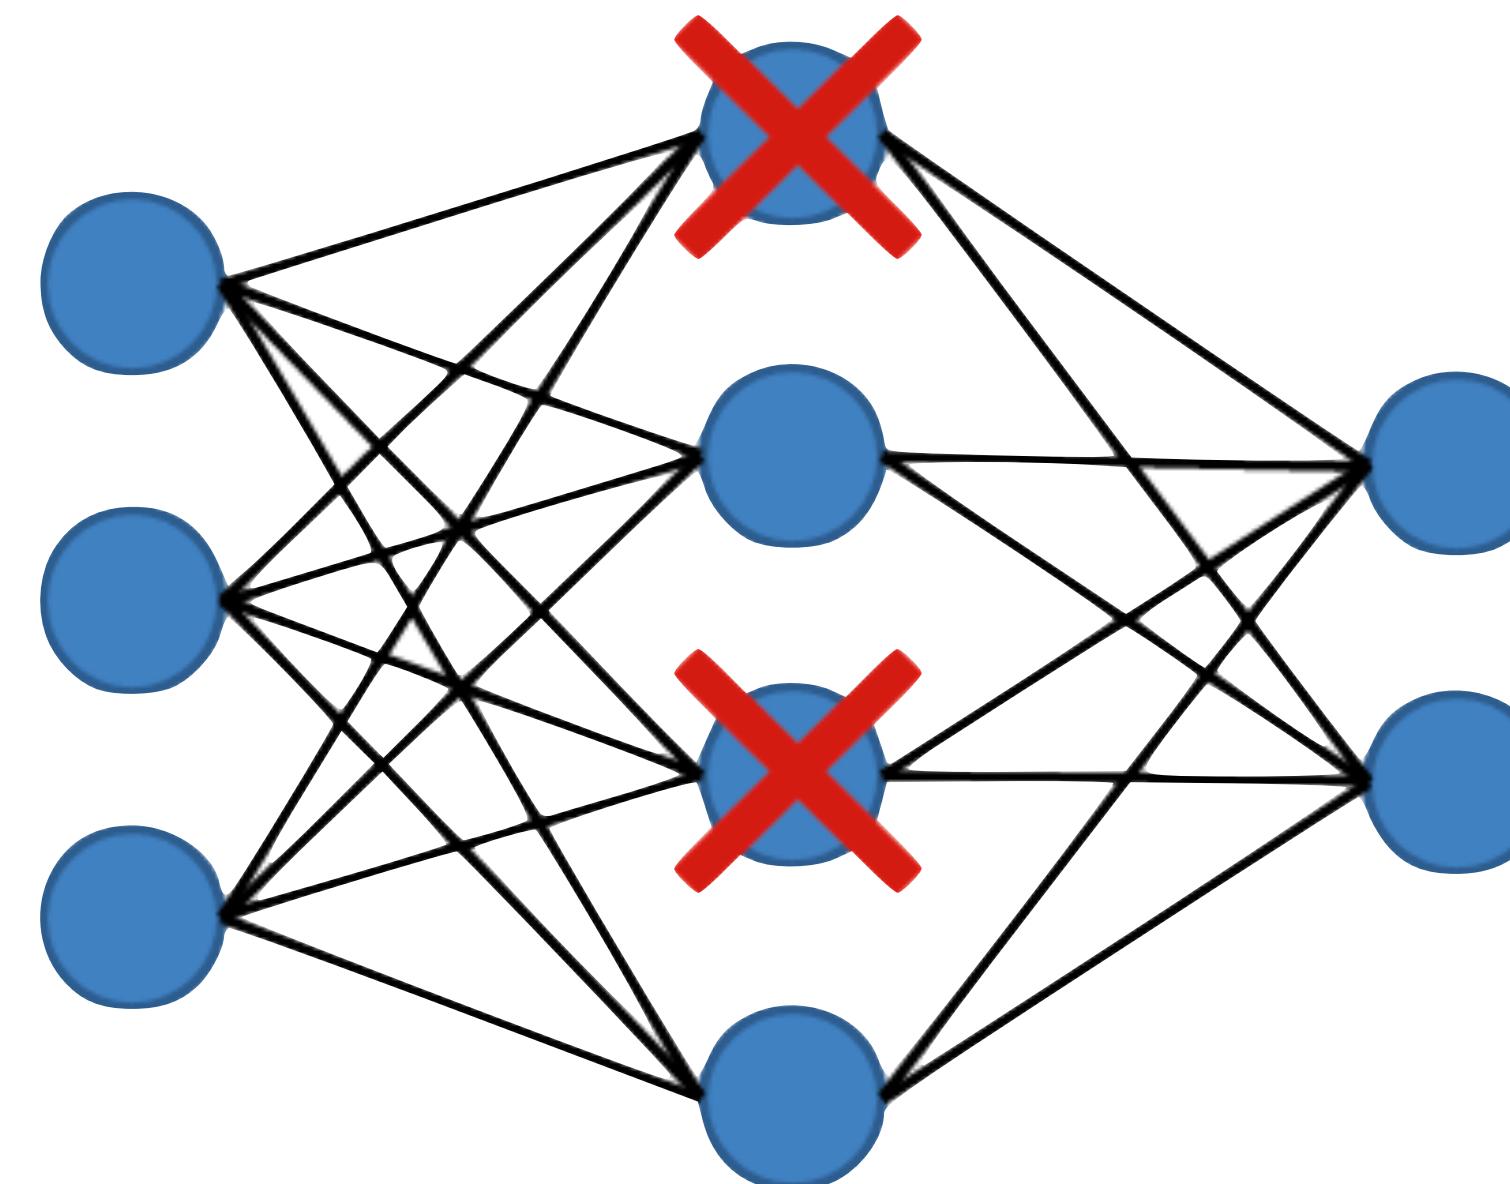
**Dropout as a Bayesian Approximation:  
Representing Model Uncertainty in Deep Learning**

---

[arXiv:1506.02142](https://arxiv.org/abs/1506.02142)

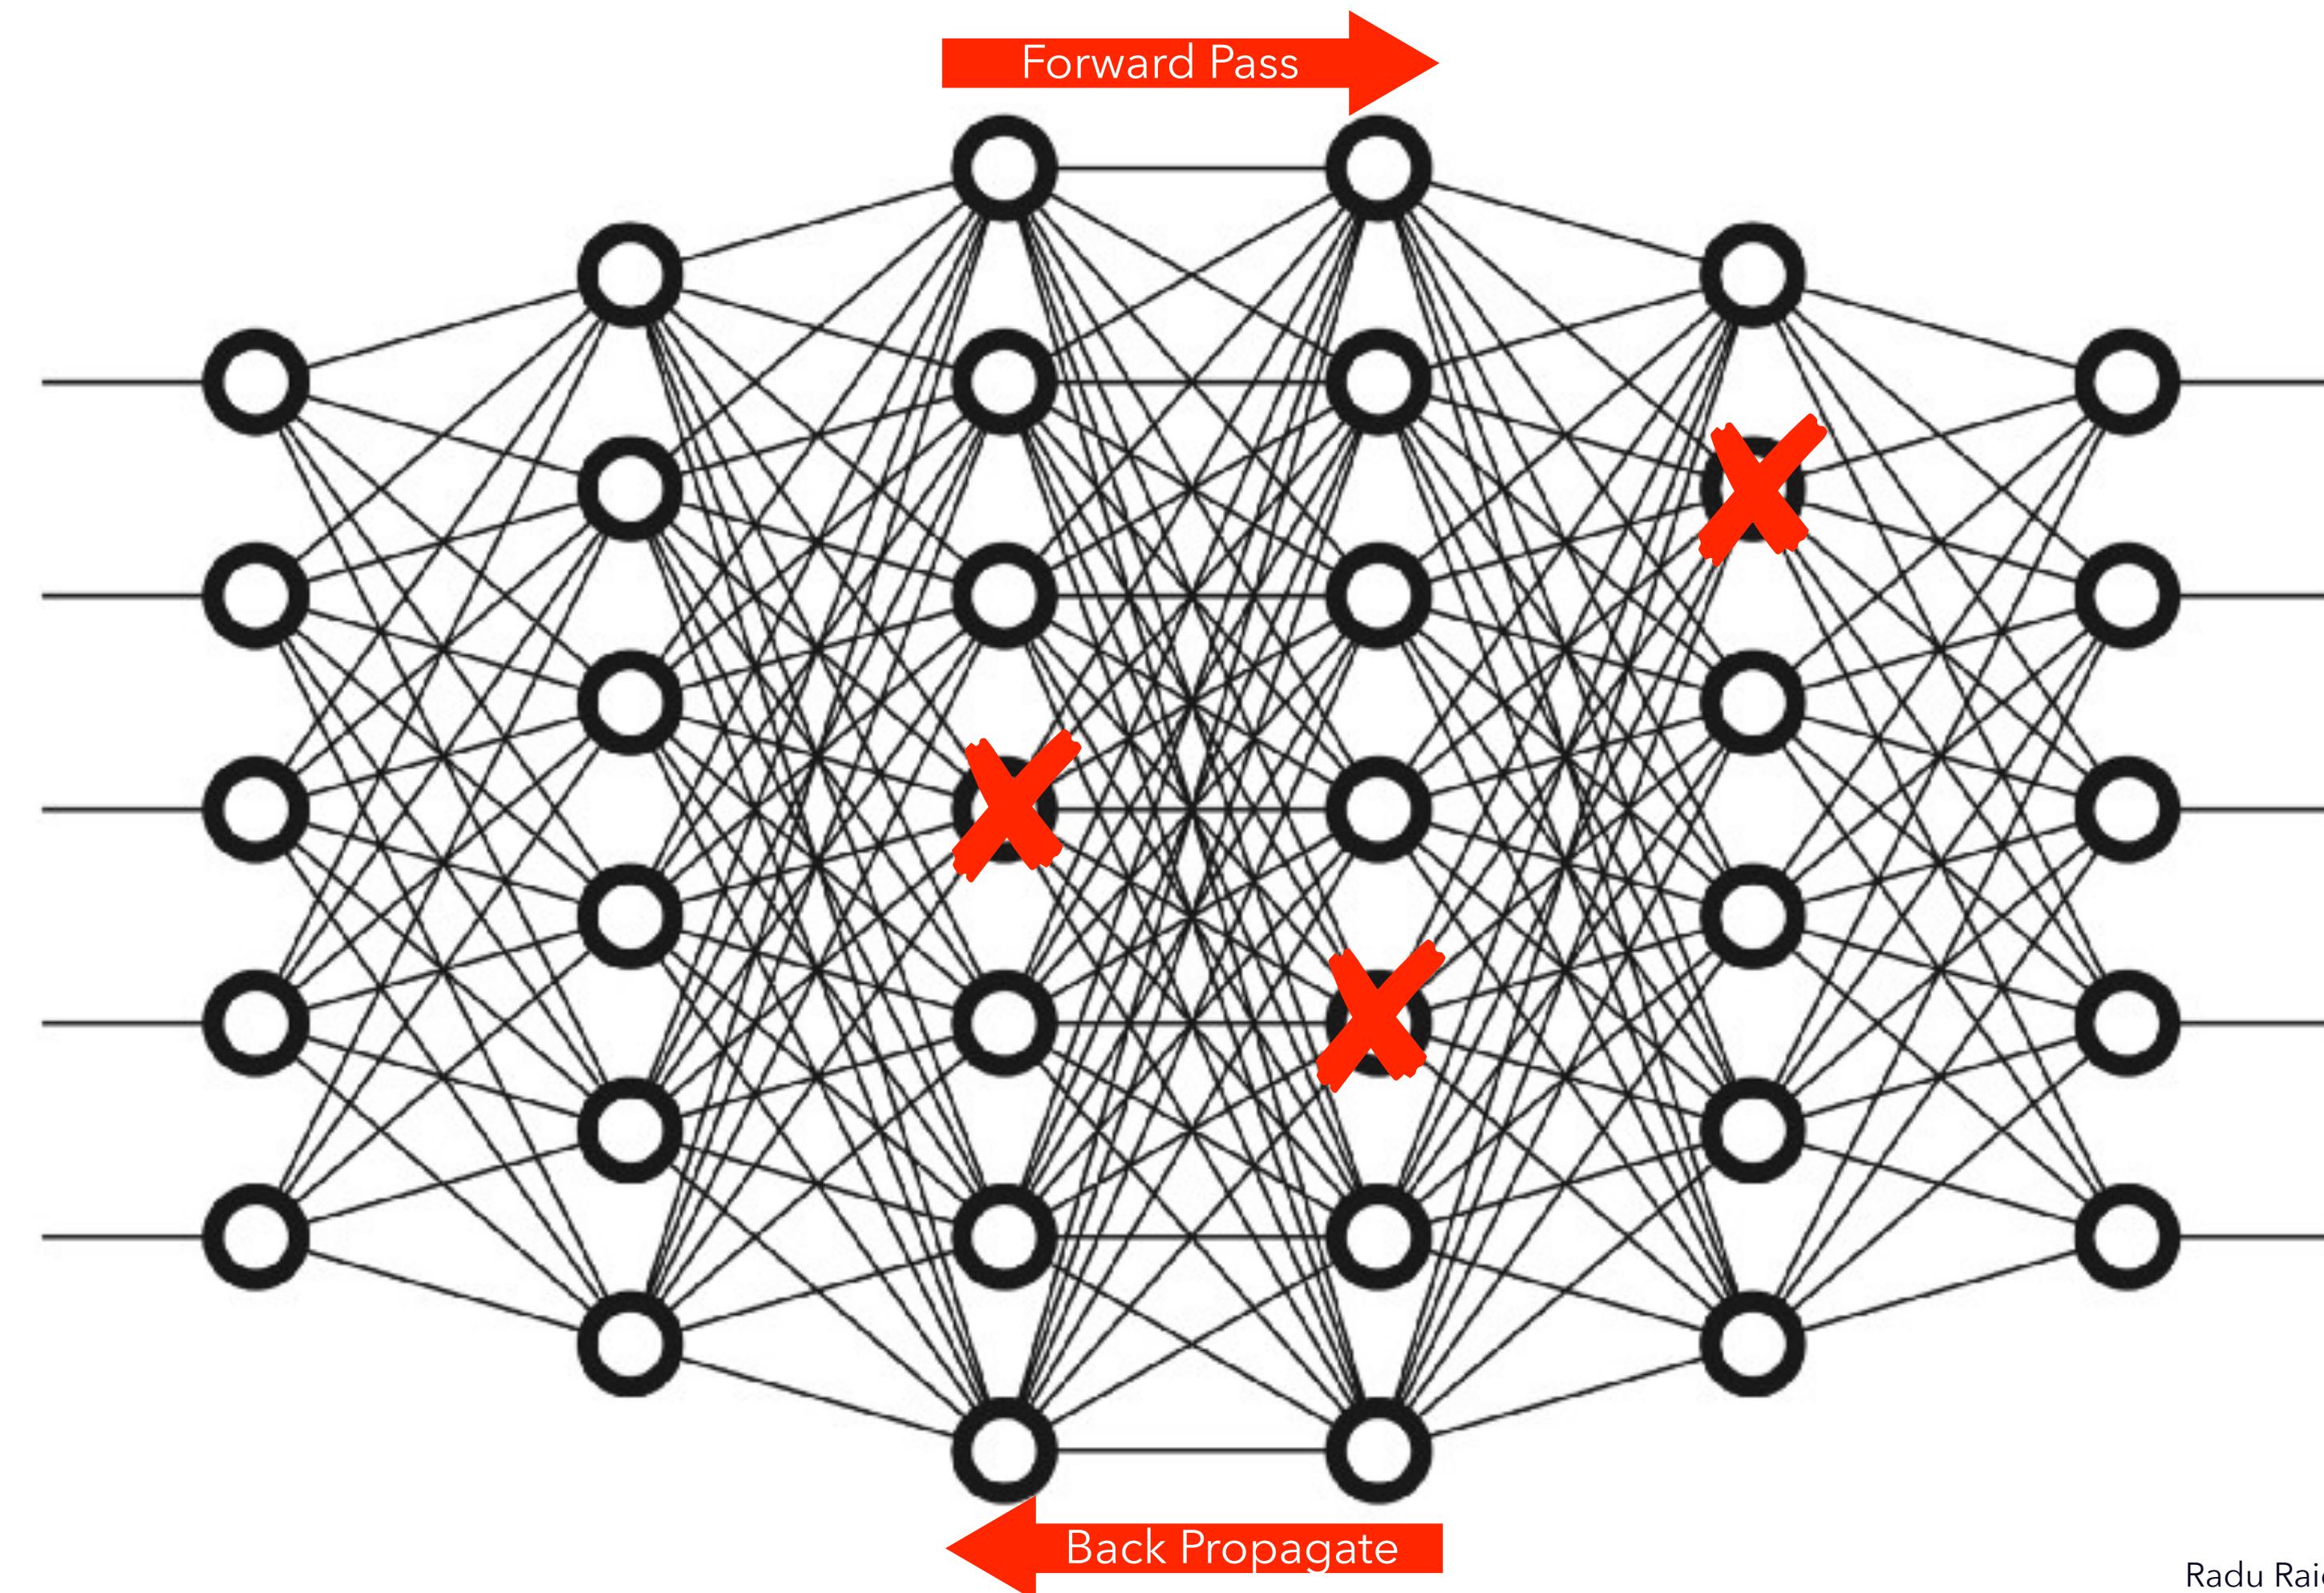
**Yarin Gal**  
**Zoubin Ghahramani**  
University of Cambridge

YG279@CAM.AC.UK  
ZG201@CAM.AC.UK



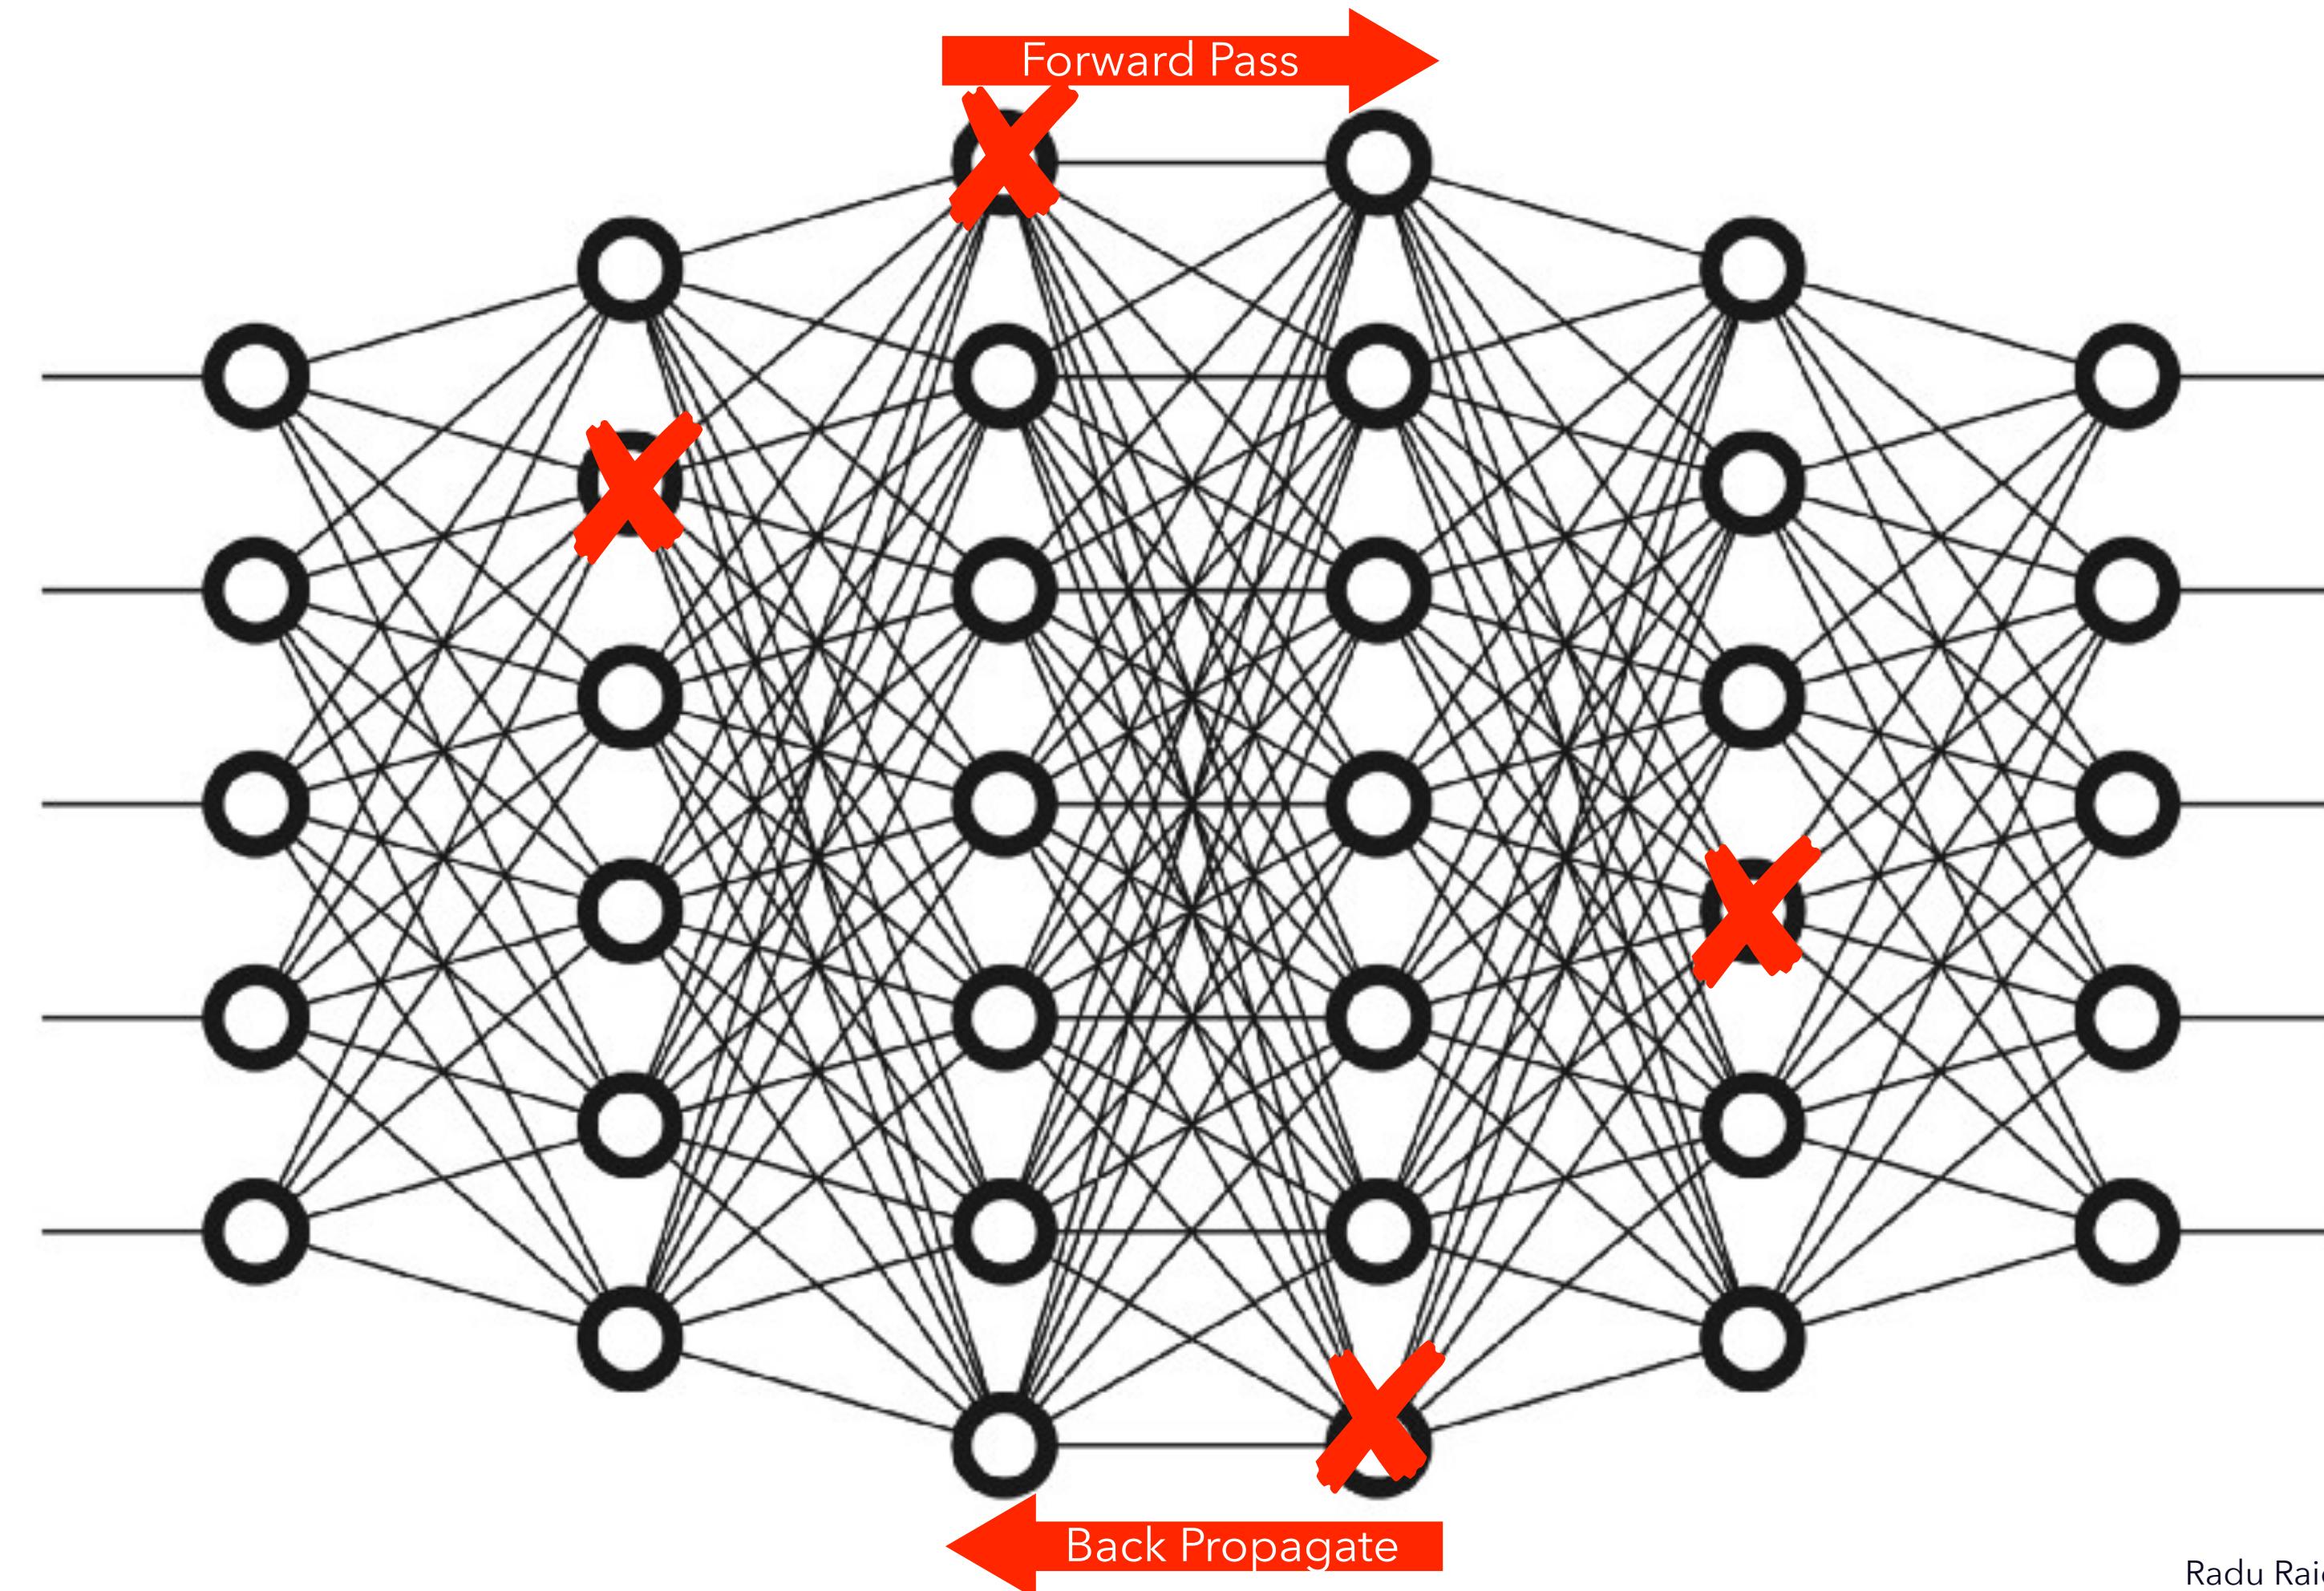
They show that a NN with arbitrary depth and non-linearities, with dropout applied before every weight layer, is mathematically equivalent to an approximation to the probabilistic deep Gaussian process.

# Dropout: train



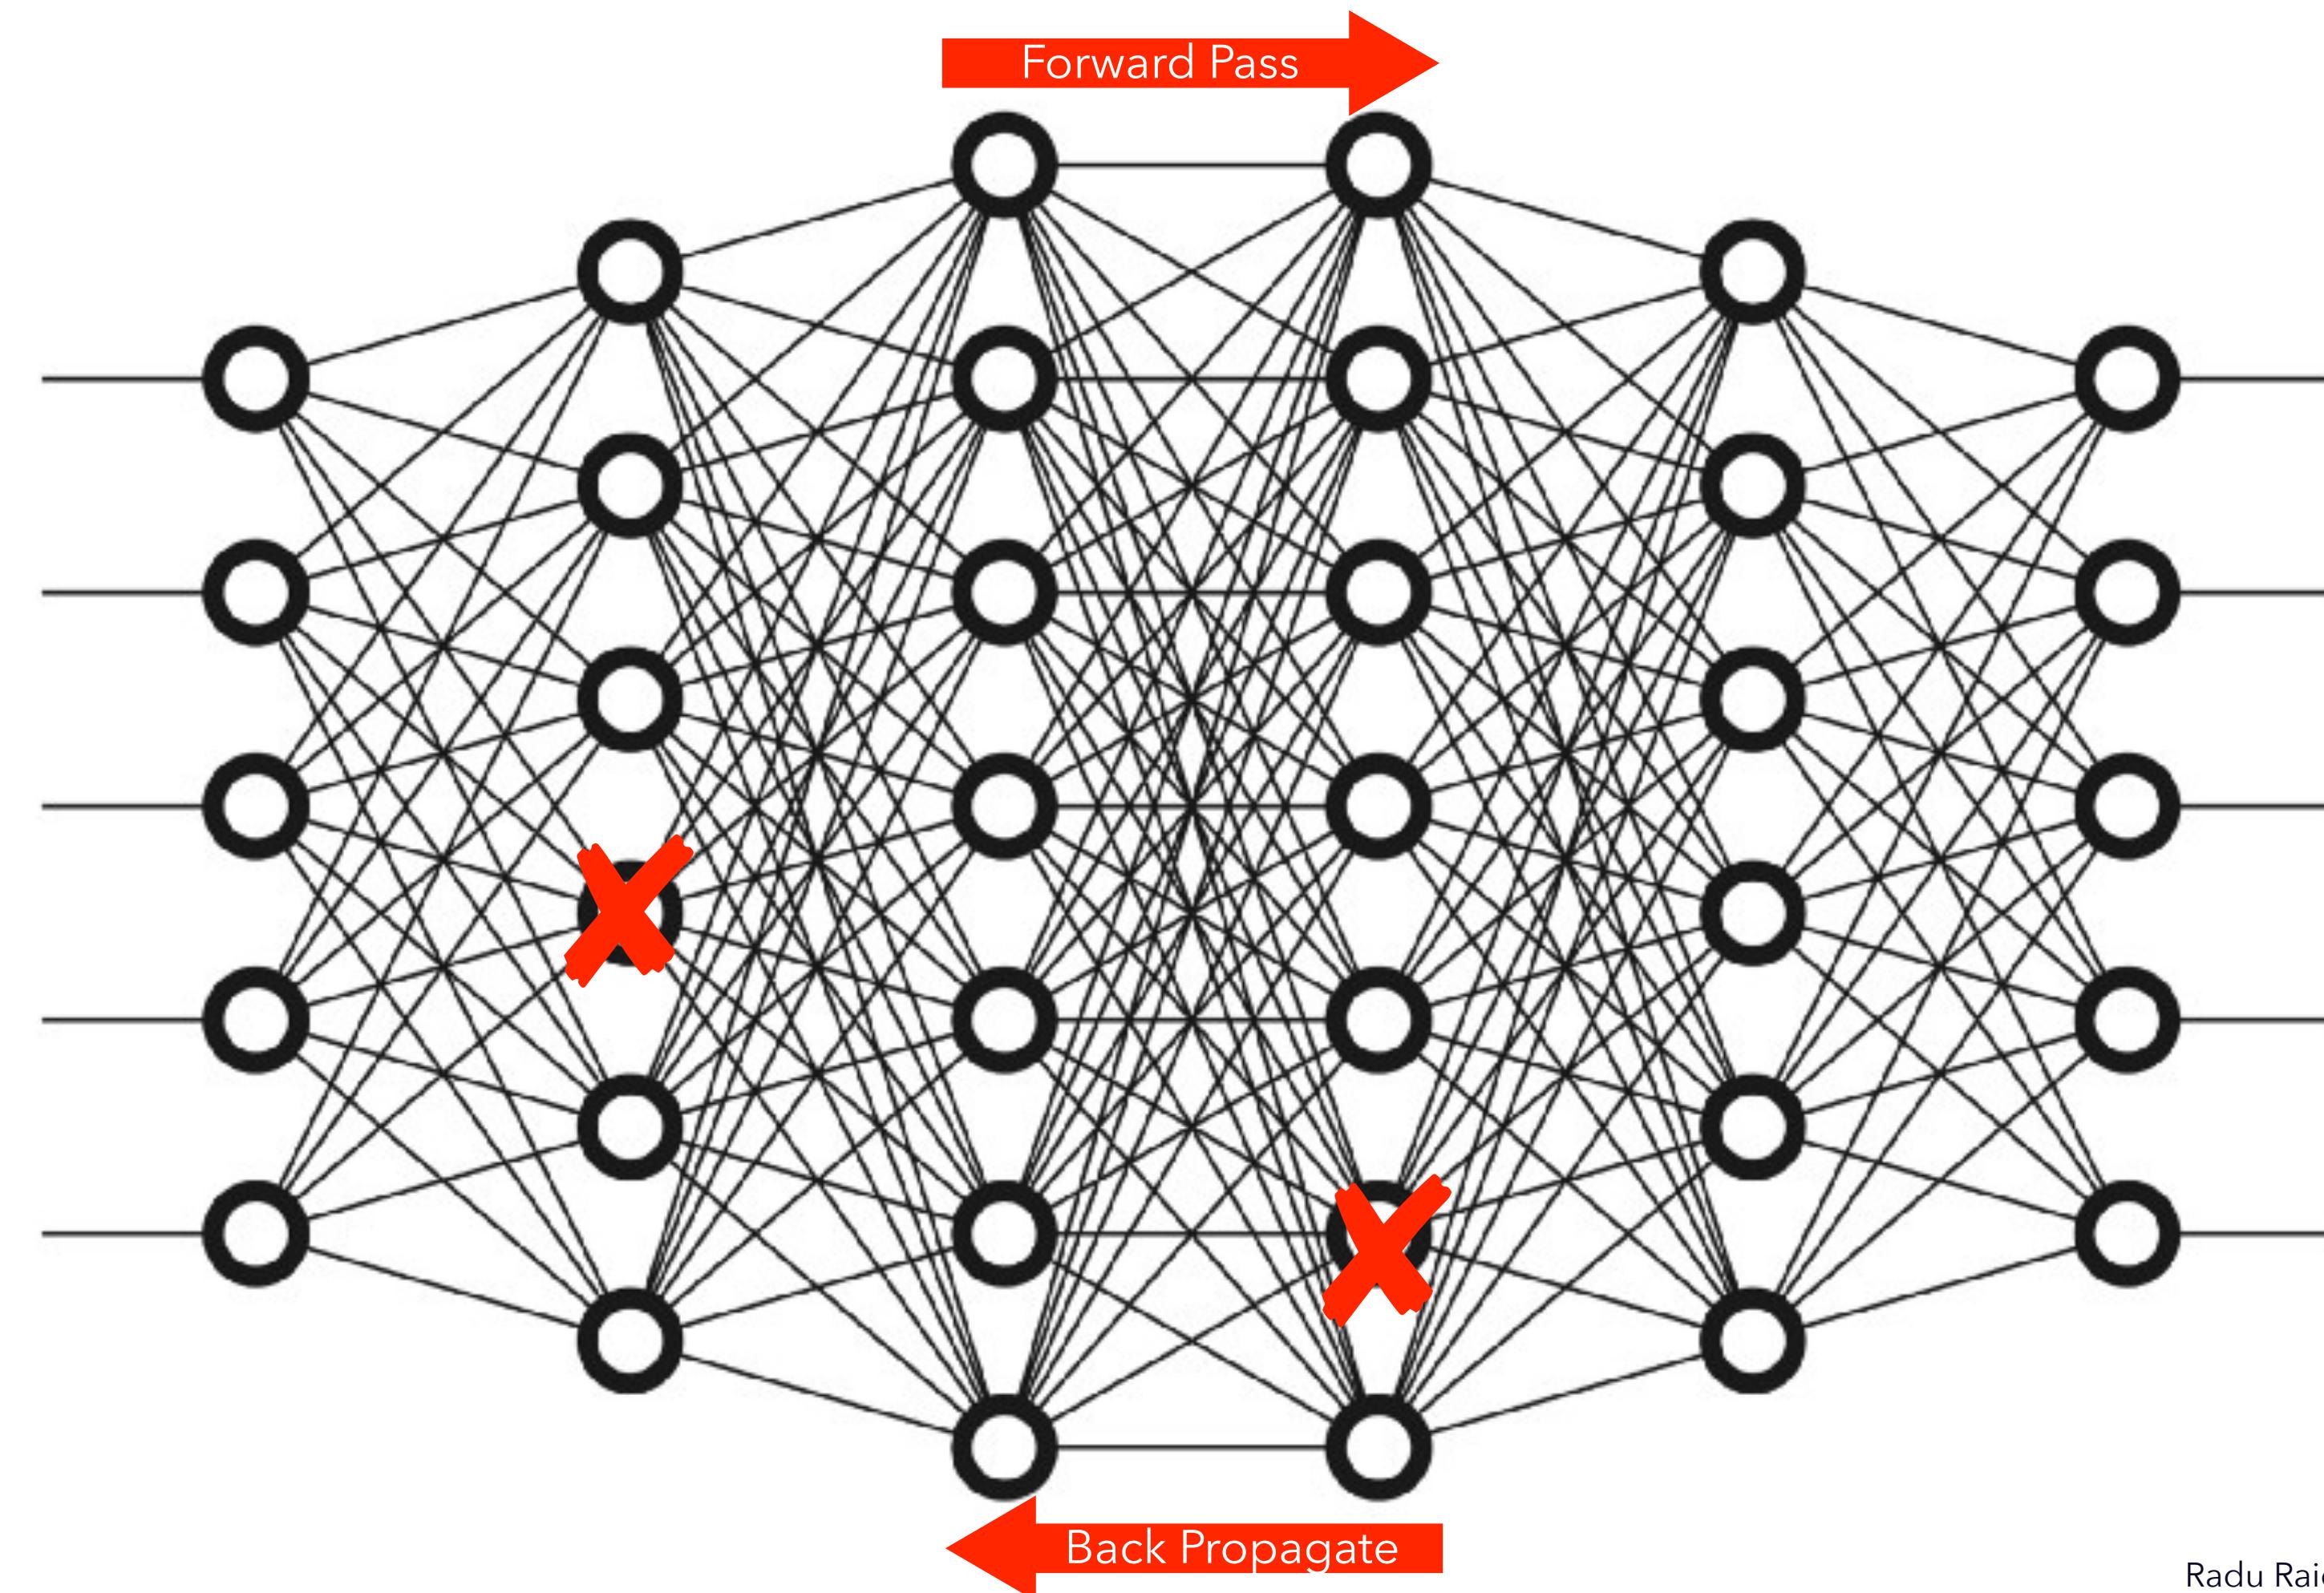
Radu Raicea

# Dropout: train



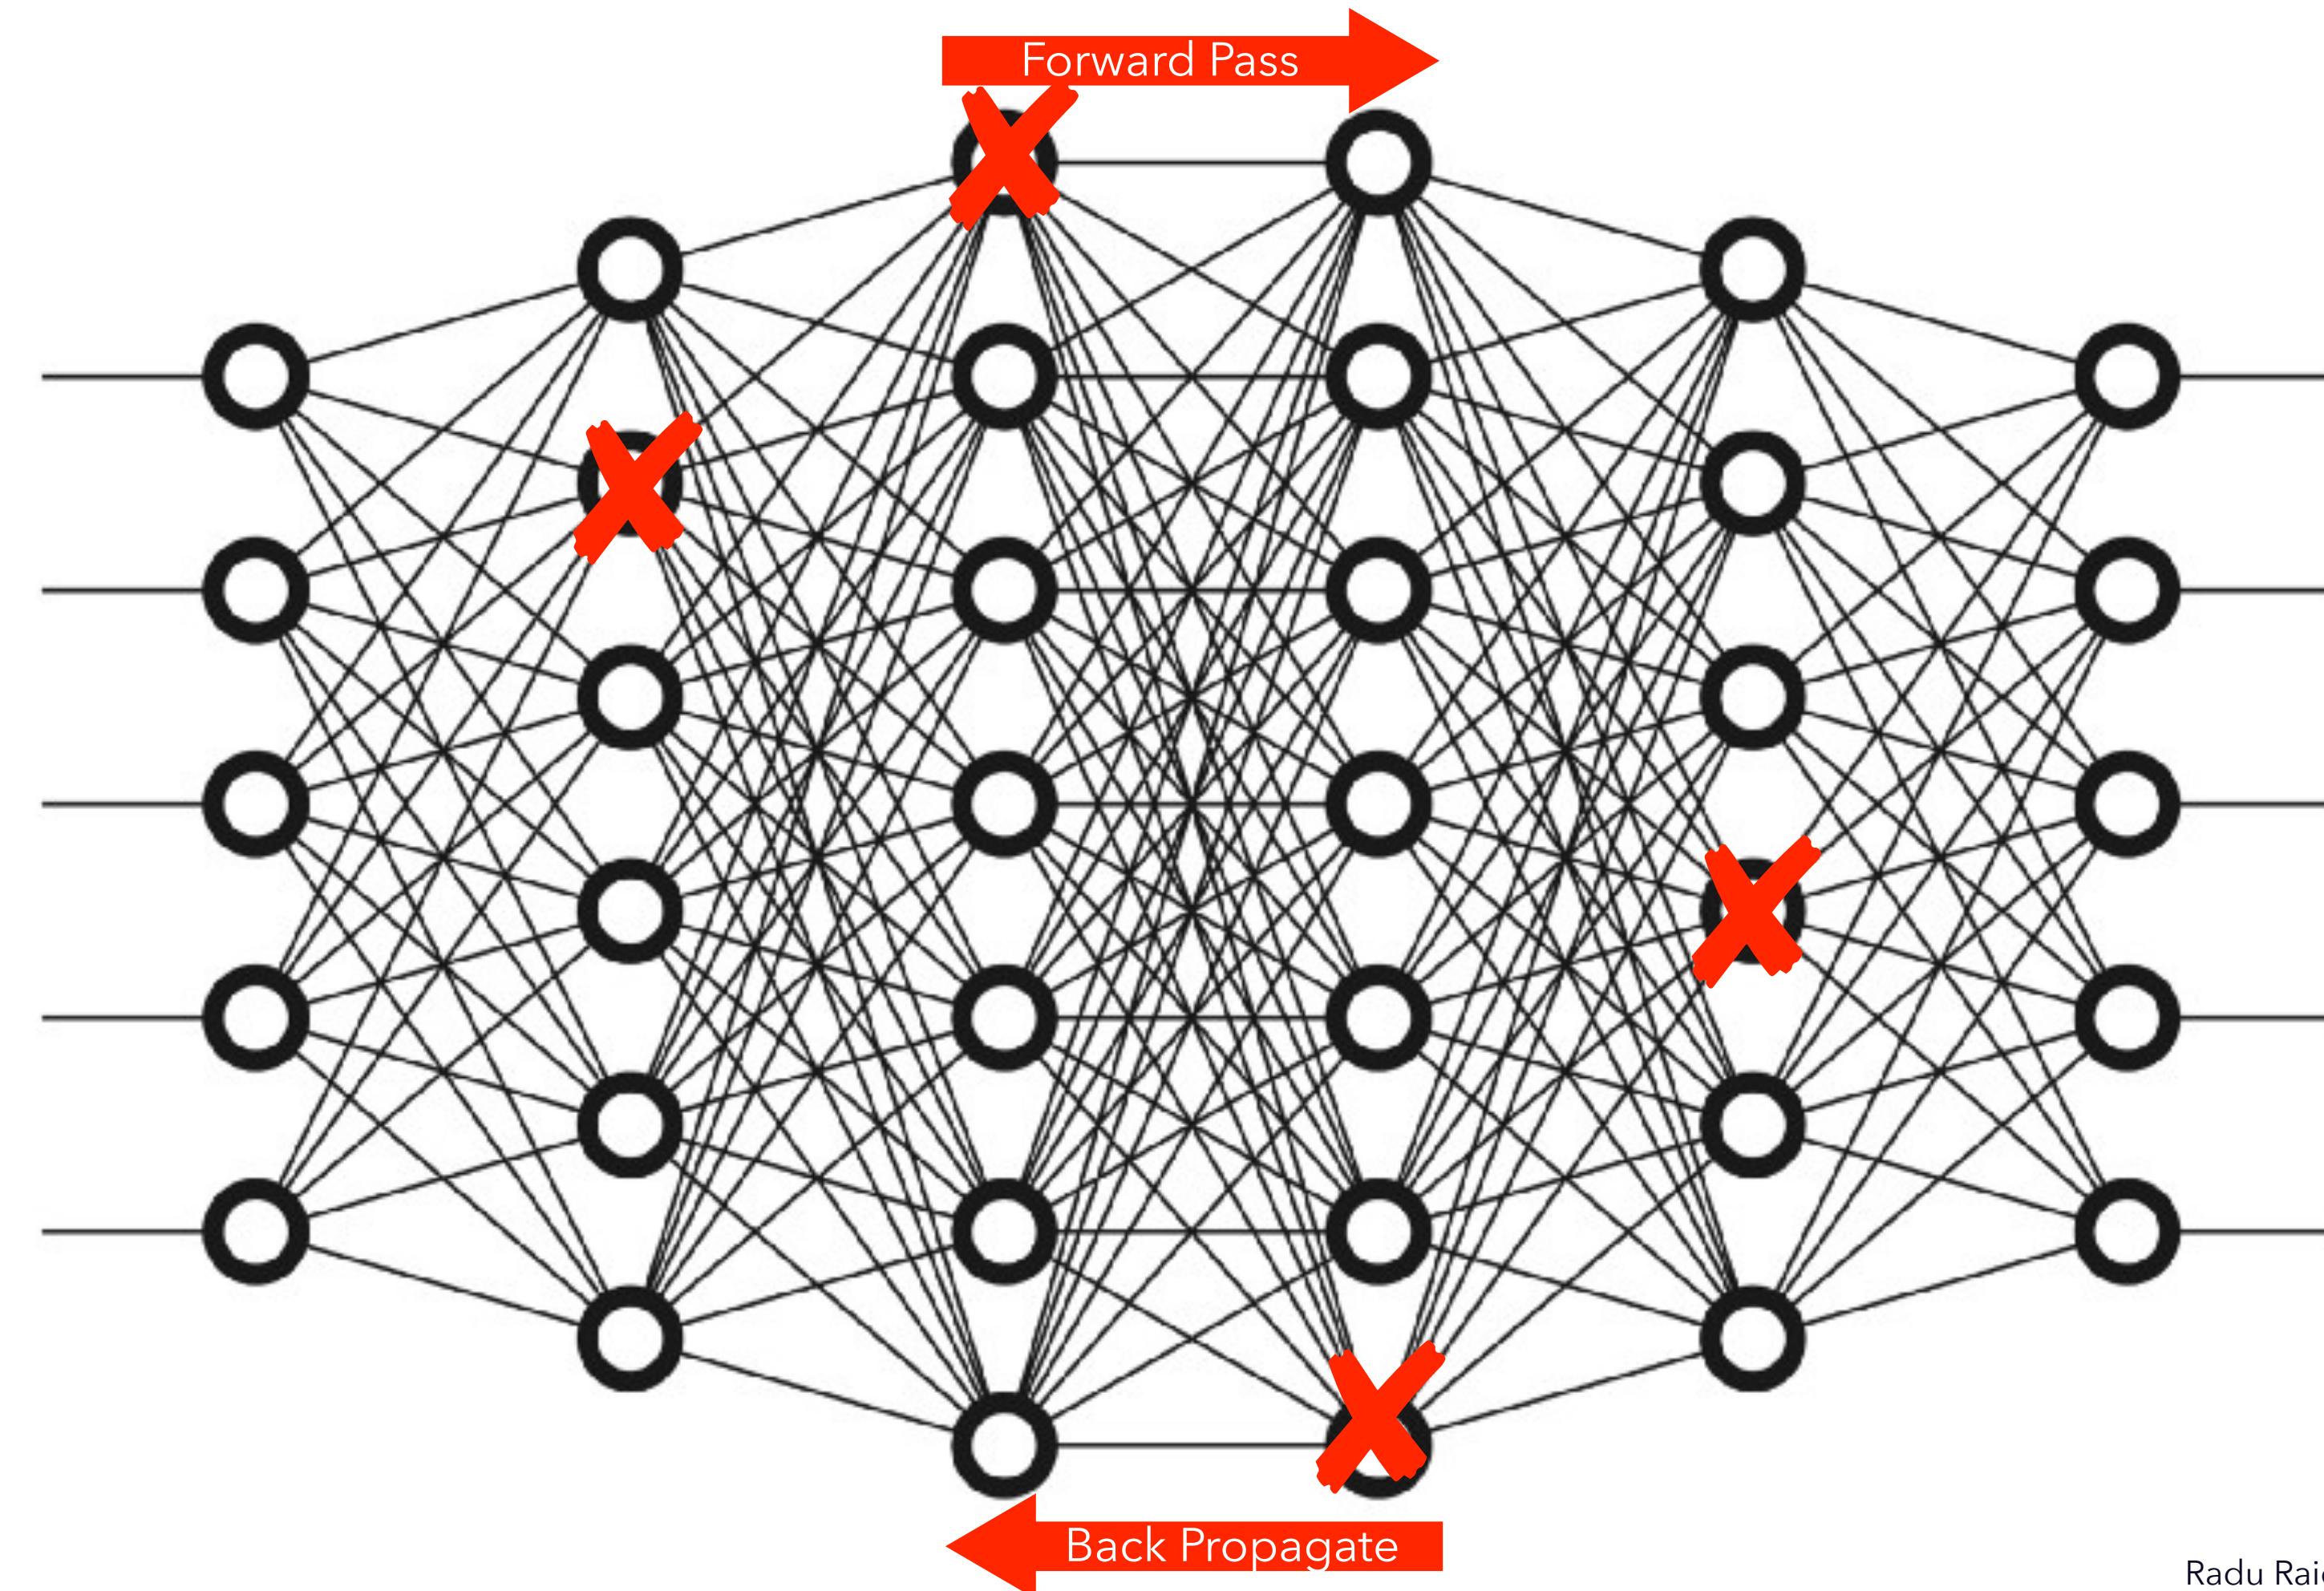
Radu Raicea

# Dropout: train



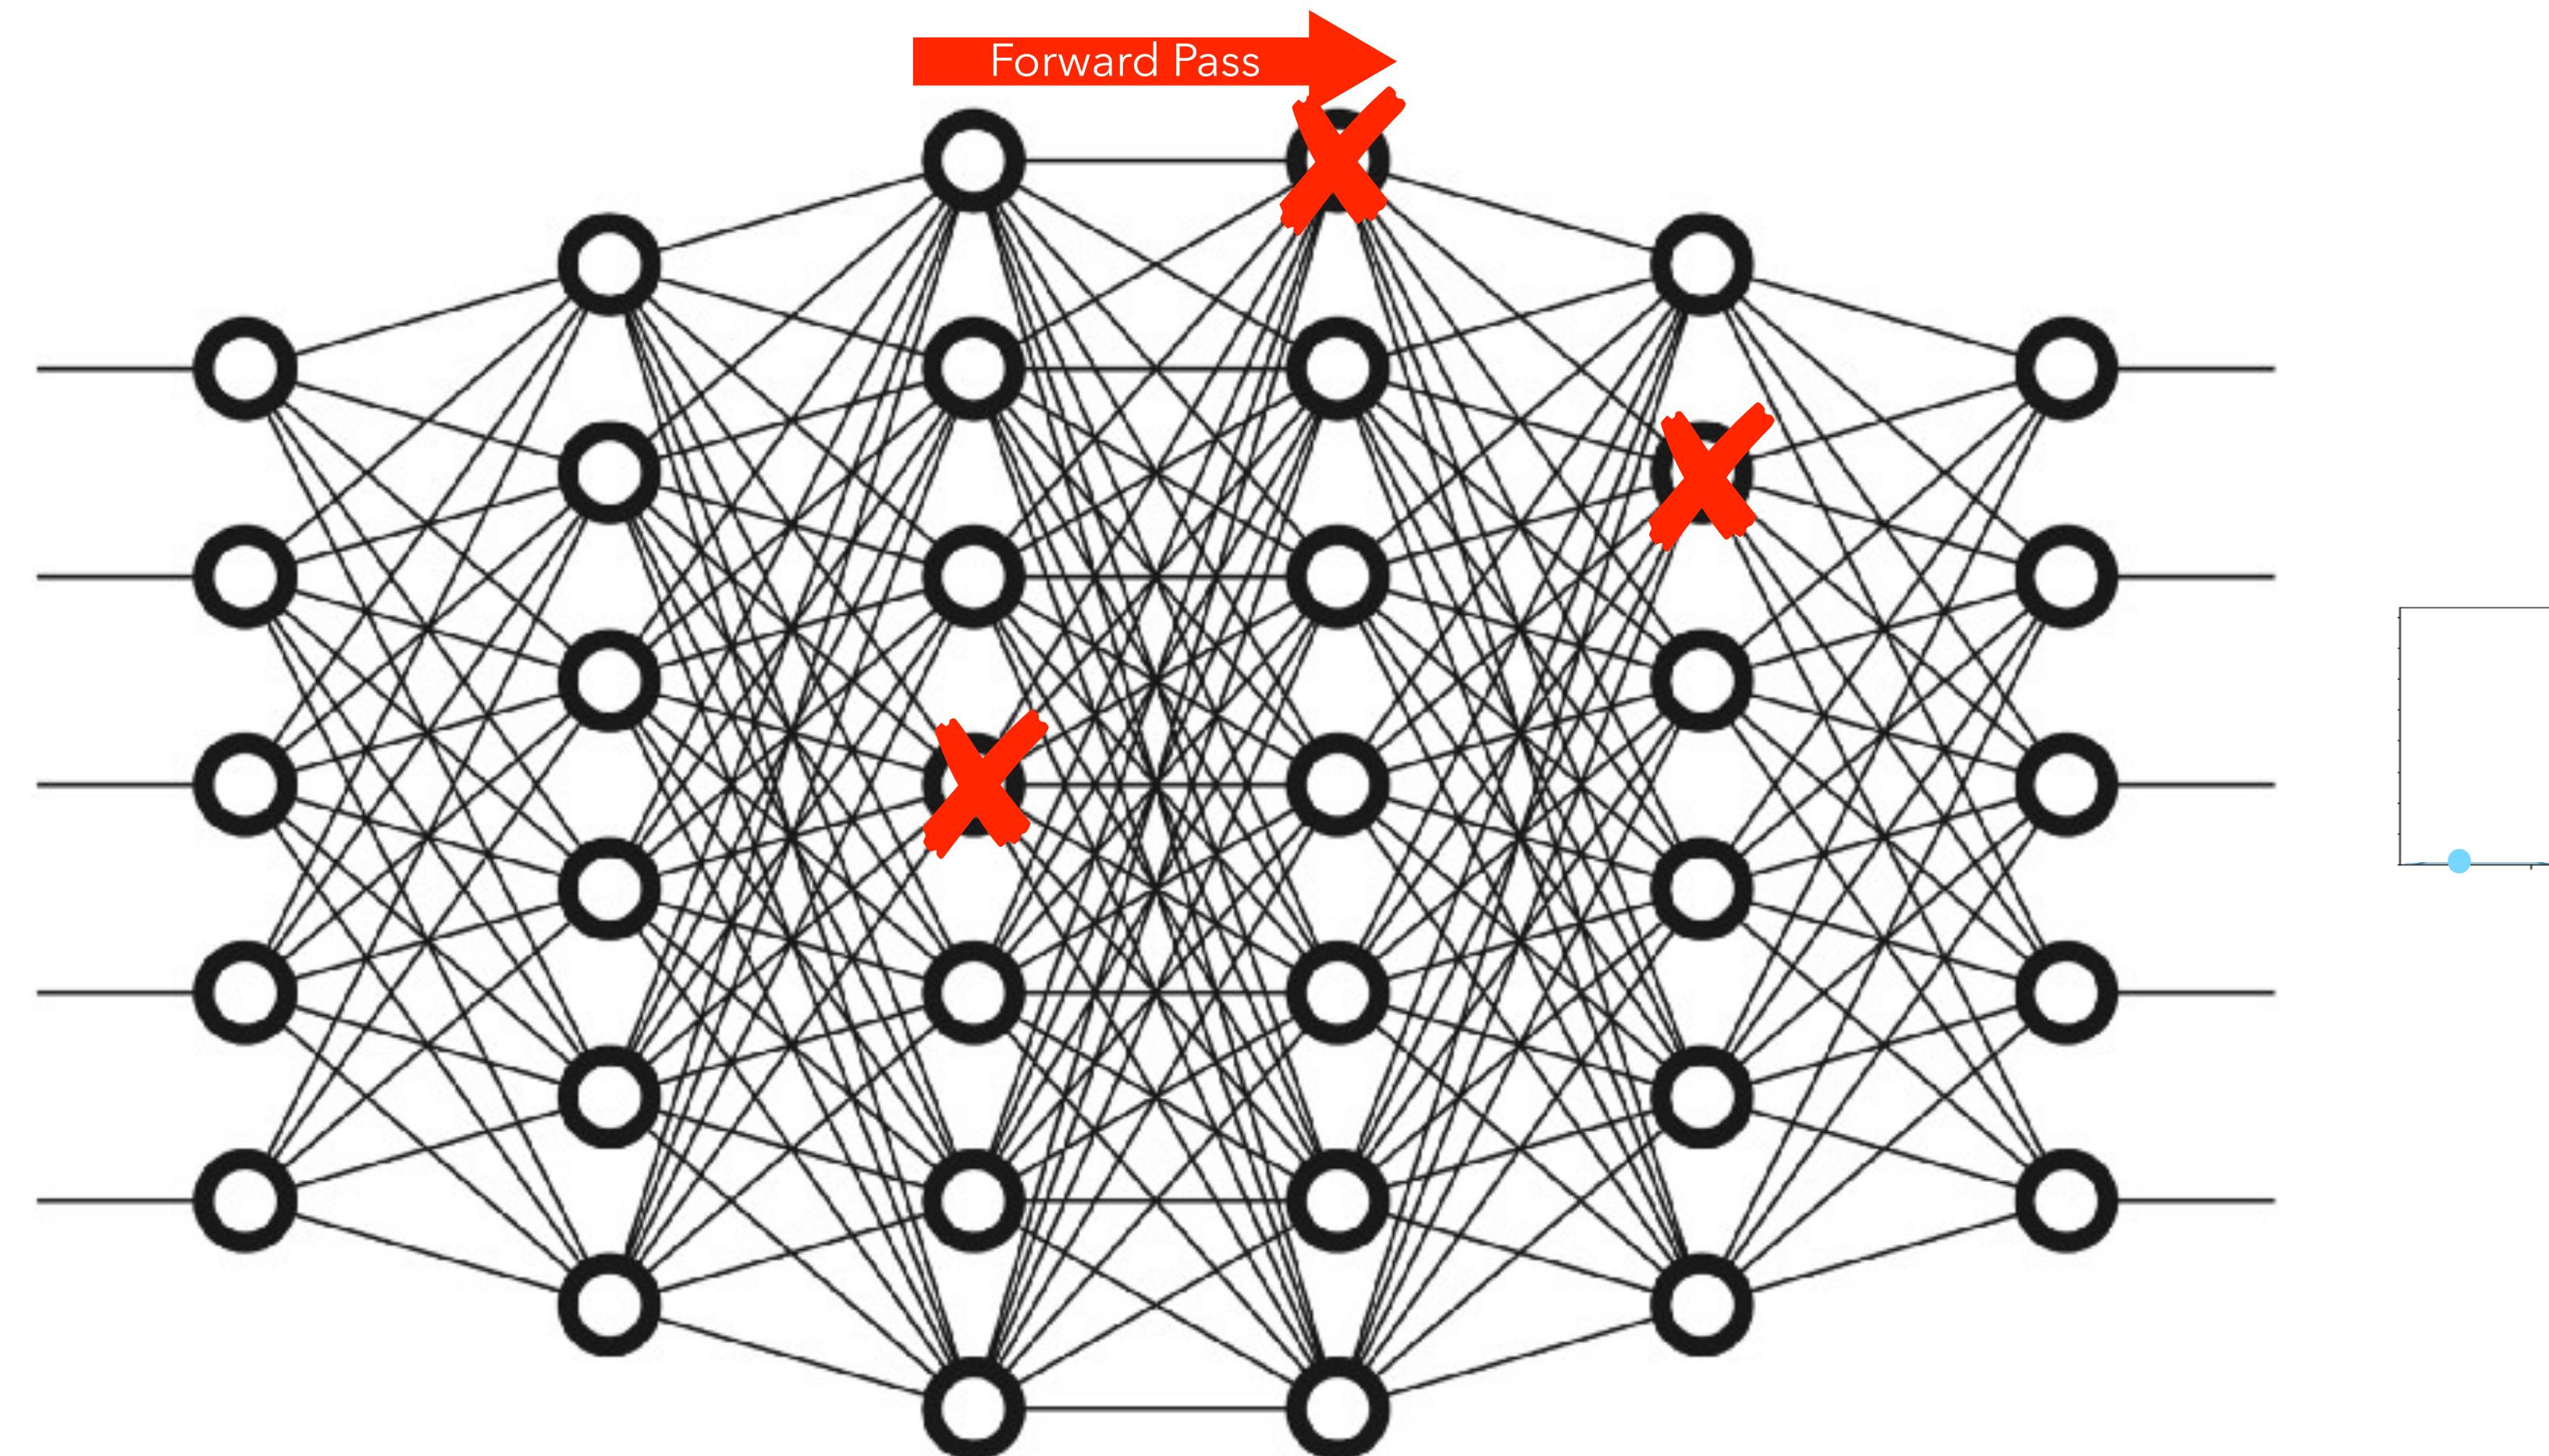
Radu Raicea

# Dropout: train



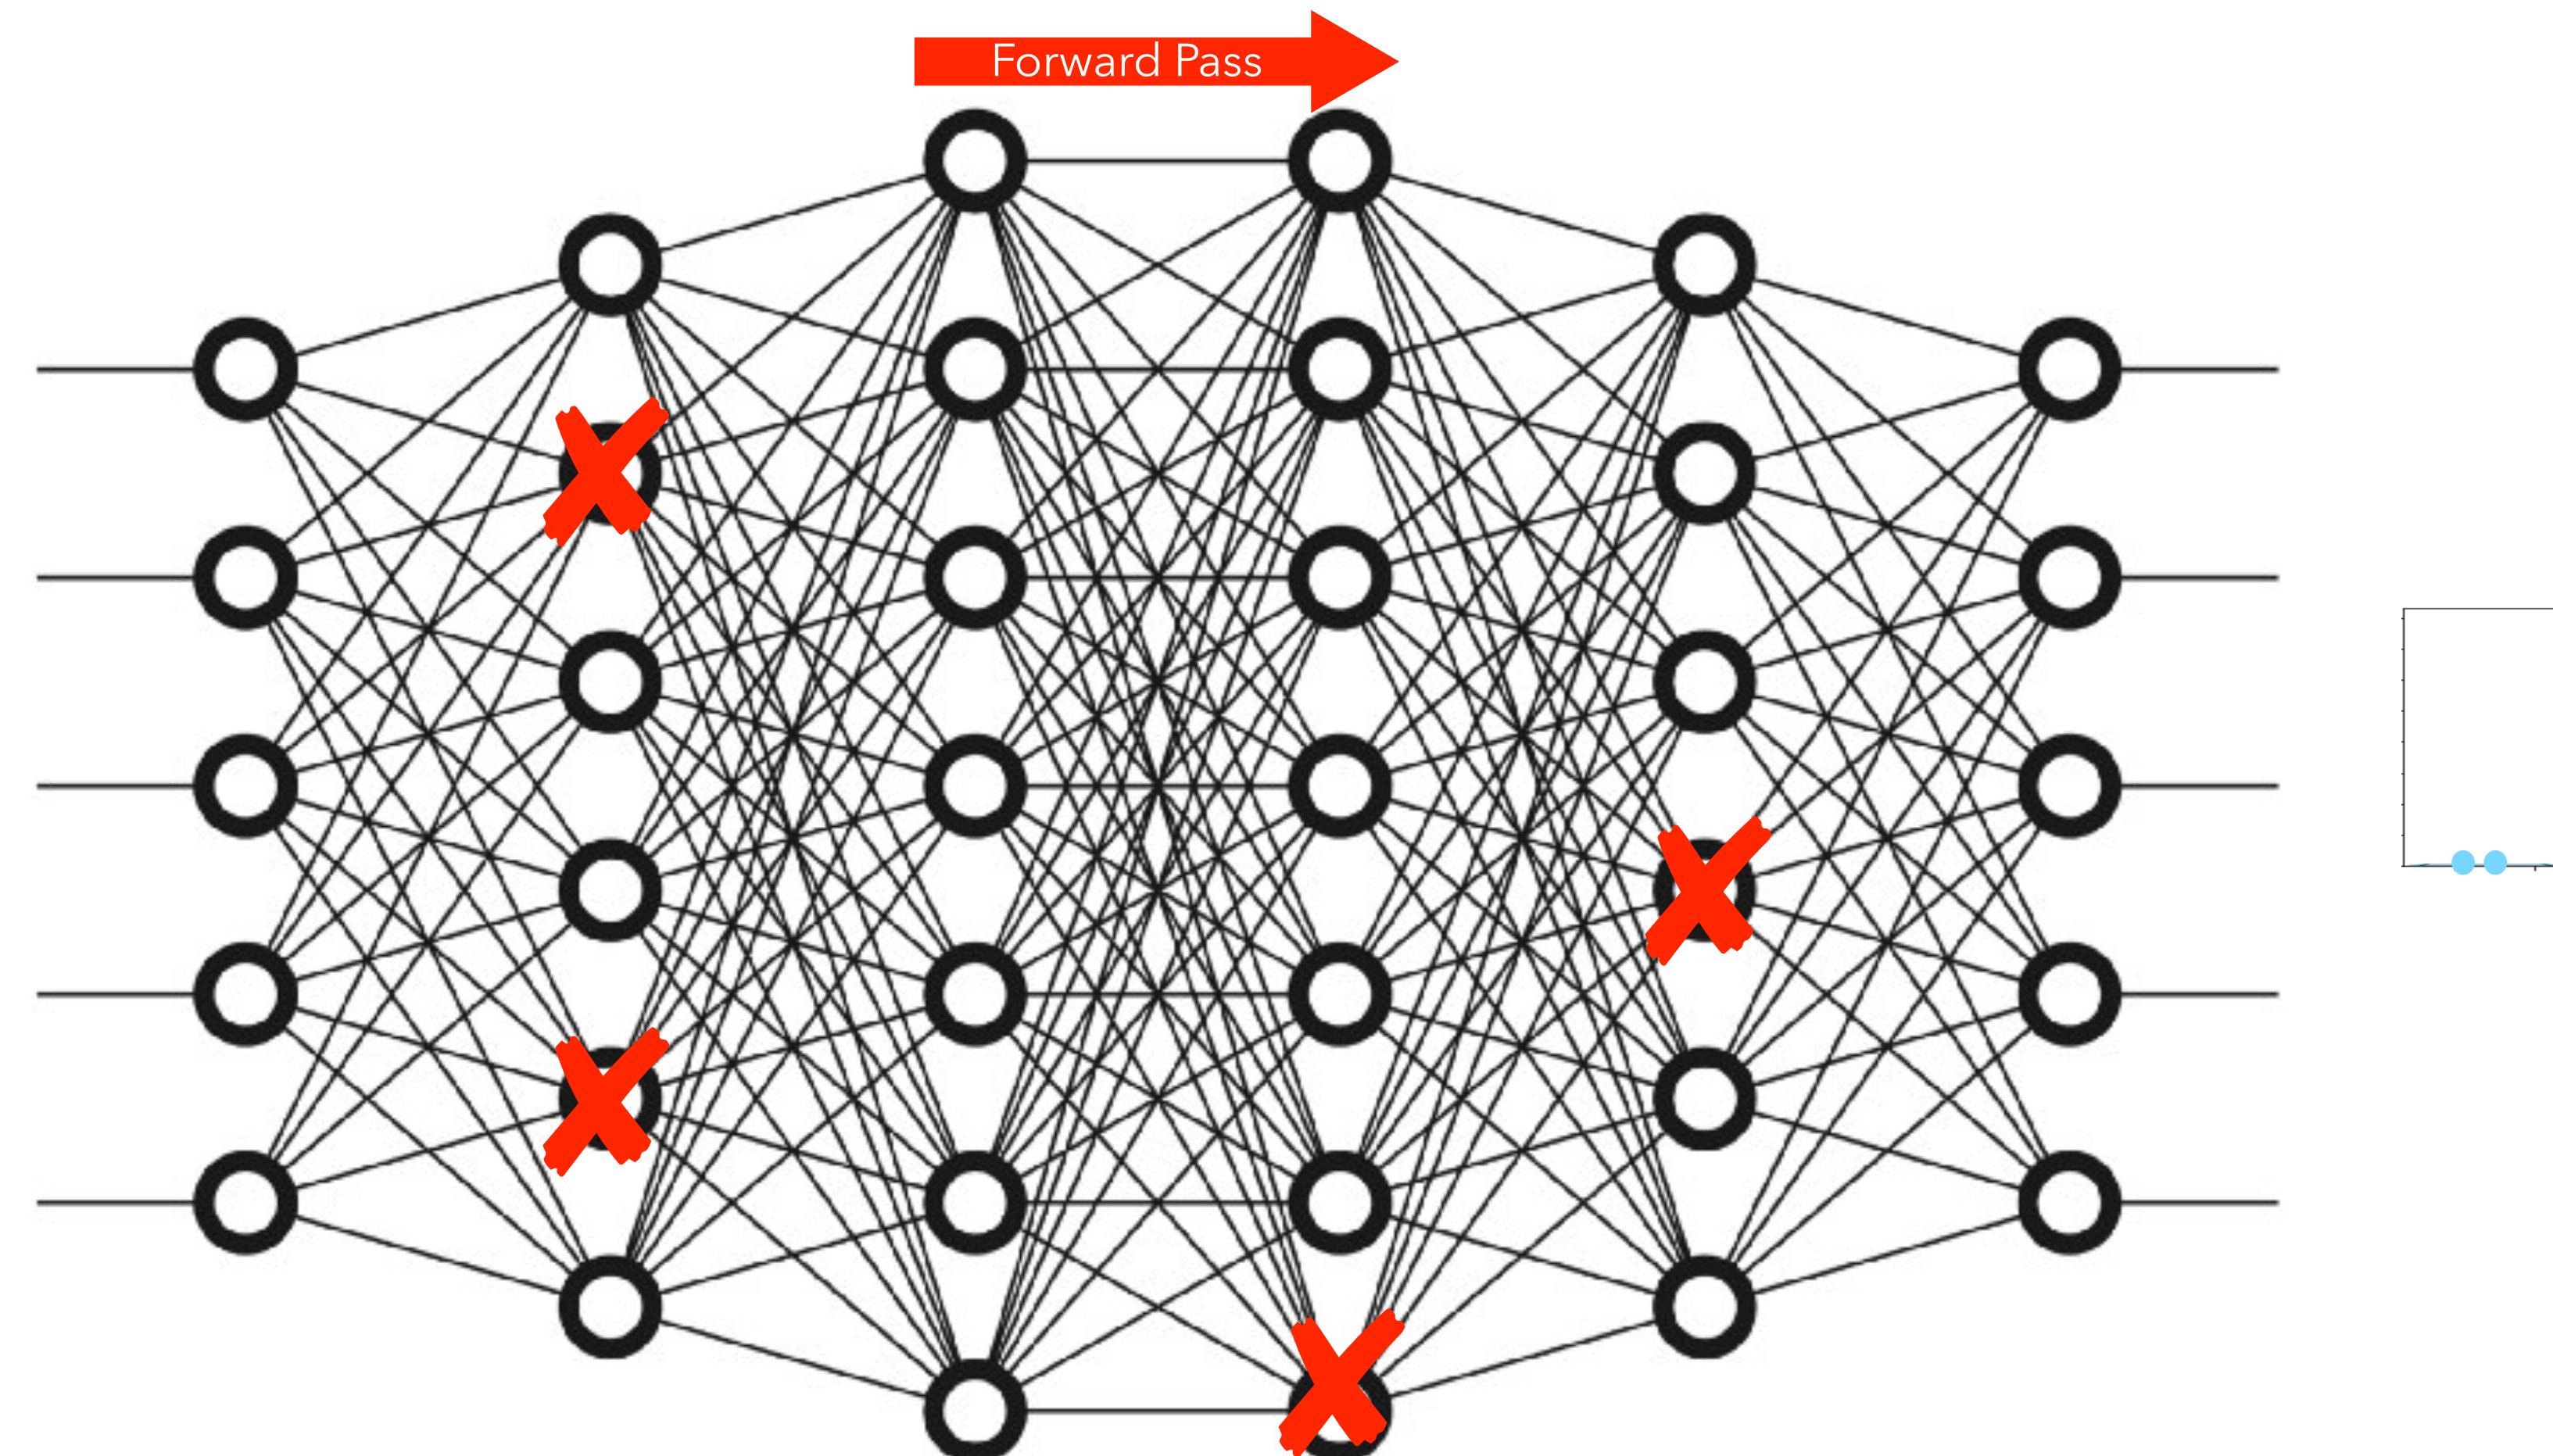
Radu Raicea

# Dropout: evaluate



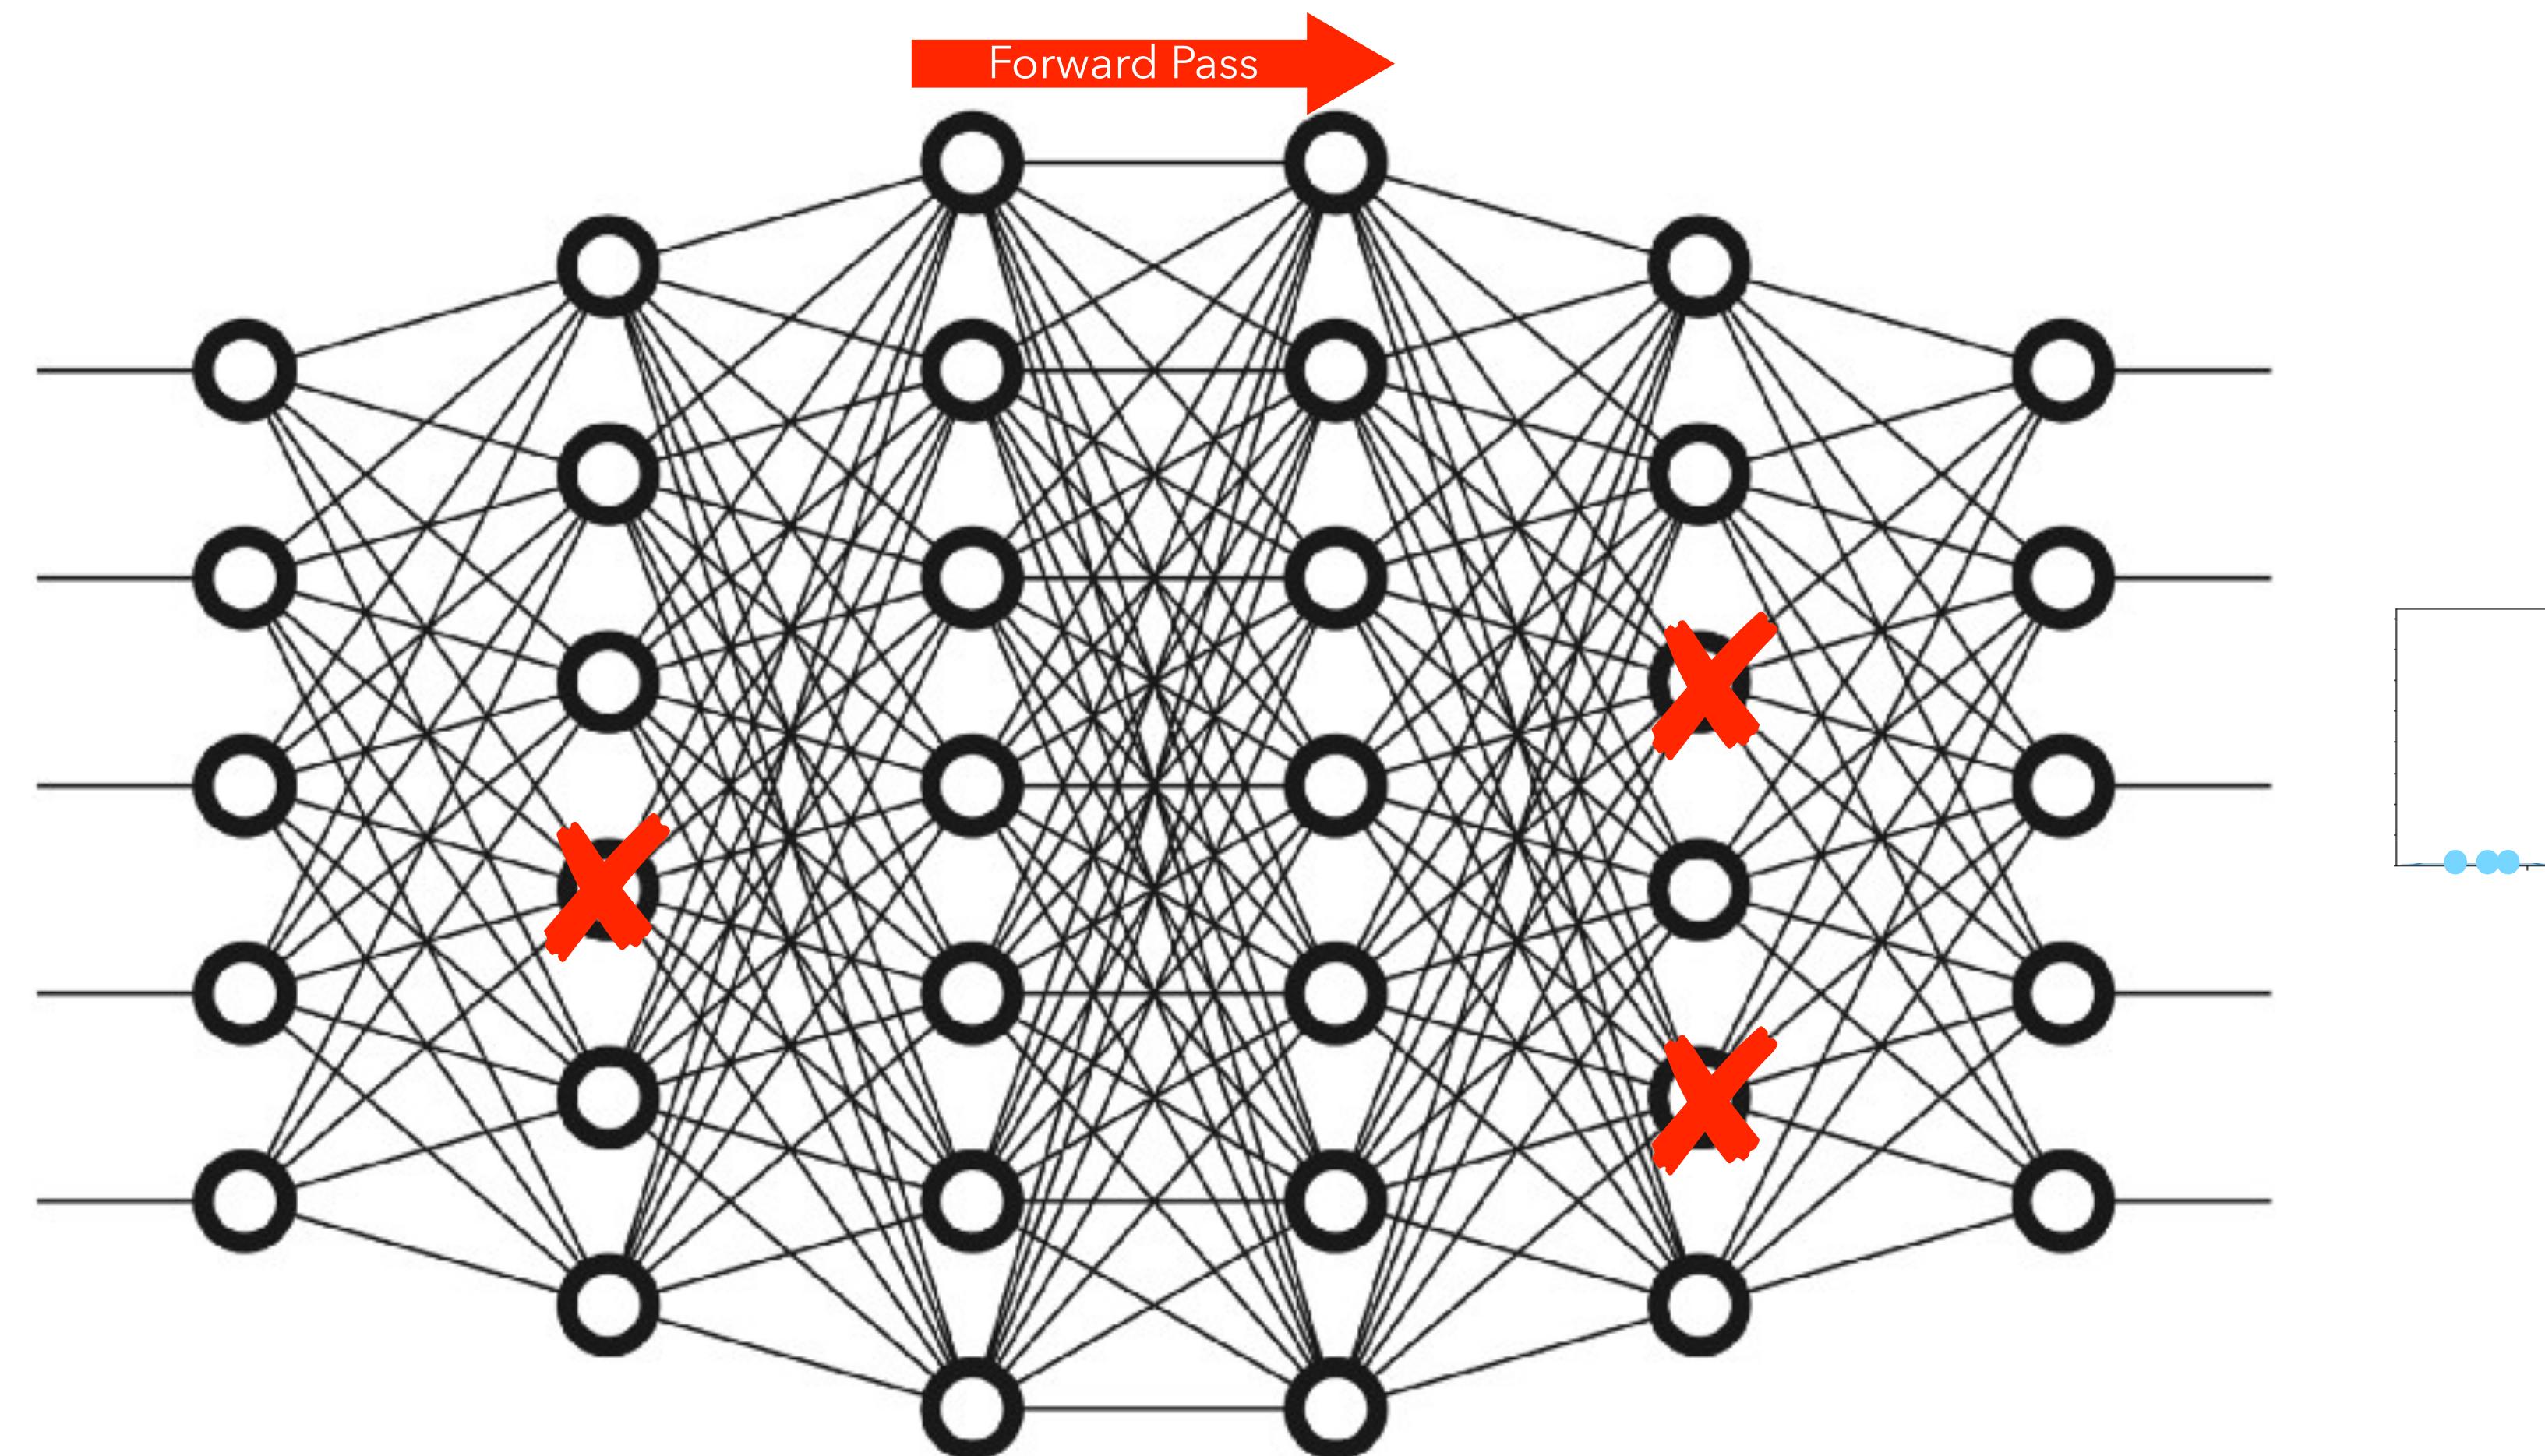
Radu Raicea

# Dropout: evaluate



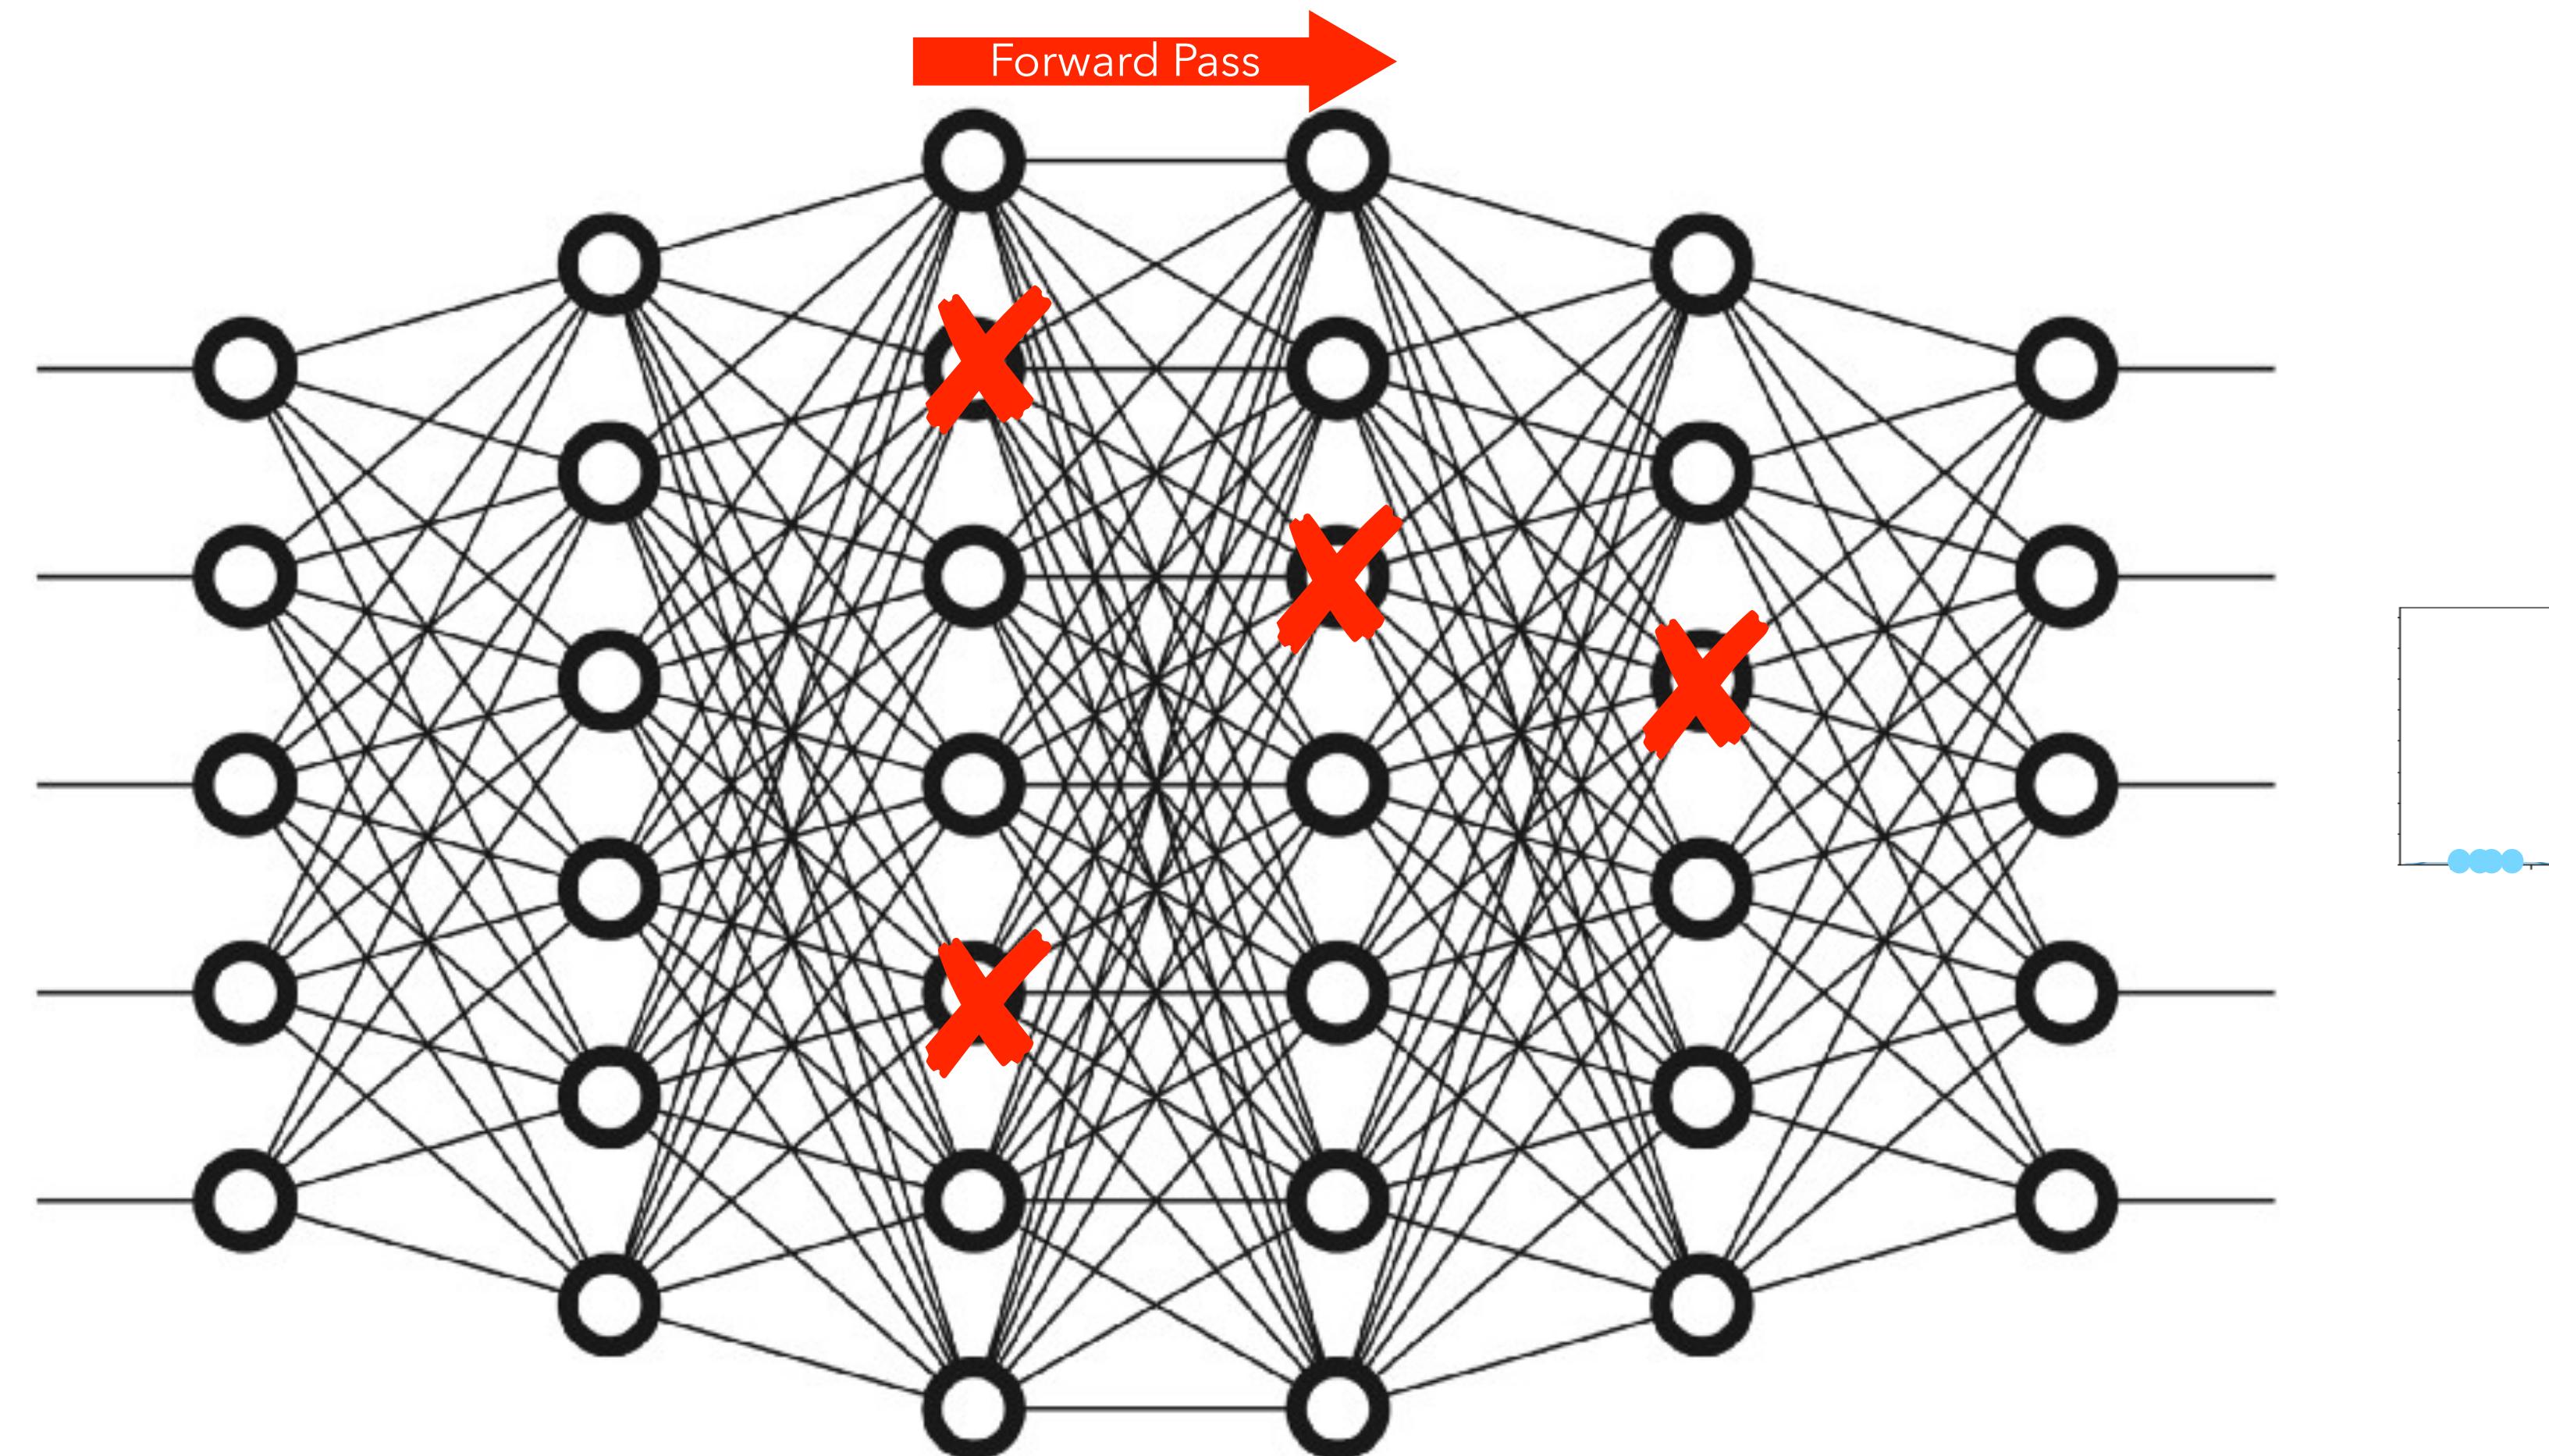
Radu Raicea

# Dropout: evaluate



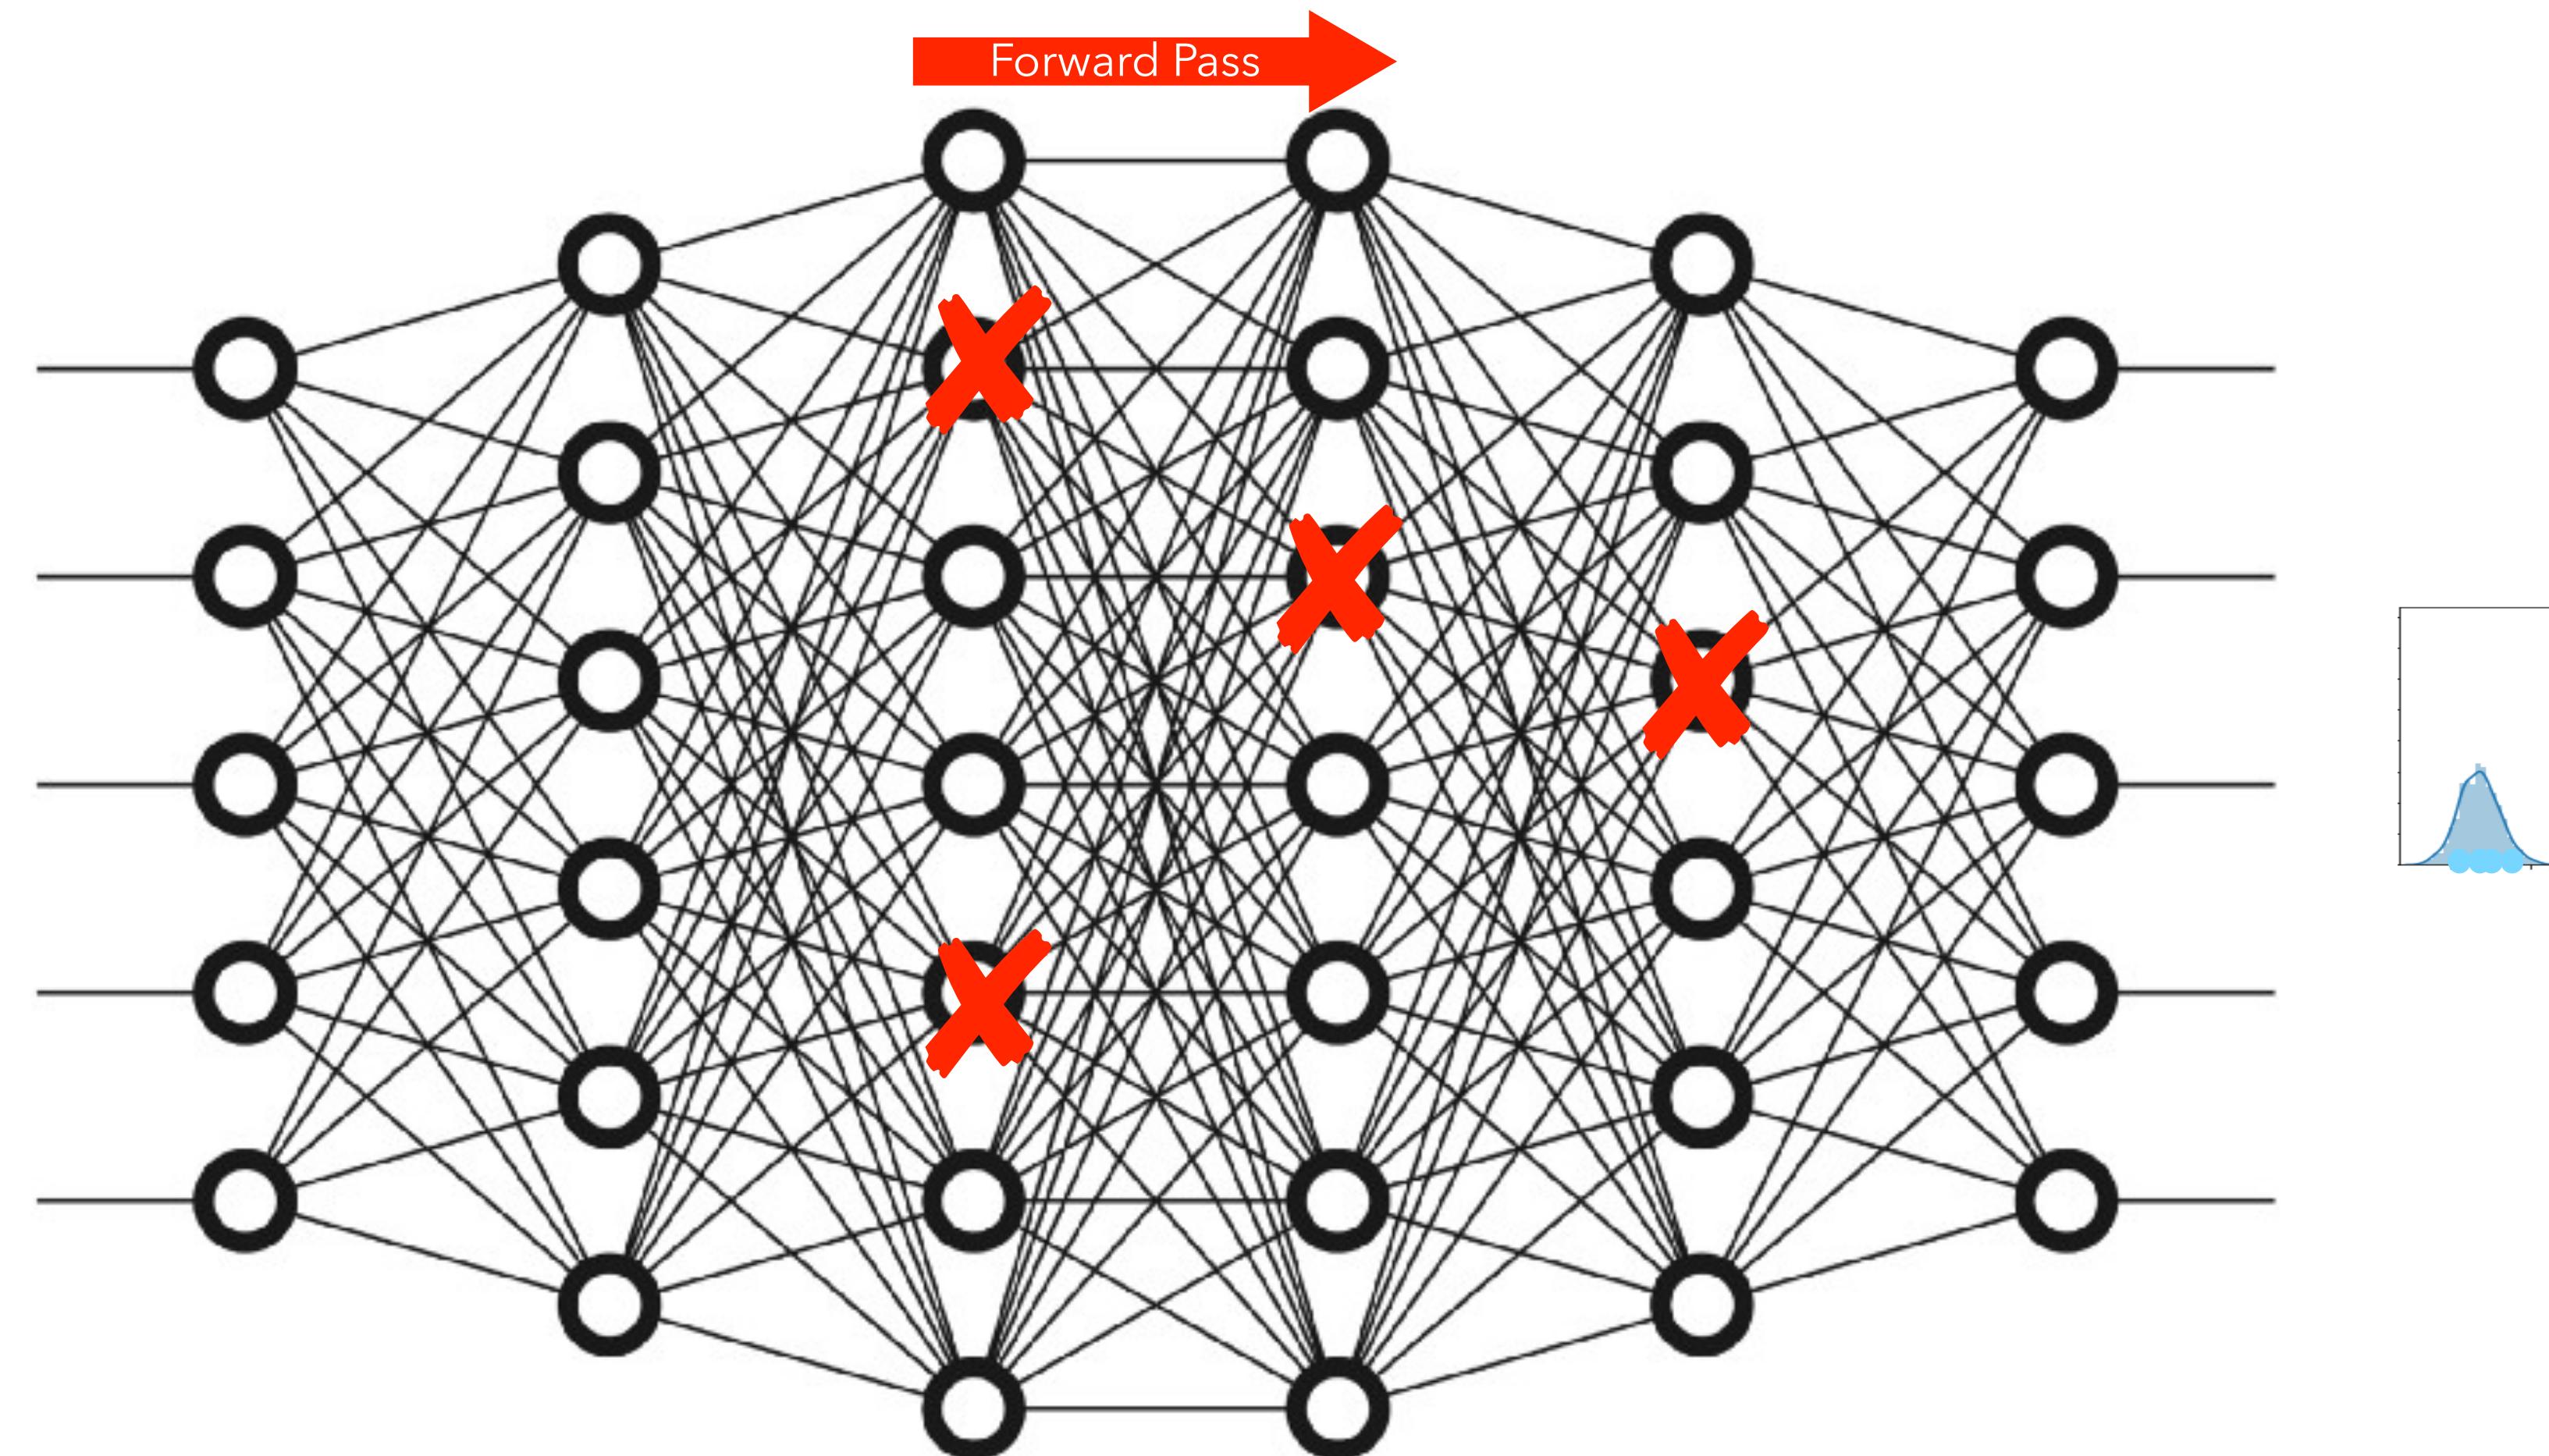
Radu Raicea

# Dropout: evaluate



Radu Raicea

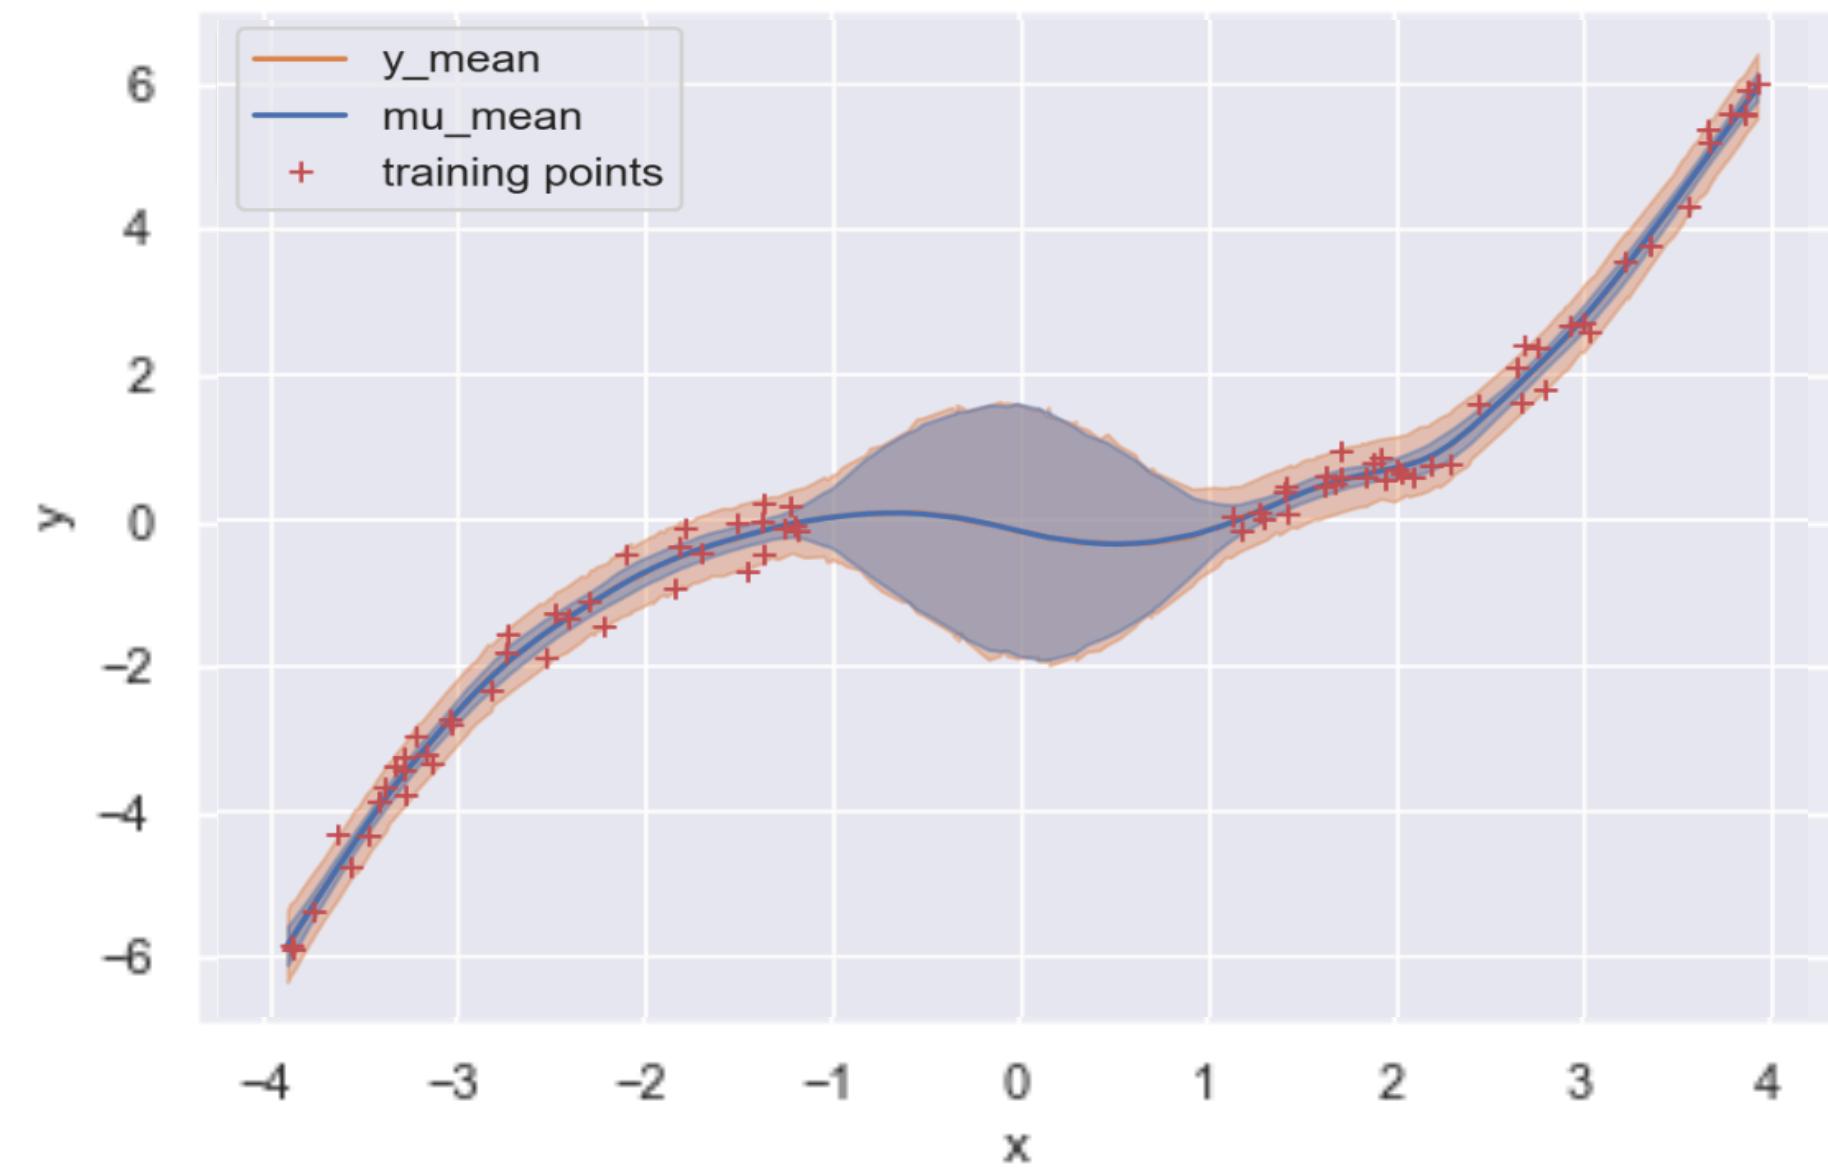
# Dropout: evaluate



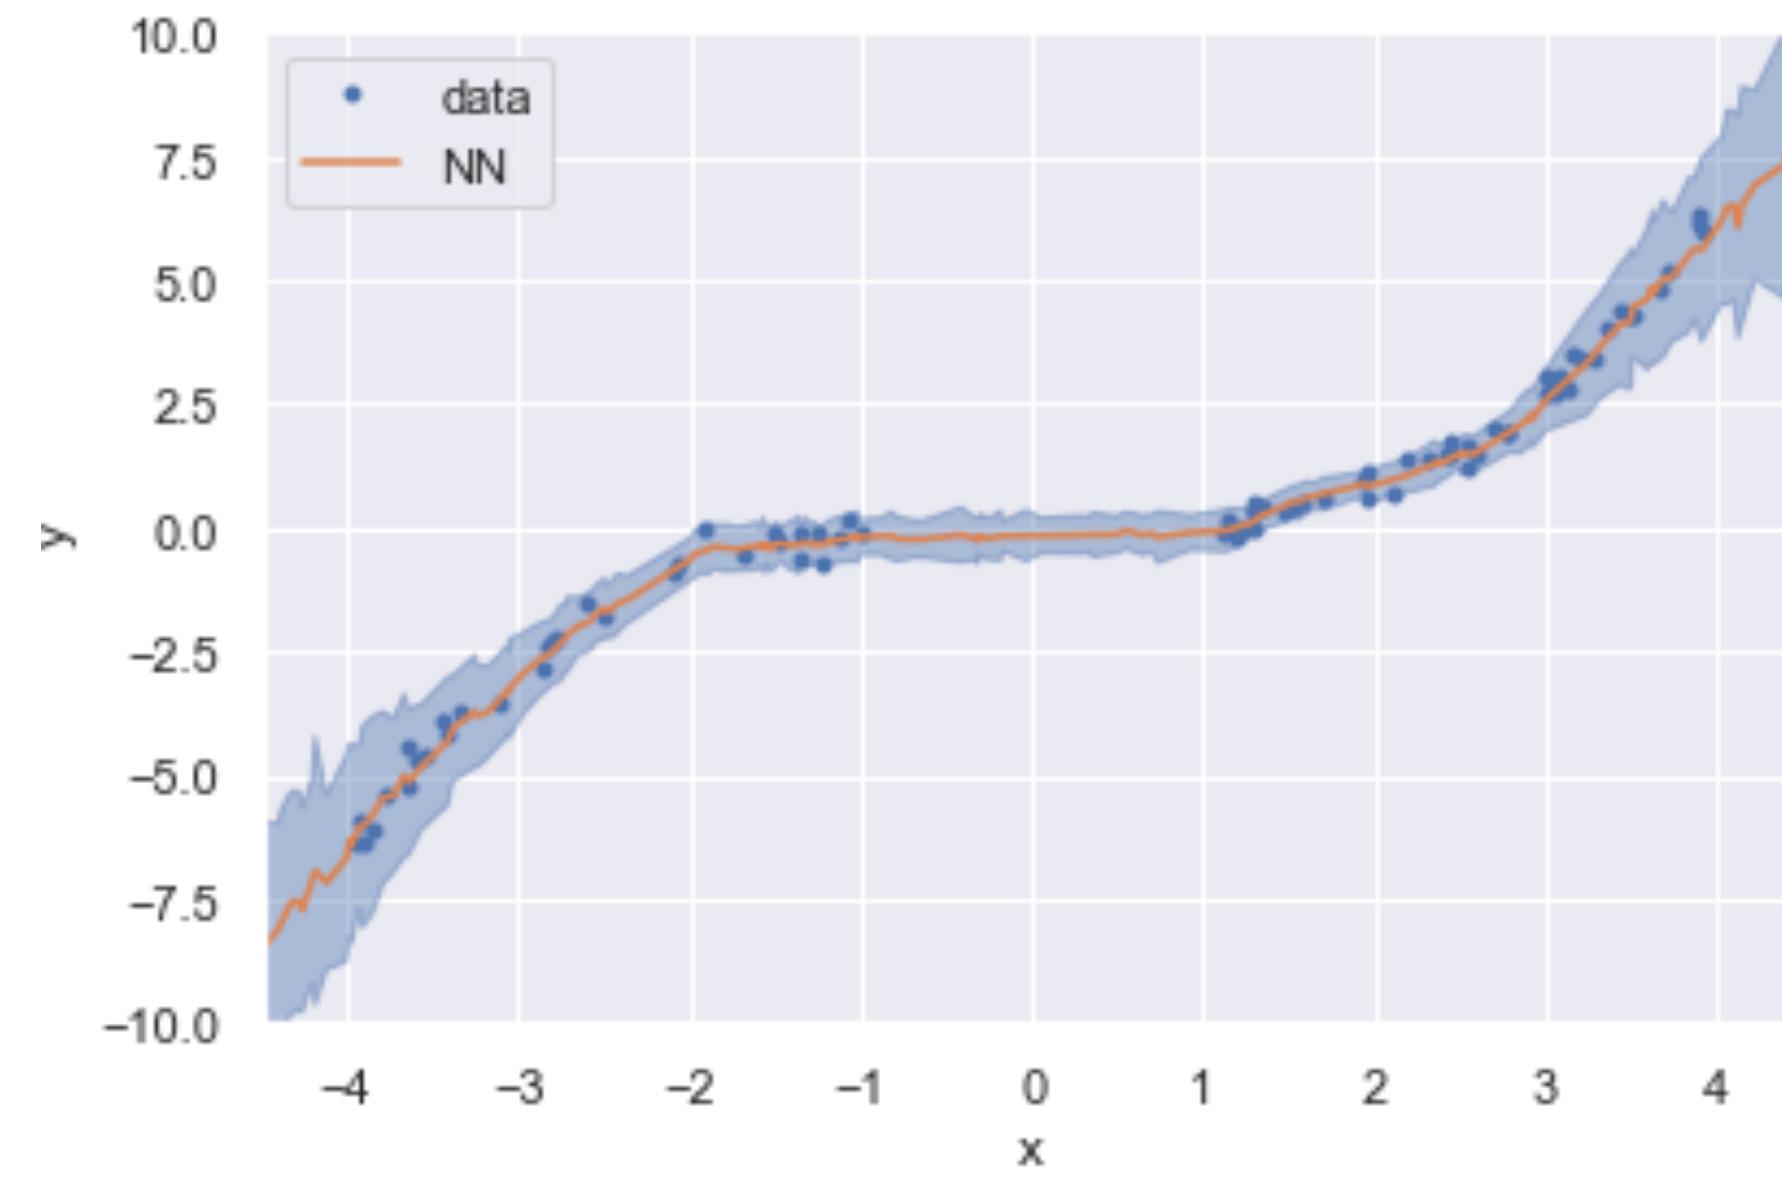
Radu Raicea

# Dropout

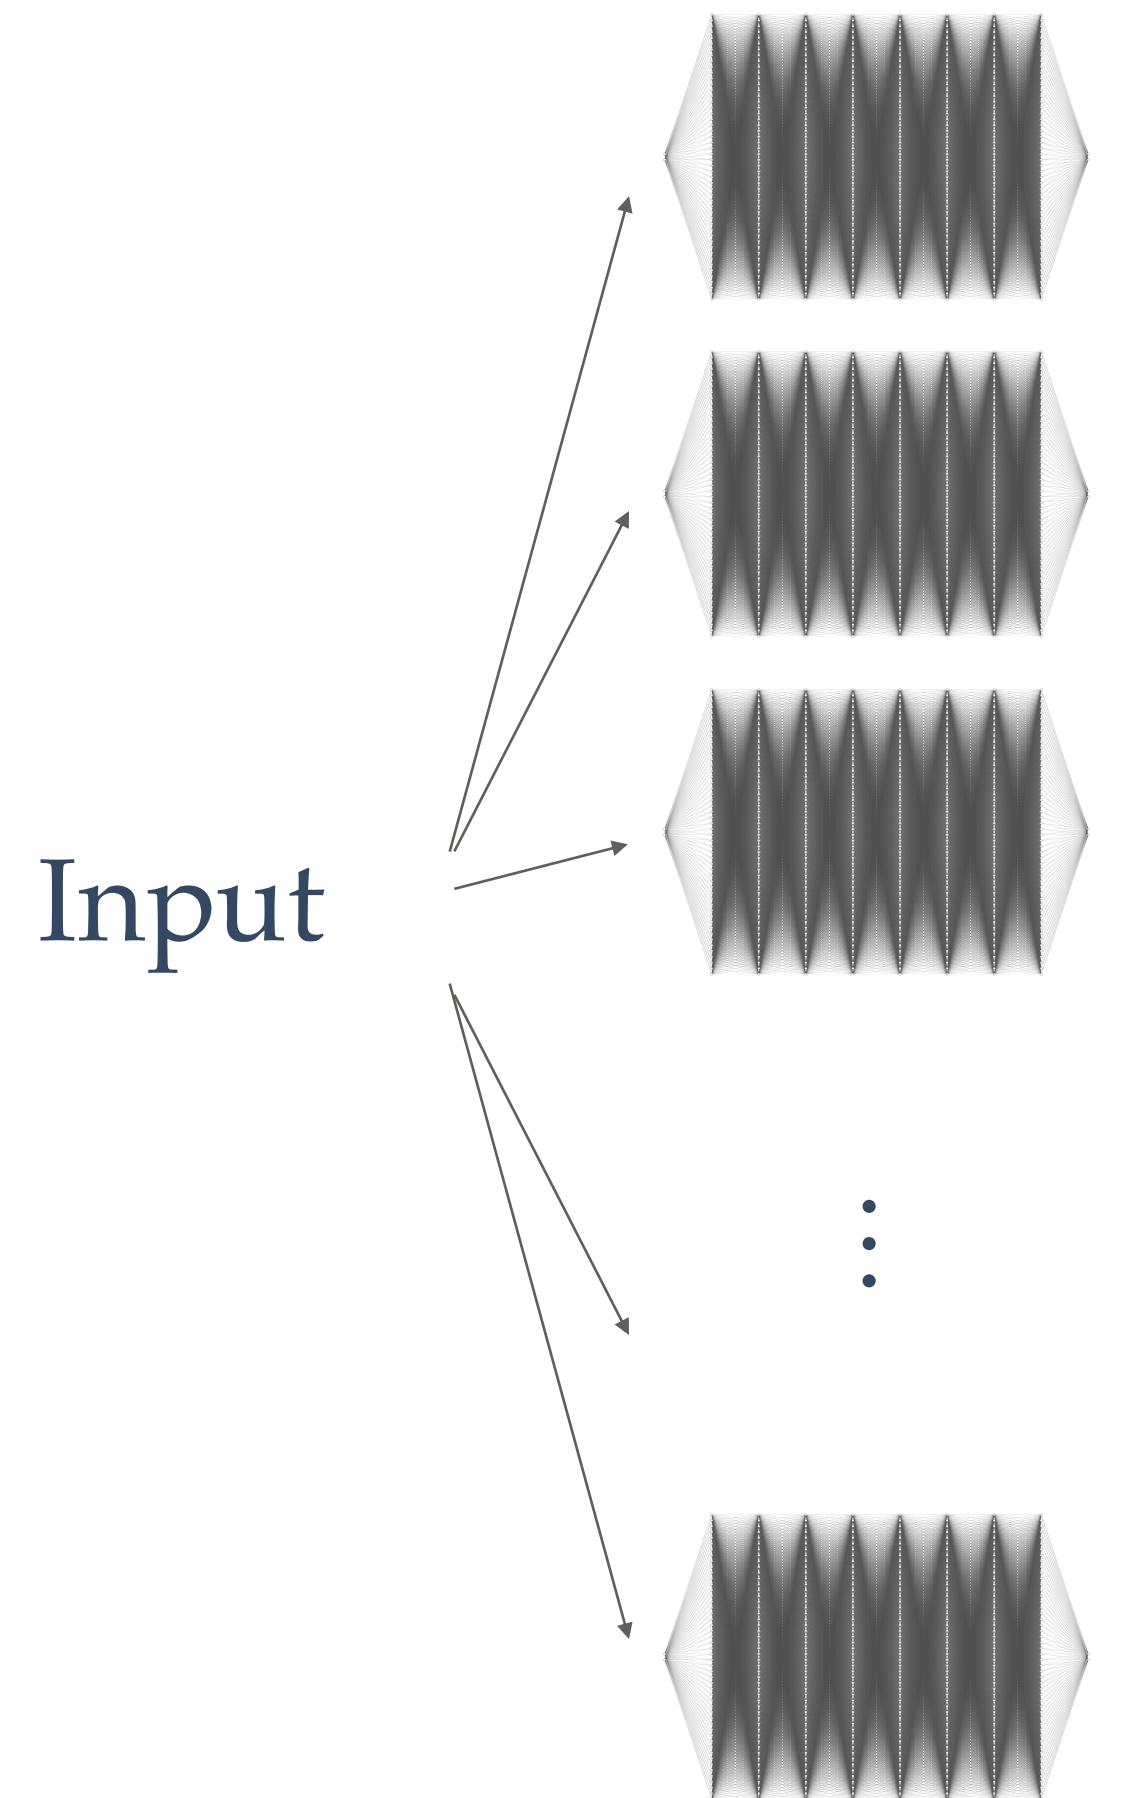
MCMC



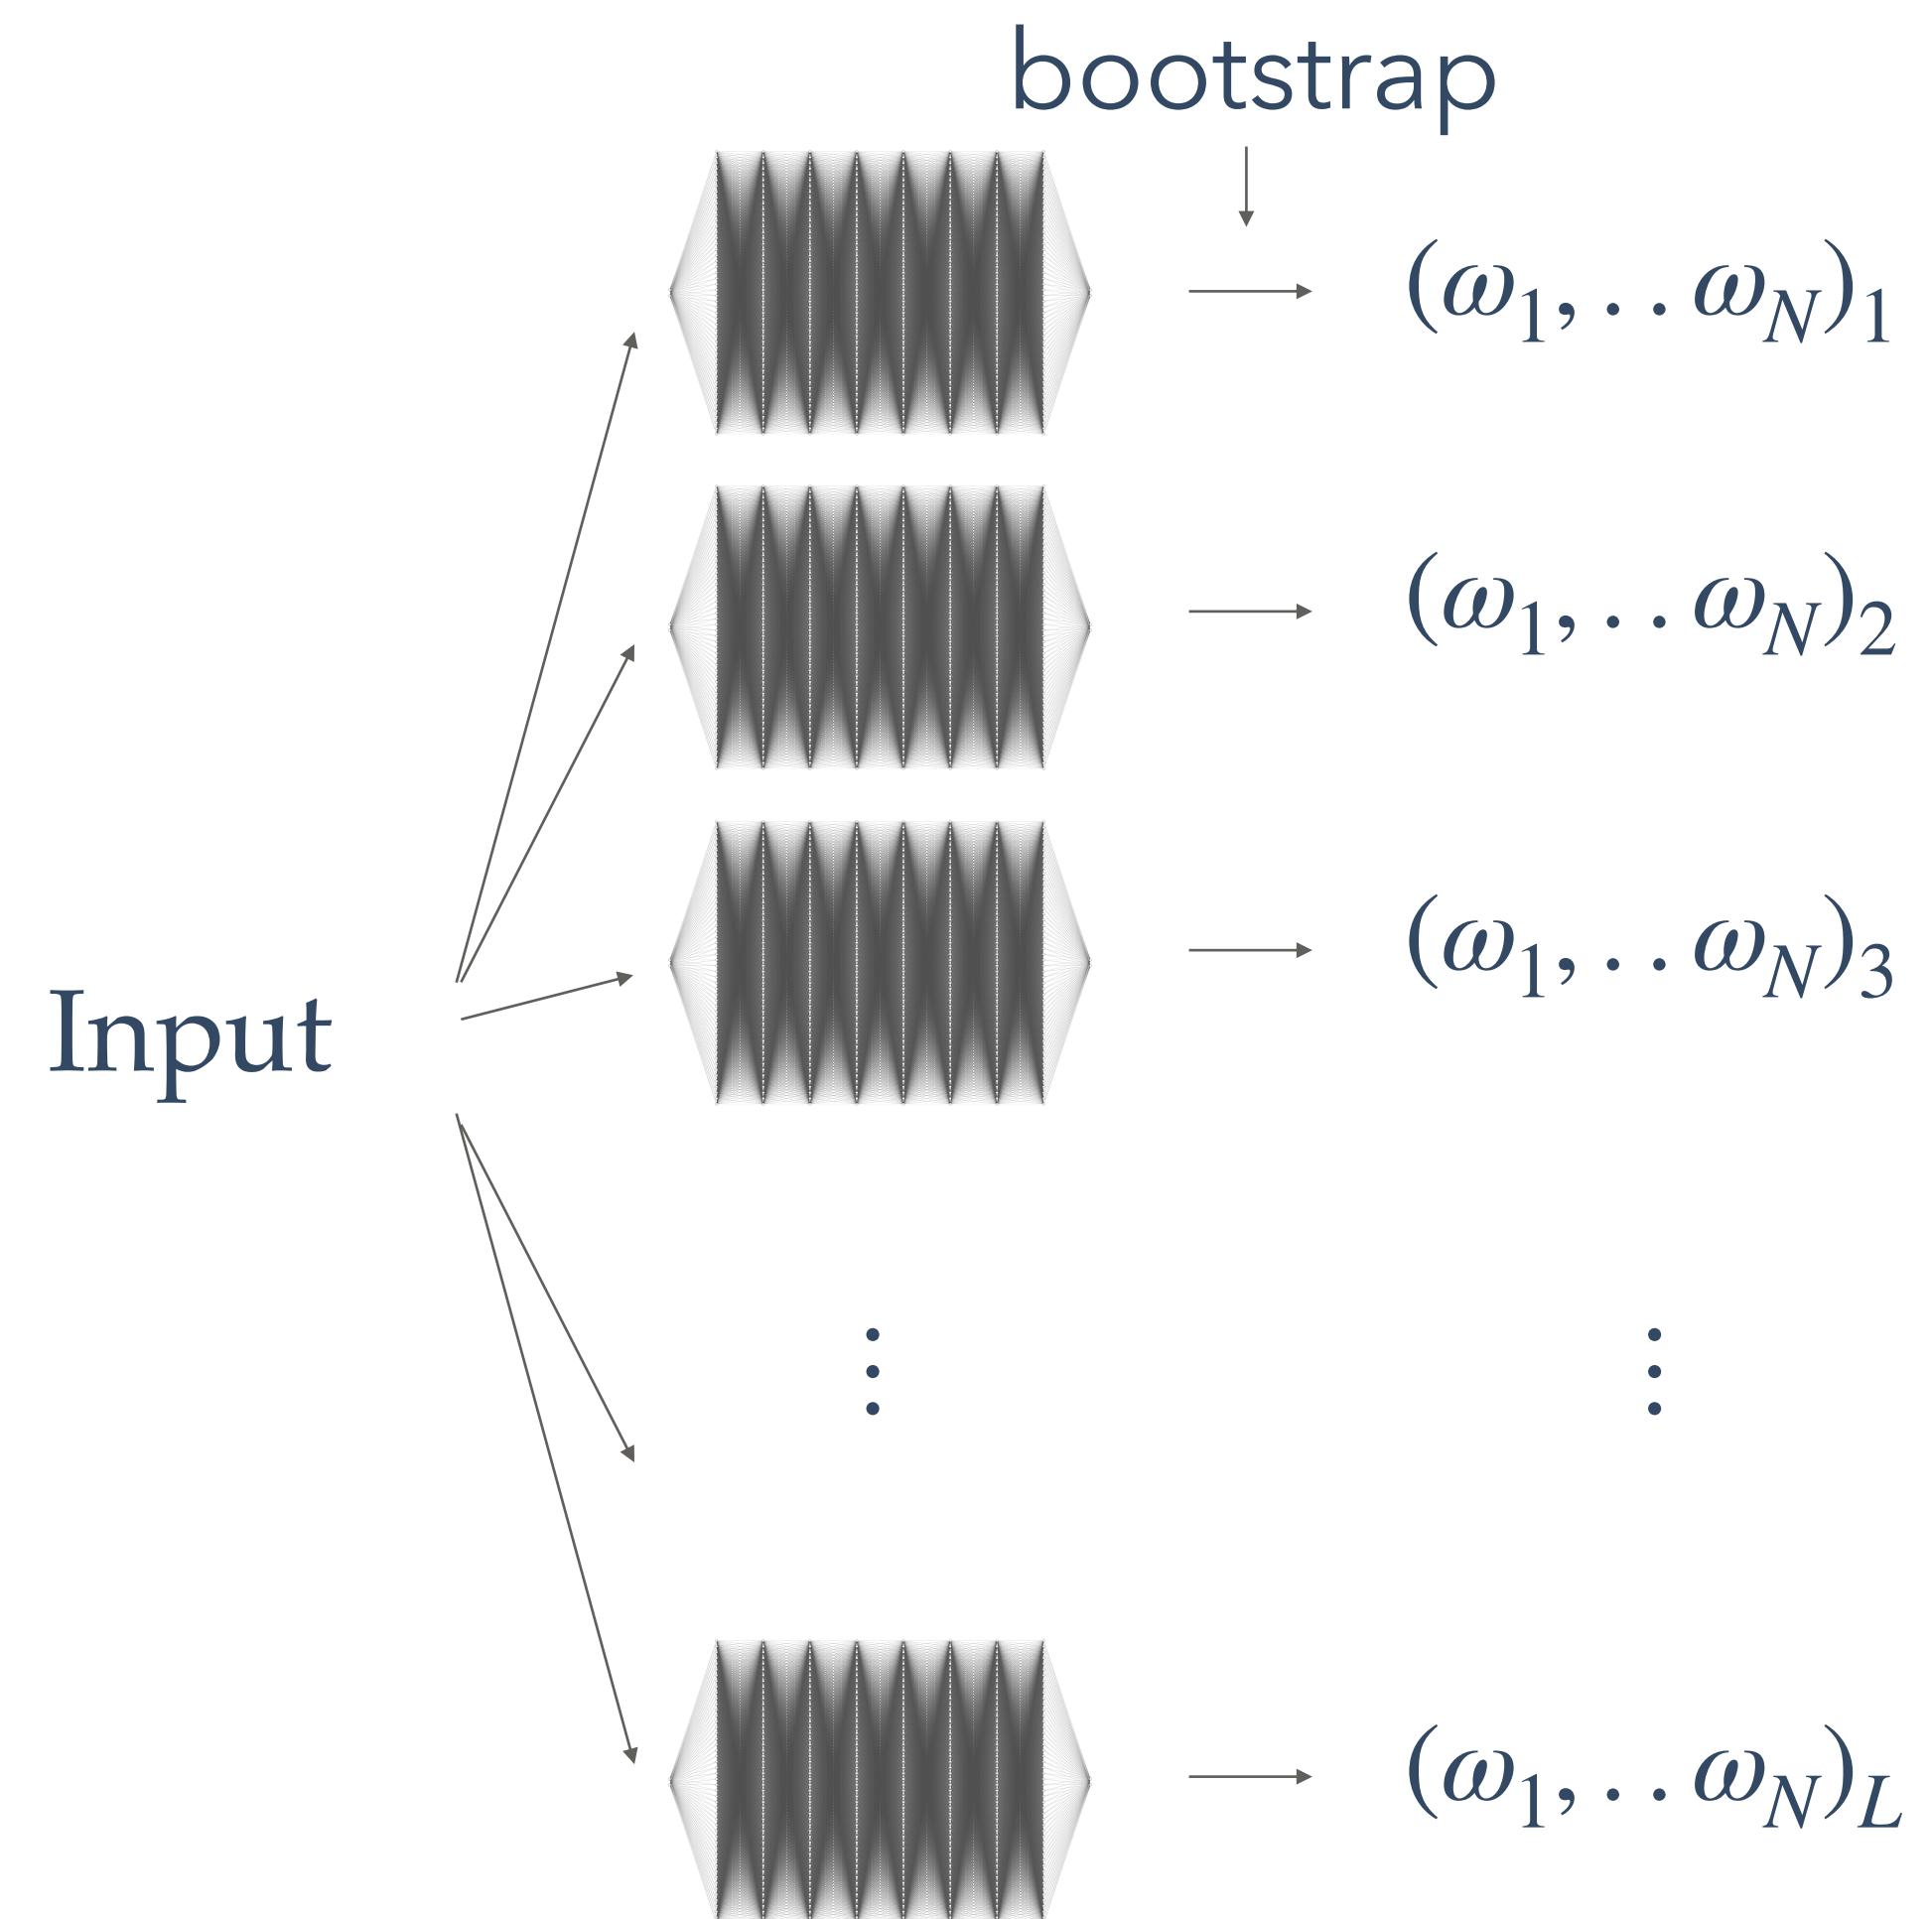
Dropout



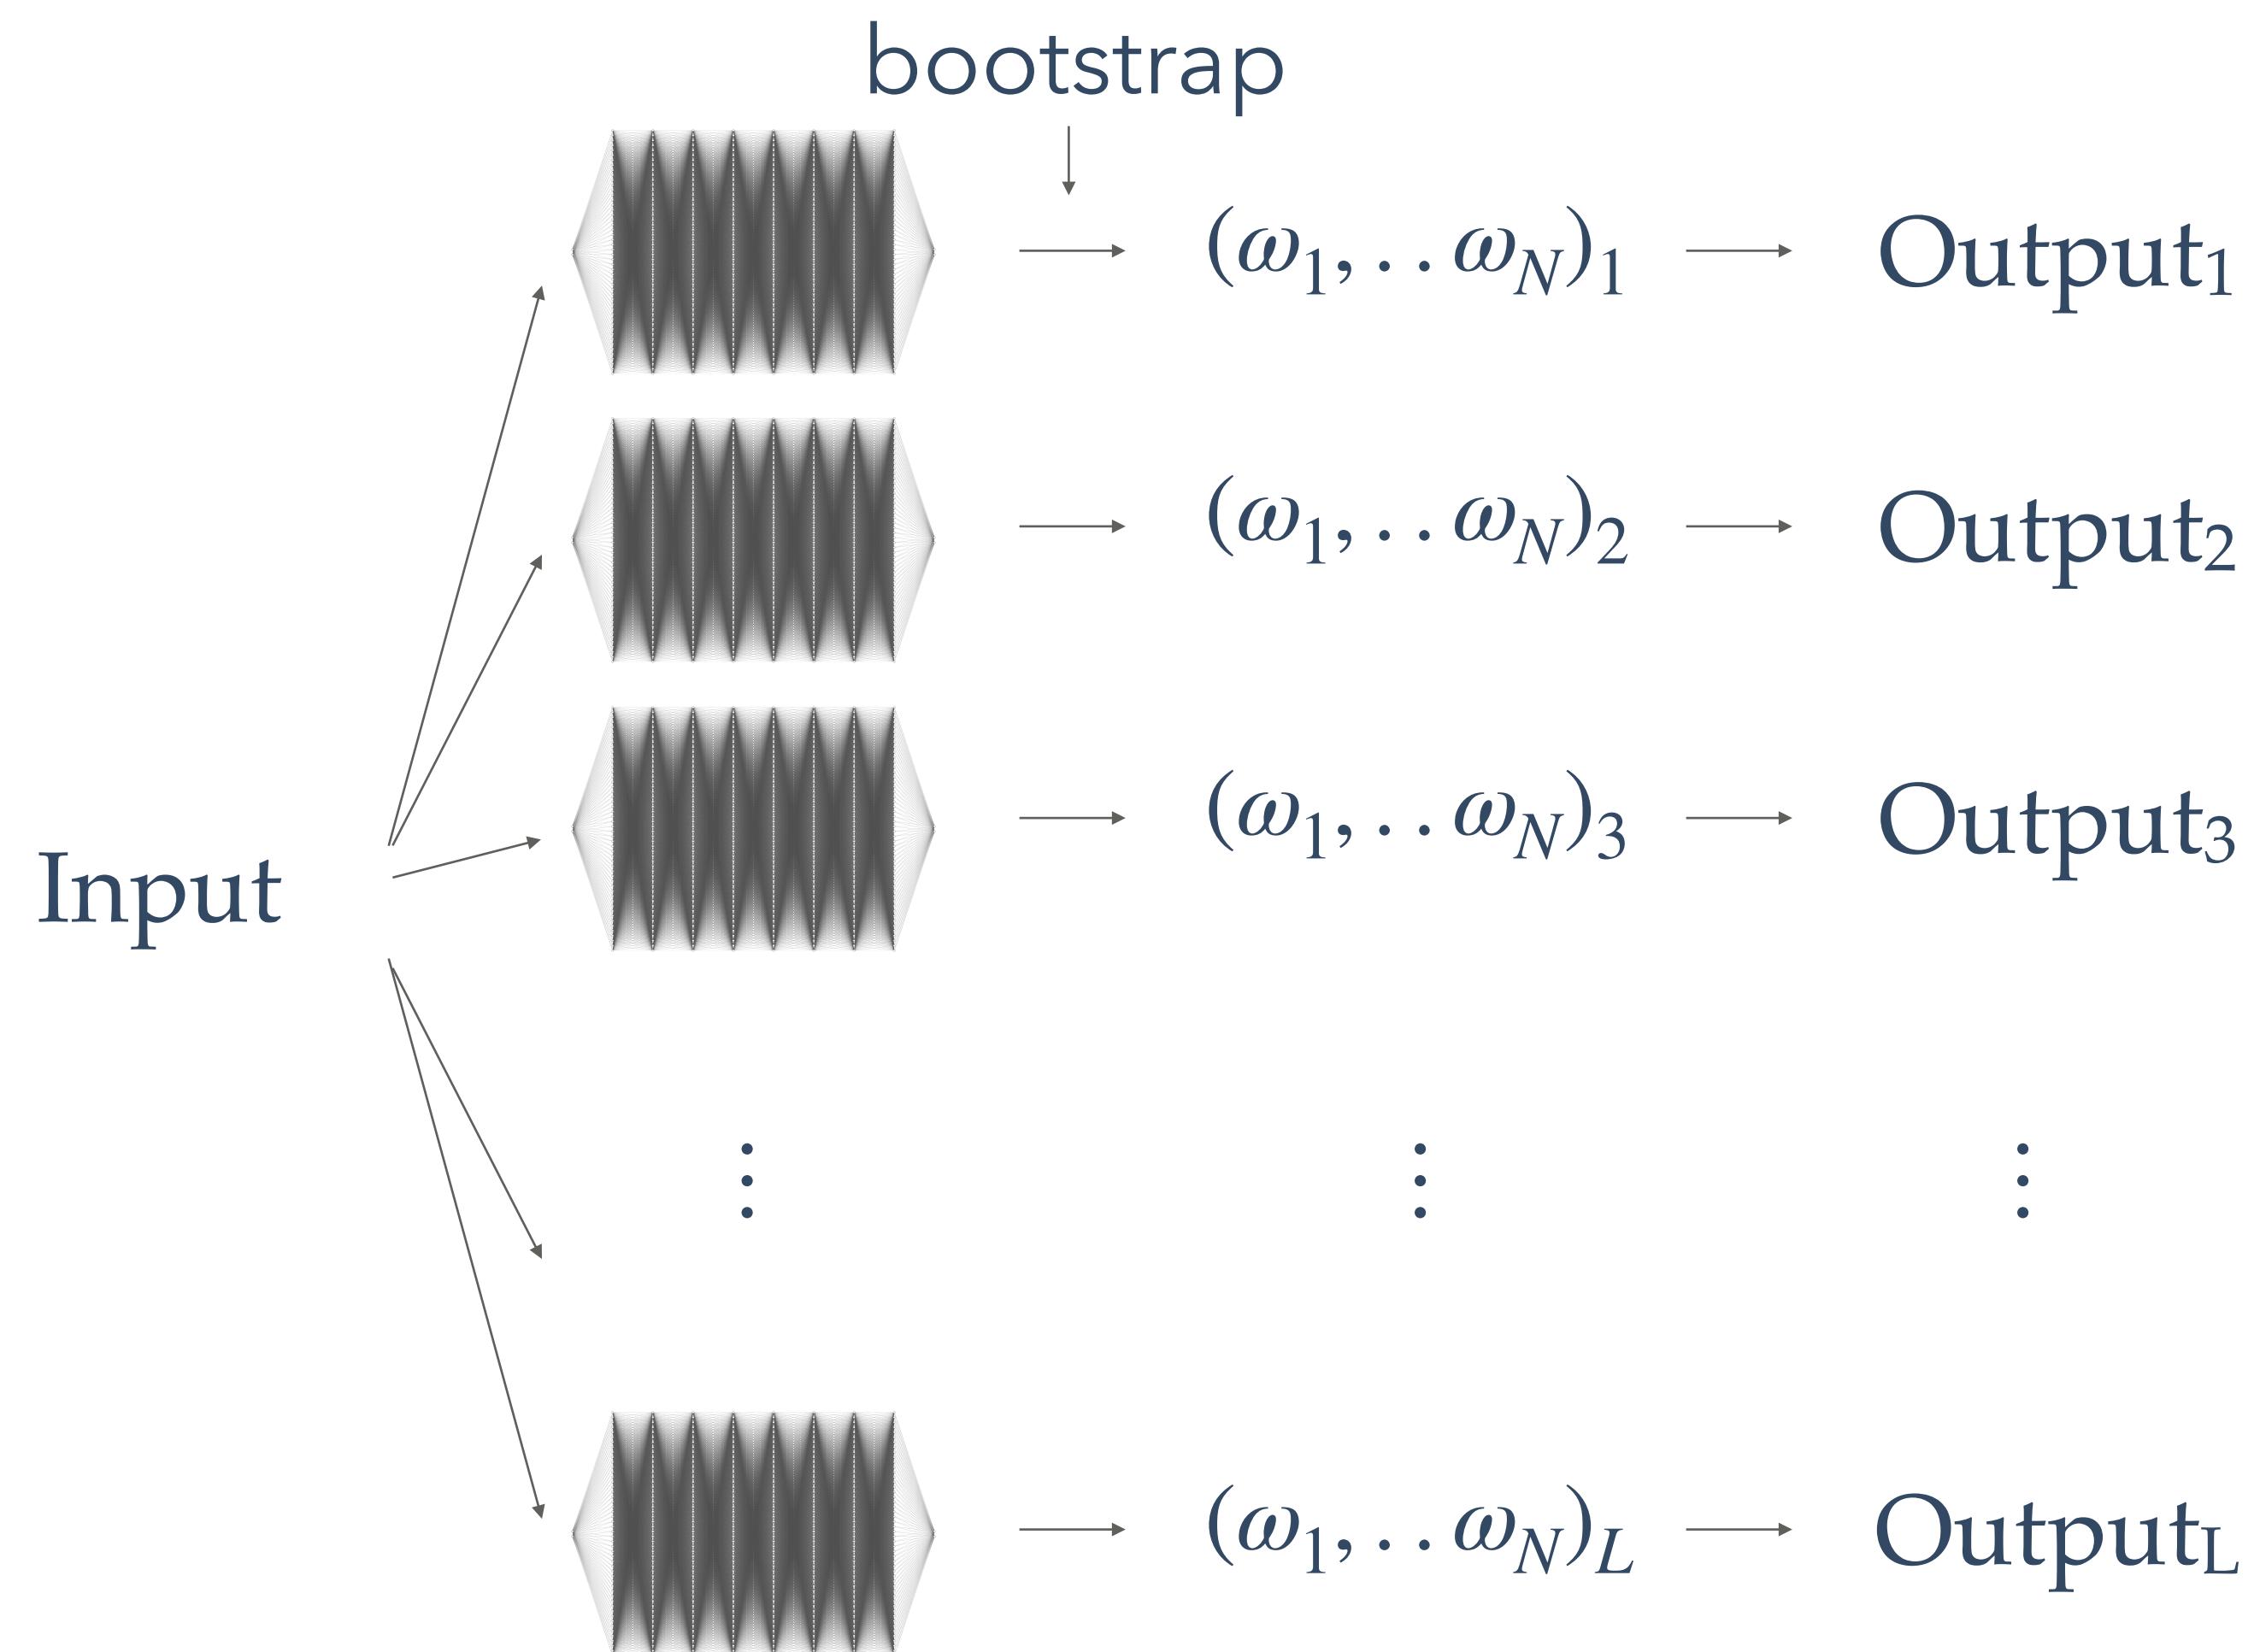
# Bootstrap: L models



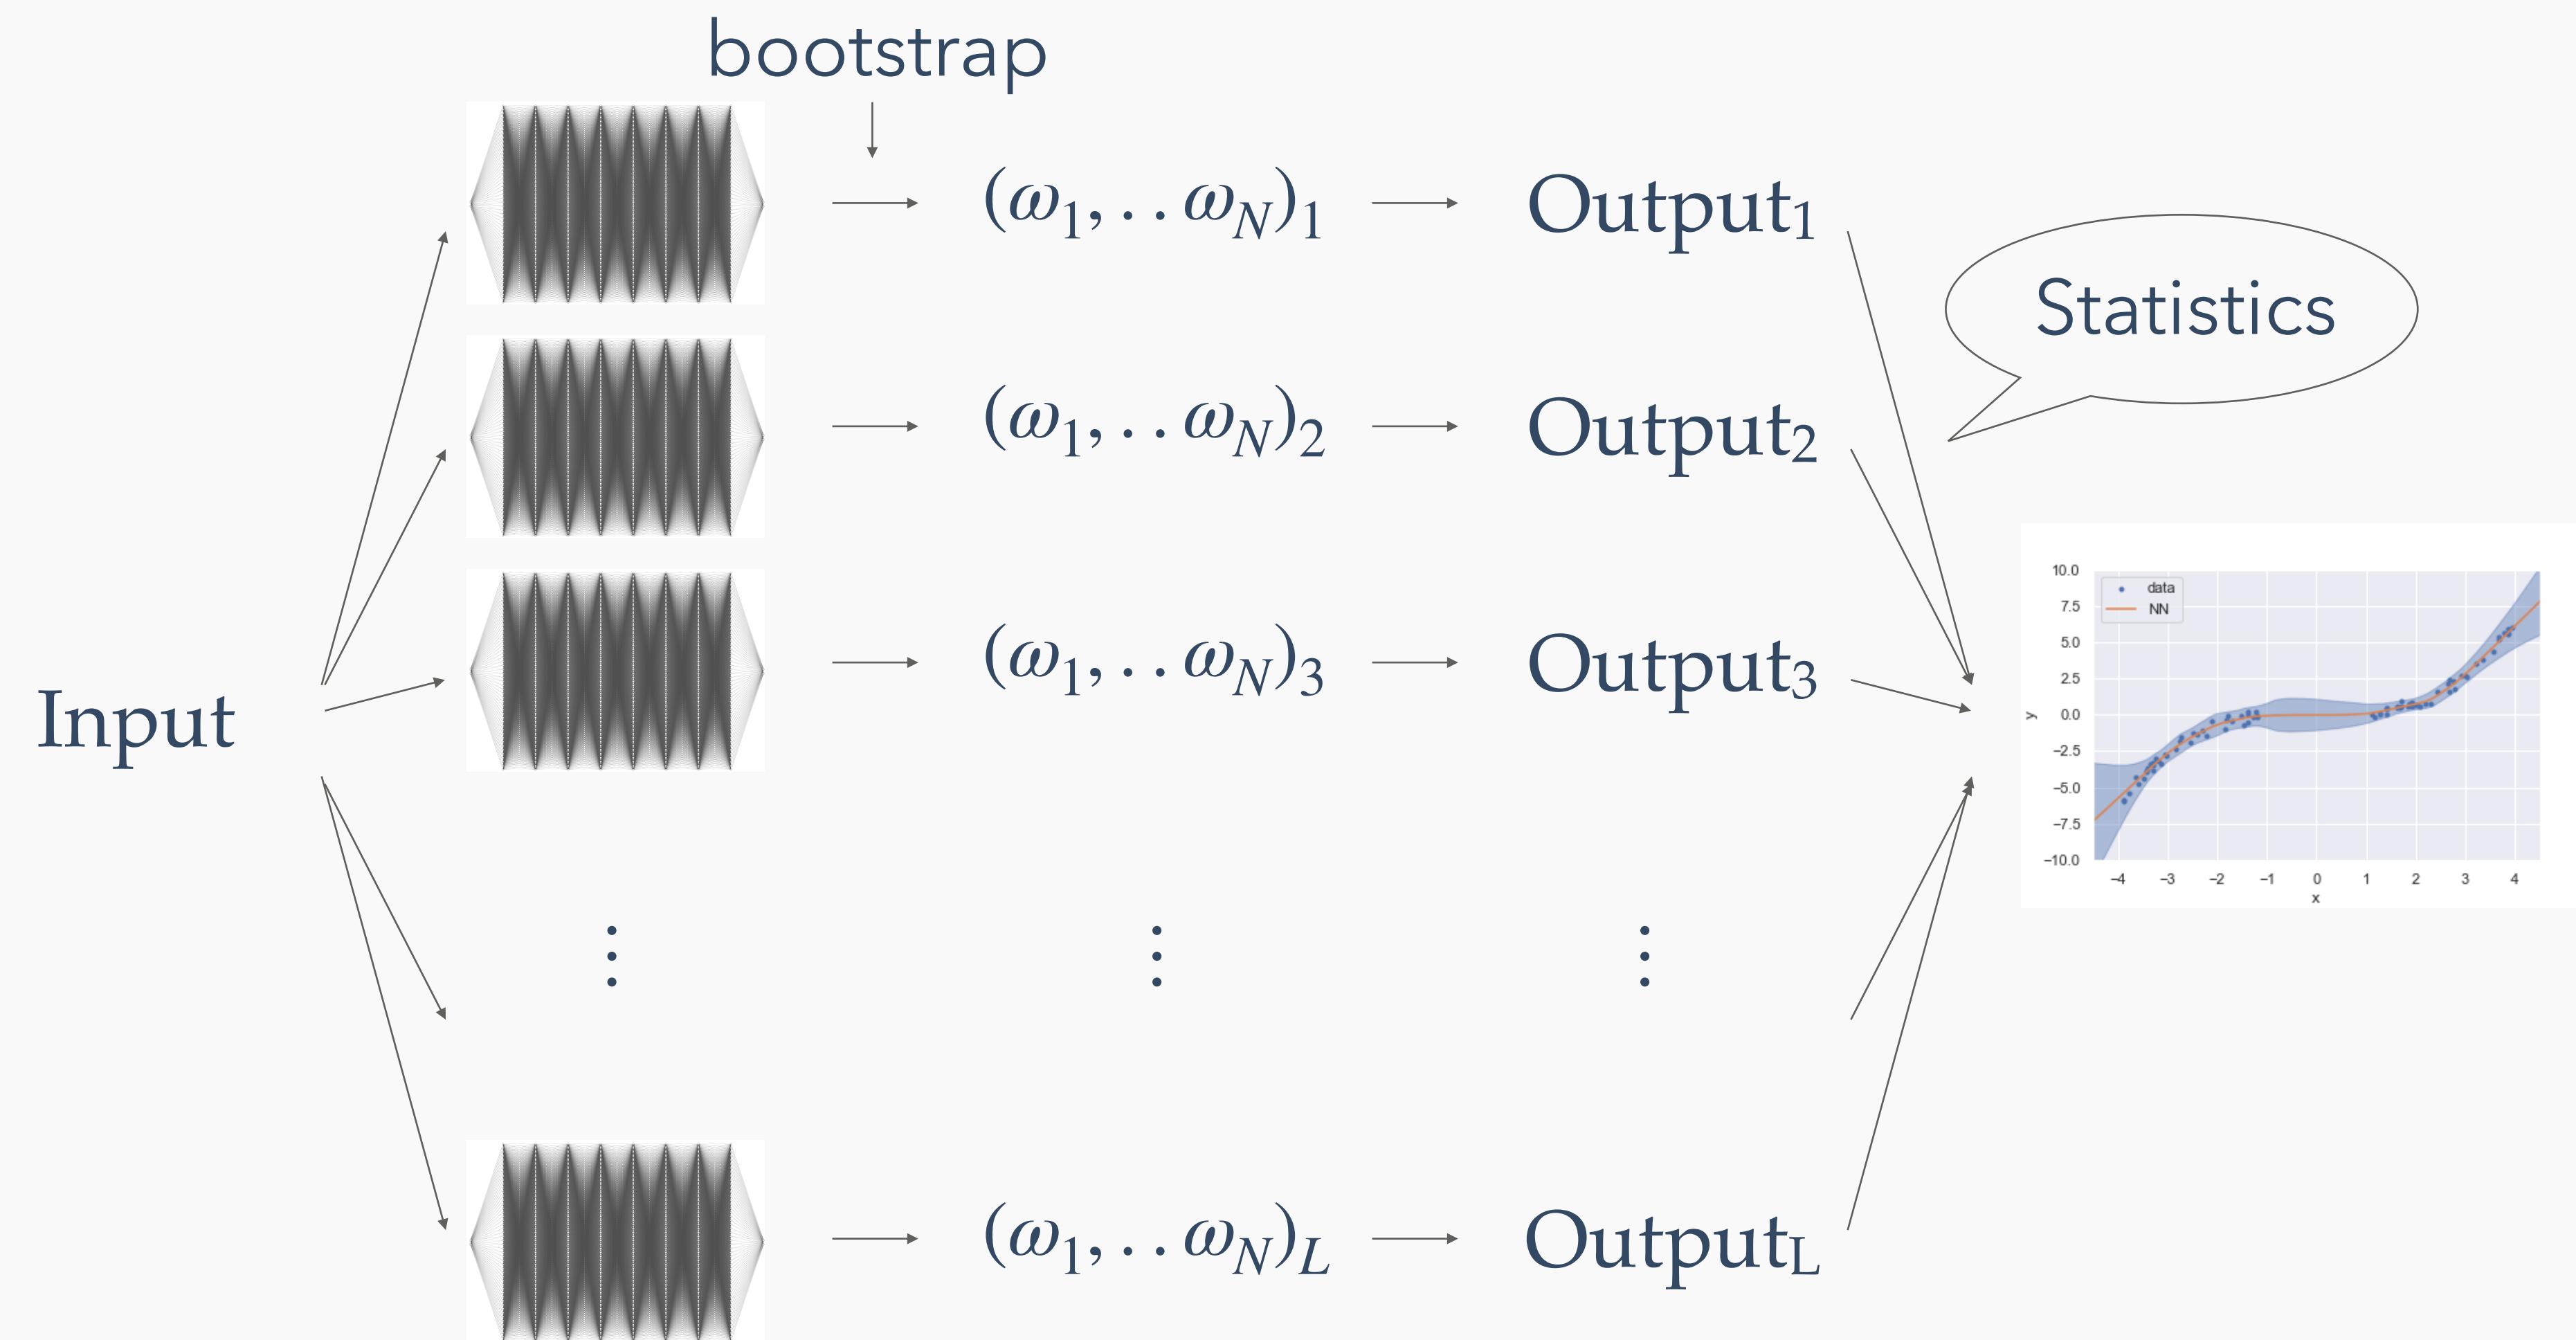
# Bootstrap: L models



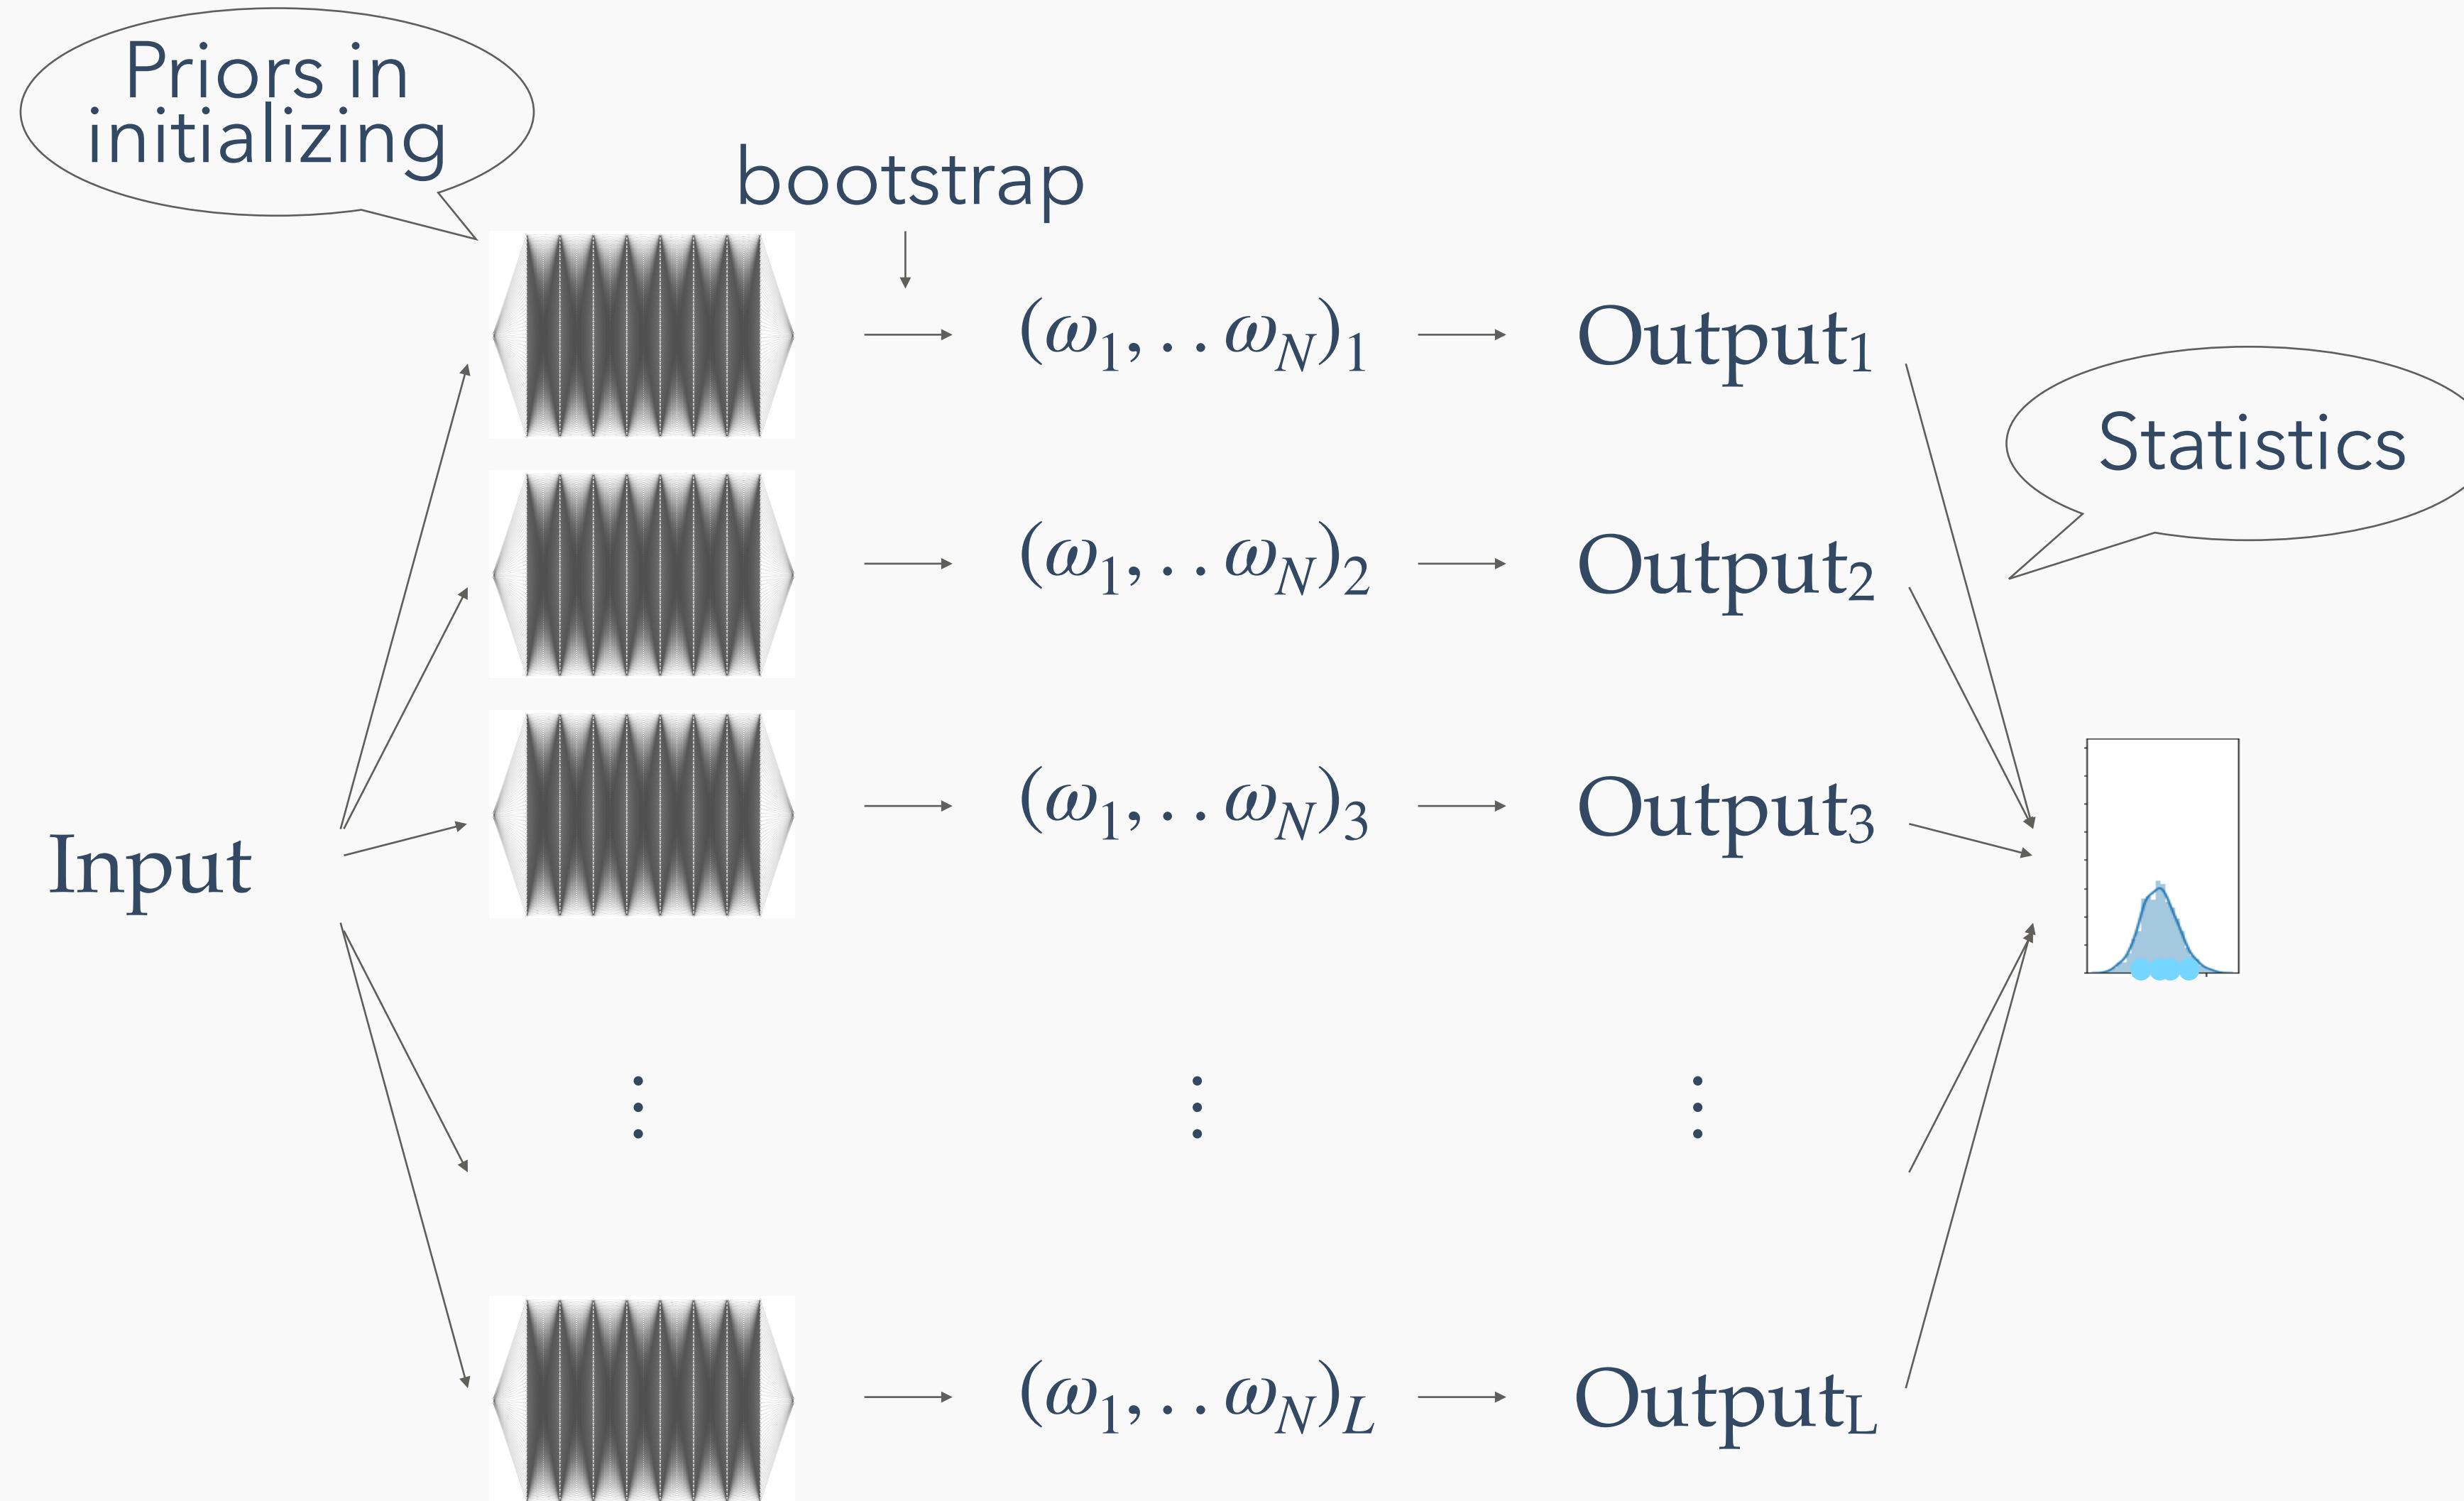
# Bootstrap: L models



# Bootstrap: L models

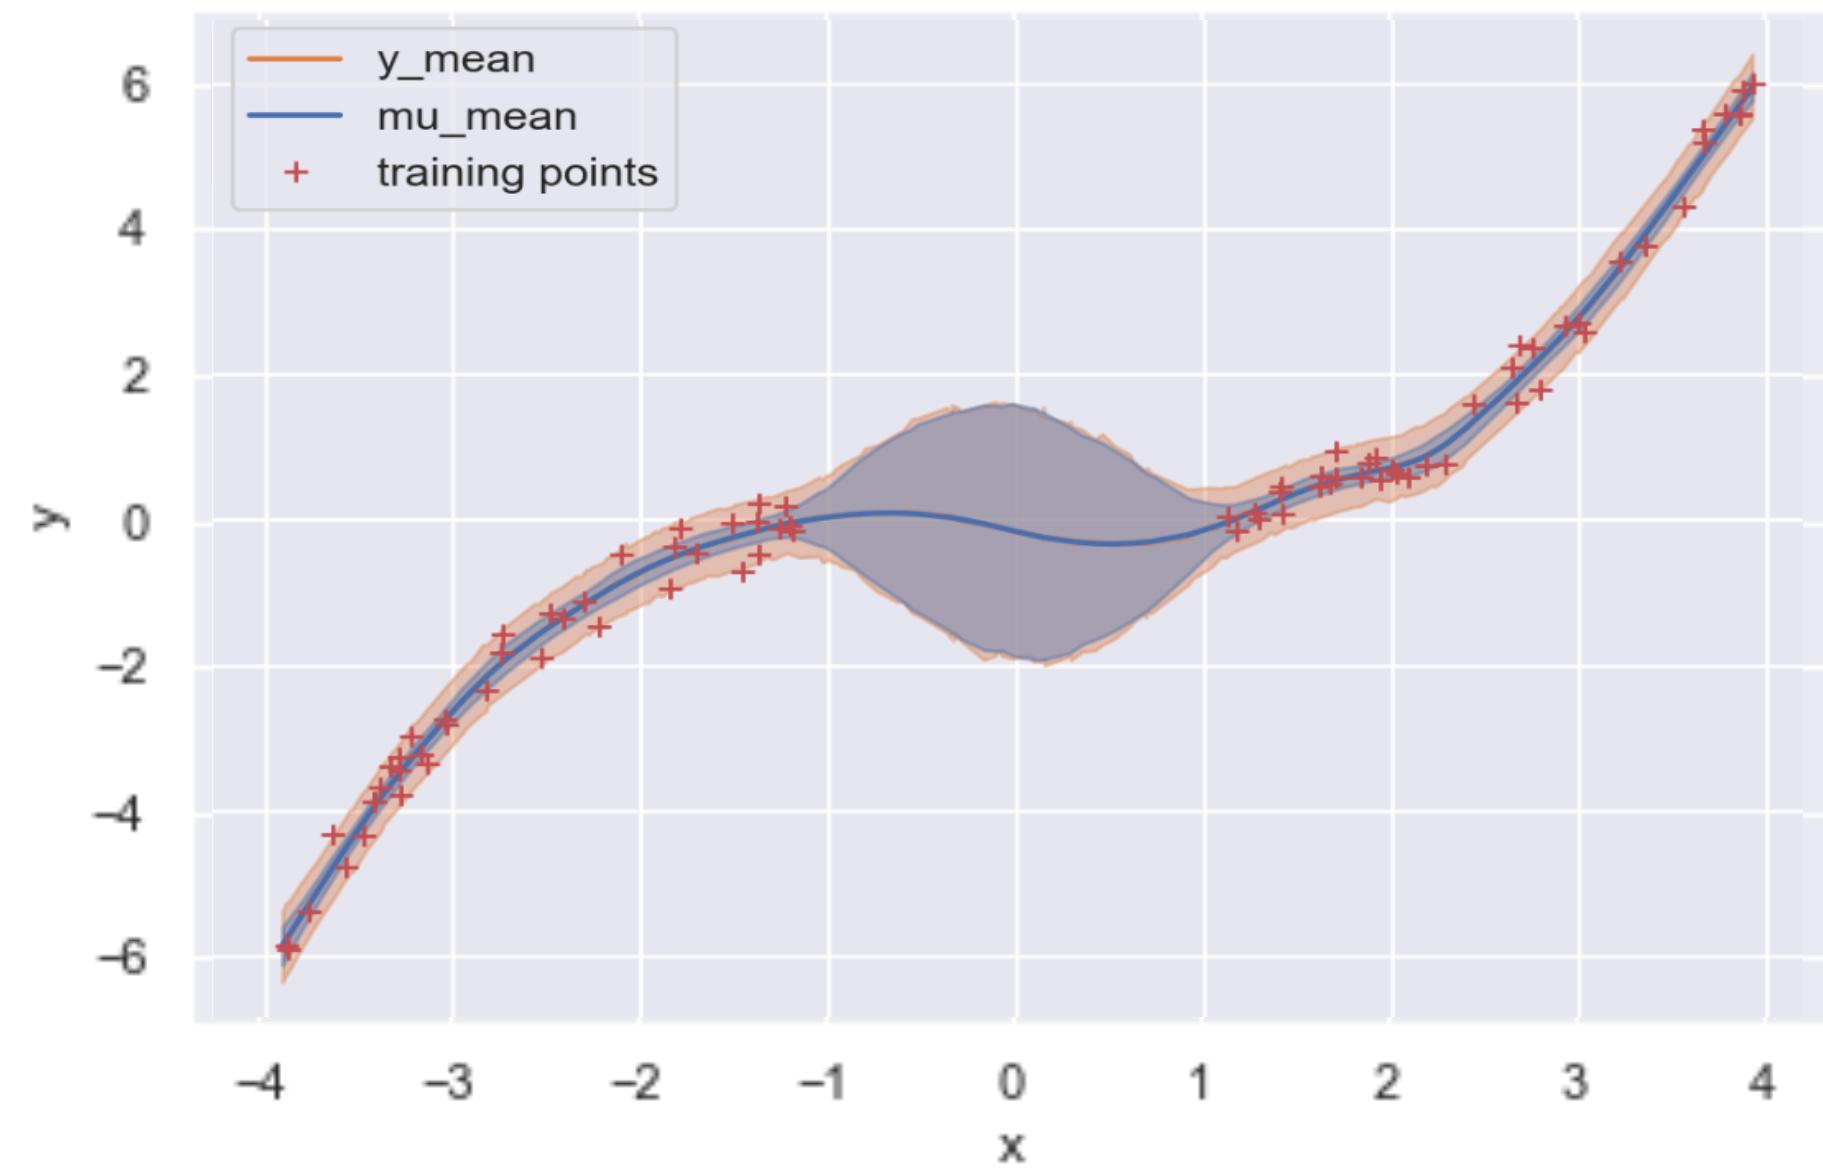


# Bootstrap: L models



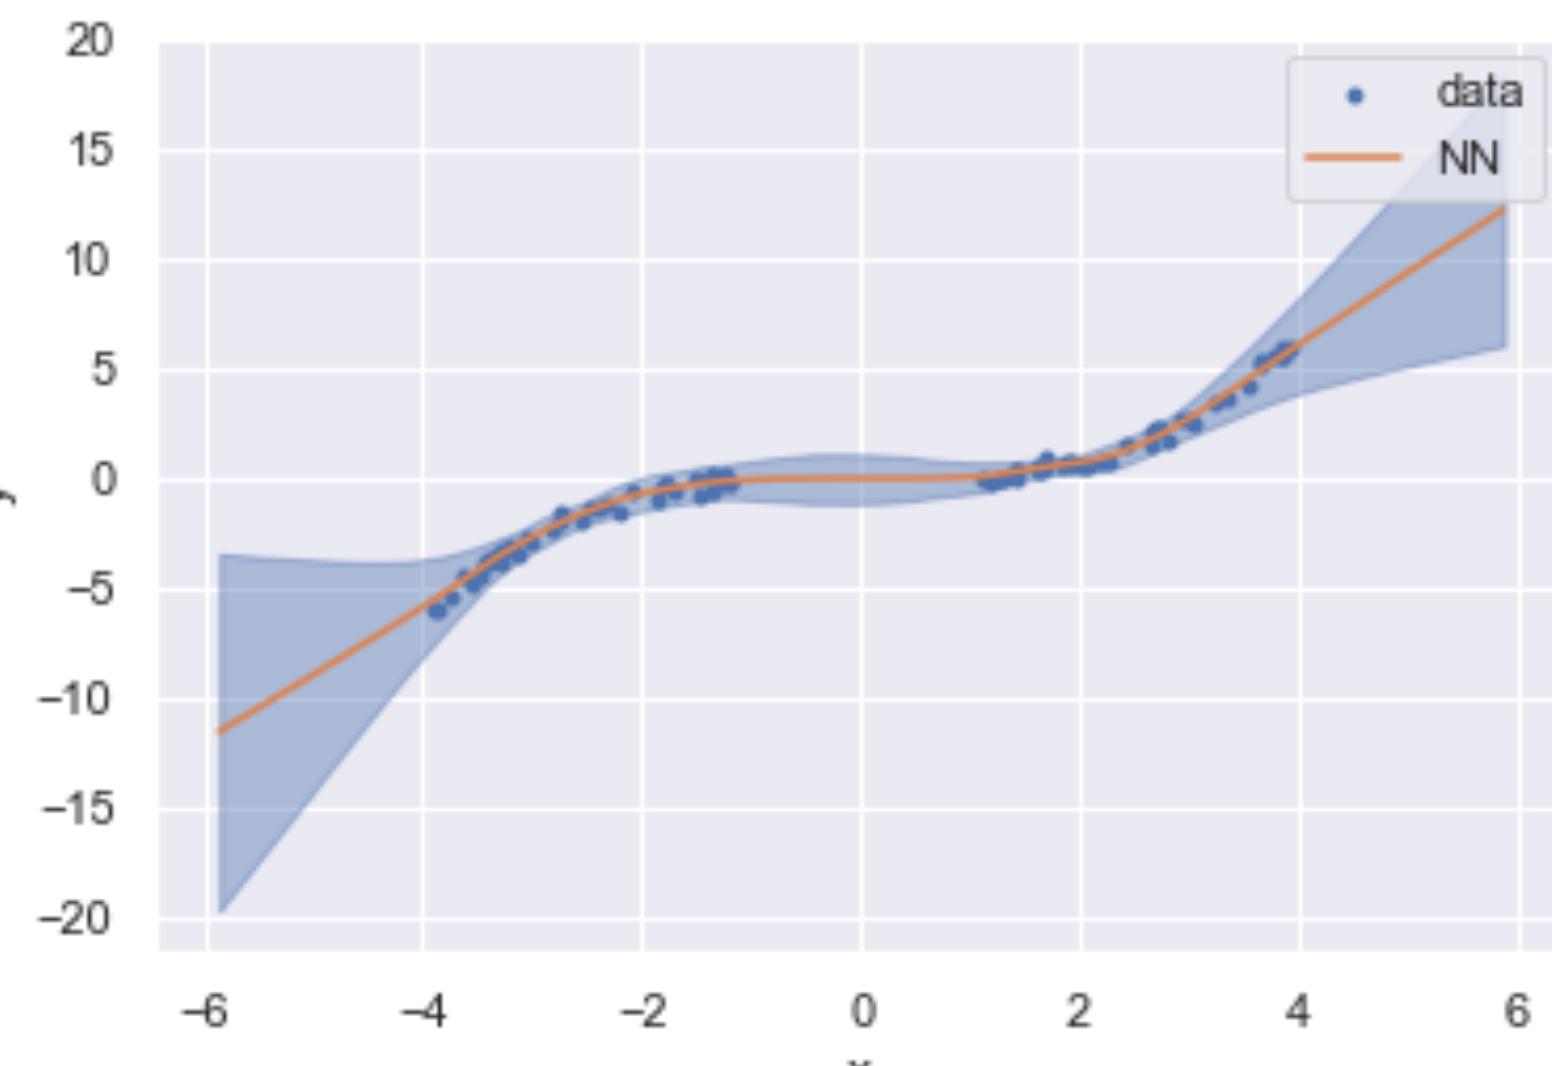
# Bootstrap: L models

MCMC

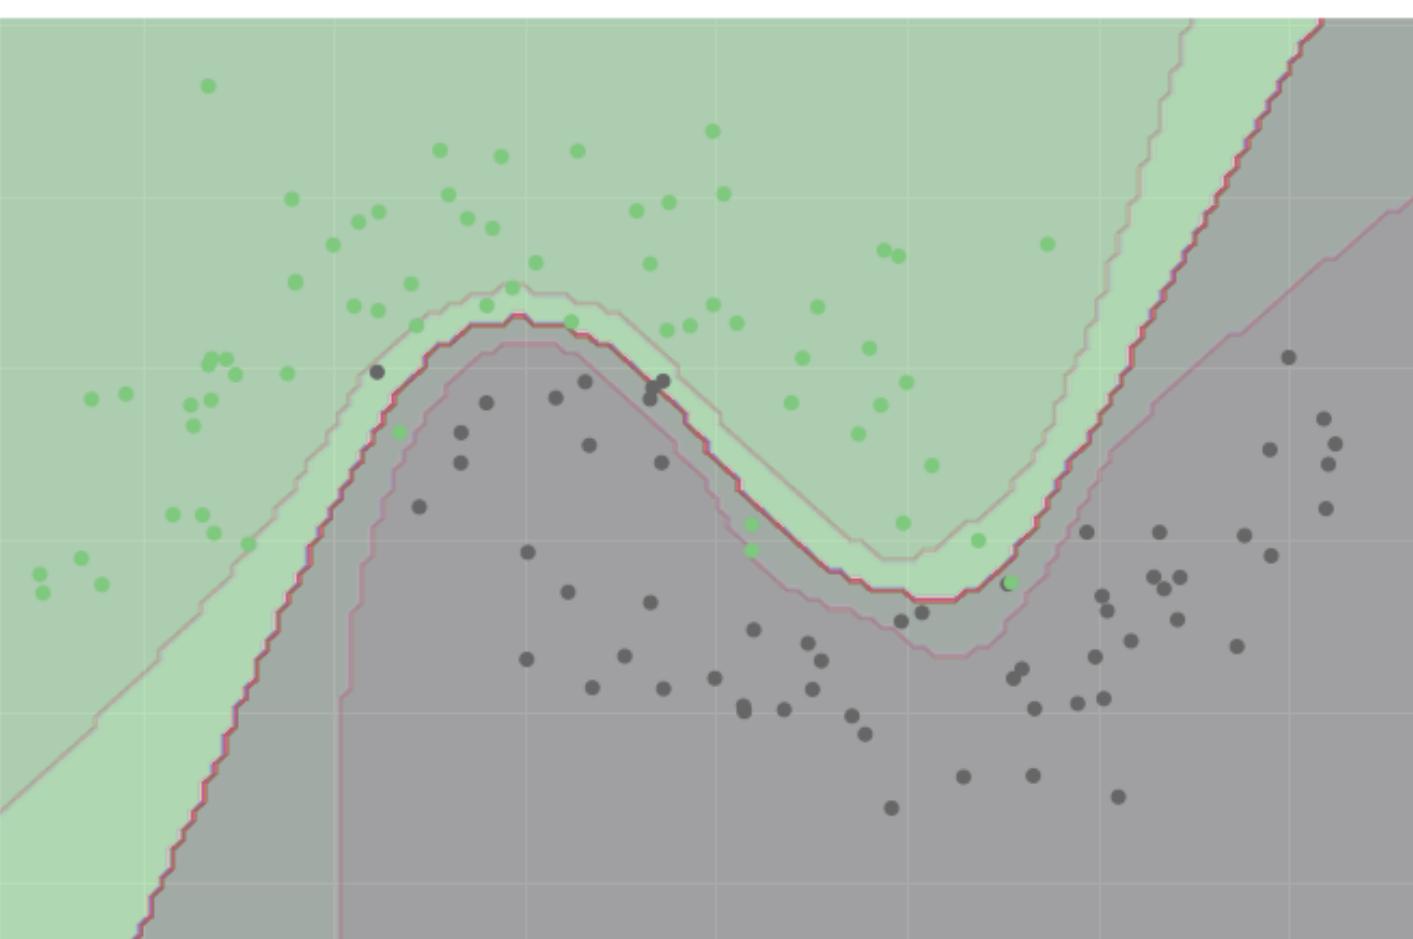


Bootstrap

Regression



Classification



Model Mean  
95% models

# Comparing Models of Uncertainty Quantification: Reading material

---

## **Quality of Uncertainty Quantification for Bayesian Neural Network Inference**

---

**Jiayu Yao<sup>\* 1</sup> Weiwei Pan<sup>\* 1</sup> Soumya Ghosh<sup>2</sup> Finale Doshi-Velez<sup>1</sup>**

<https://arxiv.org/pdf/1906.09686.pdf>

"Frequently in literature, high test log likelihood is used as evidence that the inference procedure has more faithfully captured the true posterior. However, here we argue that while test log likelihood may be a good criteria for model selection, it is not a reliable criteria for determining how well an approximate posterior aligns with the true posterior."

They compare 10 commonly used approximate inference procedures for Bayesian NNs. They find that approximate Bayesian inference methods typically do not capture true posteriors, and that non-Bayesian methods often do not capture the desired predictive uncertainty.

We need more careful metrics for evaluating the performance of methods on uncertainty quantification

The End

---

Questions?